

Homework 8: Spark Framework

Correct answers are labeled red bold.

Some necessary explanations are labeled blue.

Q1: What are the features of Spark RDD?

- A. In-memory computation
- B. Lazy evaluation
- C. Fault tolerance
- D. All of the above**

Q2: How many Spark Context can be active per JVM?

- A. Two
- B. One**
- C. Unlimited

(Only one SparkContext may be active per JVM. You must stop() the active SparkContext before creating a new one.)

Q3: Which of the following is not a transformation?

- A. flatMap(func) (flatMap(func) is a transformation operation and it's similar to map(func), but each input item can be mapped to 0 or more output items (so func should return a Seq rather than a single item).)
- B. map(func) (map(func) is a transformation operation and returns a new distributed dataset formed by passing each element of the source through a function func)
- C. reduce(func)** (reduce(func) is an action operation and it aggregates the elements of the dataset using a function func (which takes two arguments and returns one). The function should be commutative and associative so that it can be computed correctly in parallel.)
- D. filter(func) (filter(func) is a transformation and returns a new dataset formed by selecting those elements of the source on which func returns true.)

Q4: Which one of the following operations does NOT trigger an eager evaluation?

- A. take(n) (take(n) is an action operation and returns an array with the first n elements of the dataset.)
- B. collect() (collect() is an action and returns all the elements of the dataset as an array at the driver program. This is usually useful after a filter or other operation that returns a sufficiently small subset of the data.)
- C. count() (count() is an action and returns the number of elements in the dataset.)
- D. join(otherDataset, [numPartitions])** (join(otherDataset, [numPartitions]) is a transformation operation and won't trigger the evaluation.)

Q5: Which of the following commands will NOT generate a shuffle of data from each executor across the cluster?

- A. `collect()`
- B. `map(func)`**
- C. `repartition(numPartitions)`
- D. `distinct([numPartitions])`

(`map()` transformation is narrow and does not trigger a shuffle. The other options, such as `repartition()`, `distinct()` and `distinct()`, typically cause data shuffling across the cluster.)