

# Cloud Computing (INFS3208)

## Lecture 3: Docker I: Container, Docker Commands, and Dockerfile

Lecturer: Dr Heming Du

School of Electrical Engineering and Computer Science

Faculty of Engineering, Architecture and Information Technology

The University of Queensland

# Re-cap Lecture 2

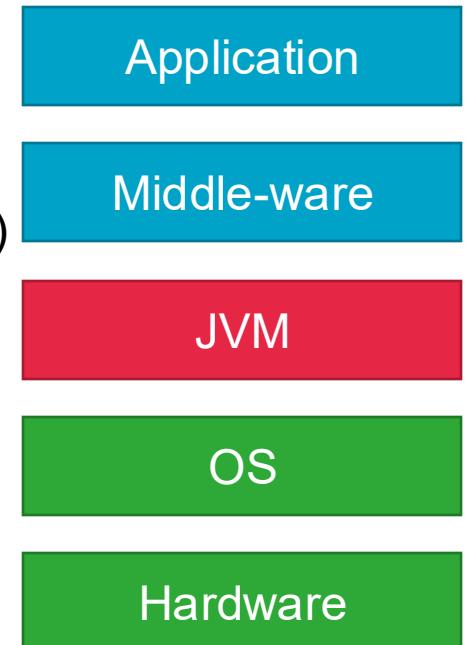
- Networking and Virtual Private Cloud
- Load Balancing
  - What & Why Load Balancing
  - Algorithms
  - LB in Cloud Architecture
  - LB in Distributed Systems
  - LB in Network Communications
  - LB in Cloud Product

# Outline

- • Container
  - What is Docker
  - Basic concepts in Docker
    - Images
    - Layer architecture
    - Containers
    - Registry
  - Docker Commands
  - Containerisation and Dockerfile
    - Dockerfile instructions

# Revisit: Virtual Machines (VMs)

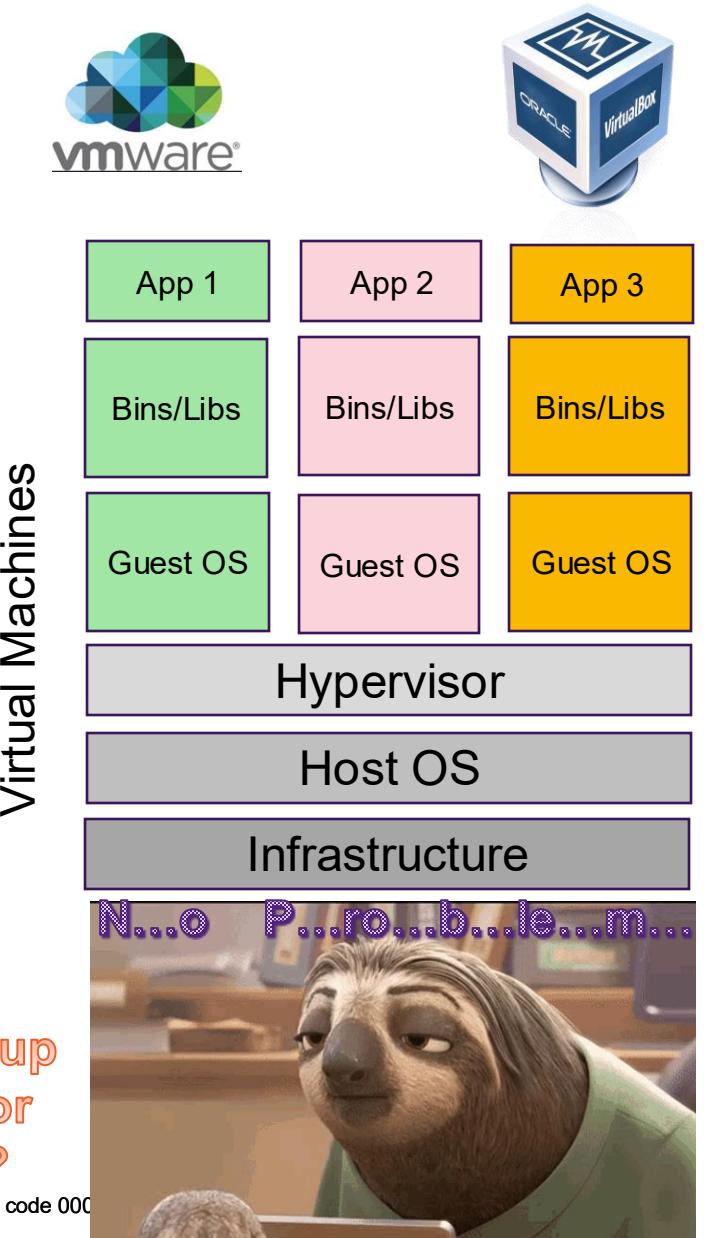
- A virtual machine (VM) is an **emulation of a computer** system.
- Virtual machines are based on **computer architectures** and provide functionality of a physical computer: software, specialised hardware, or a combination.
- **System virtual machines**: provide a substitute for a real machine.
- **Process virtual machines**
  - are designed to execute computer programs in a platform-independent environment.
  - provide a high-level abstraction (a high-level programming language abstraction)
  - Example: Java virtual machine (JVM)



# Virtual Machines (VM)

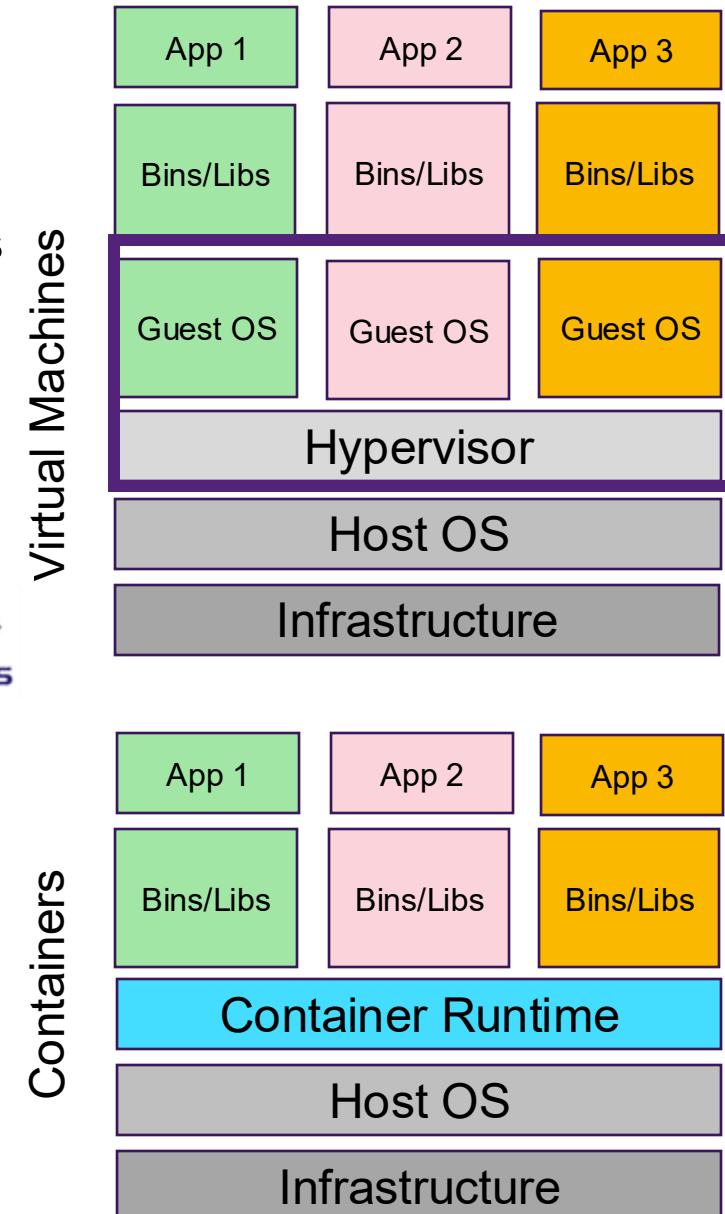
- VM could simulate almost anything:
  - Hardware (vCPU, vRAM, vHDisk, vGPU, vNetwork, etc.)
  - Software (Operating System: Ubuntu, CentOS, Windows Server)
- Environment **Isolation** and **Consistency**:
  - Isolation: VMs are **isolated** envs and **independent** to the host system (No affect).
  - Consistency: dev can create a **consistent** env. (Code MUST behaves consistently).
- **But...**
  - Each VM requires its own dedicated OS; Extra hardware and software overhead
  - Performance of VM is not very satisfied (slow to boot up)
  - Migration between physical servers is not always smooth.

Can we boot up  
more VMs for  
scalability?



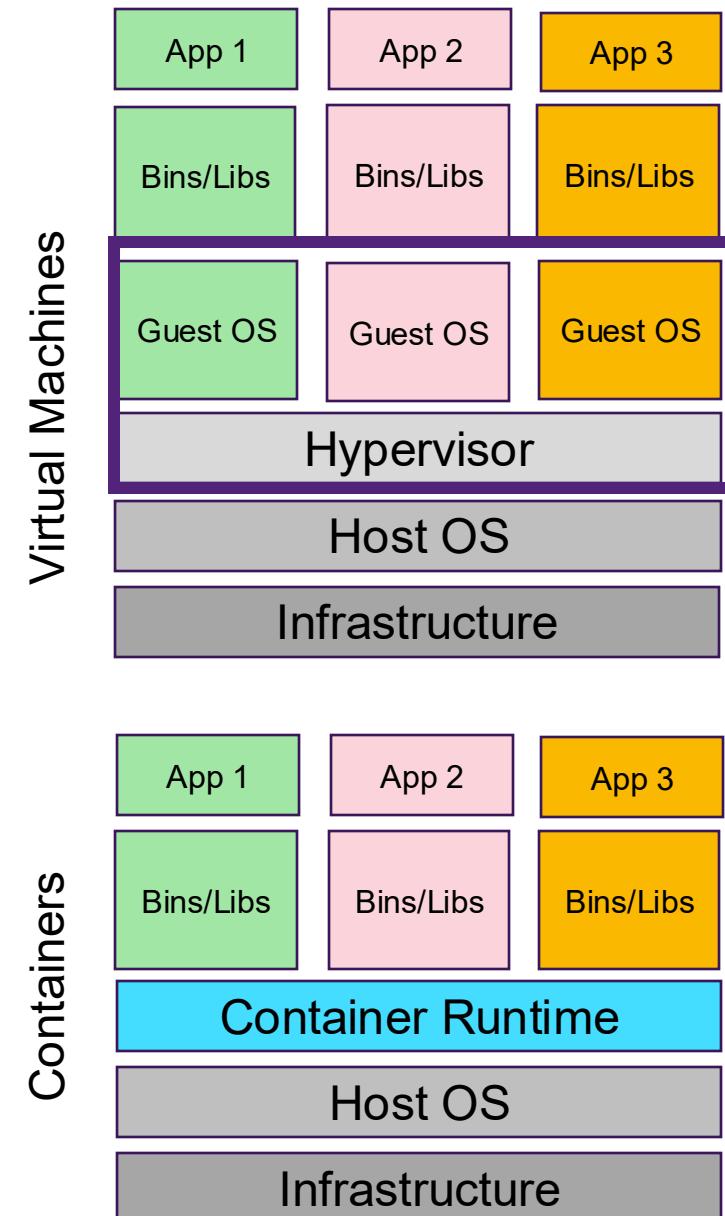
# Containers

- **Containers** offer a **logical packaging mechanism** in which applications can be abstracted from the environment in which they actually run.
- Completely isolated environments: services, network interfaces, mounts
- Multiple containers running atop the **OS kernel** directly – far more **lightweight**:
  - share the OS kernel
  - start much faster (**in seconds**)
  - use **a fraction of** the memory (compared to booting an entire OS)
  - one physical machine/server can run much **more containers** than VMs
- Running Linux Containers on Windows (not easy but possible)
  - Windows Subsystem for Linux (WSL)
  - Docker Desktop (helps you build, share, and run containers on Mac and Windows as you do on Linux)



# The History of Containers

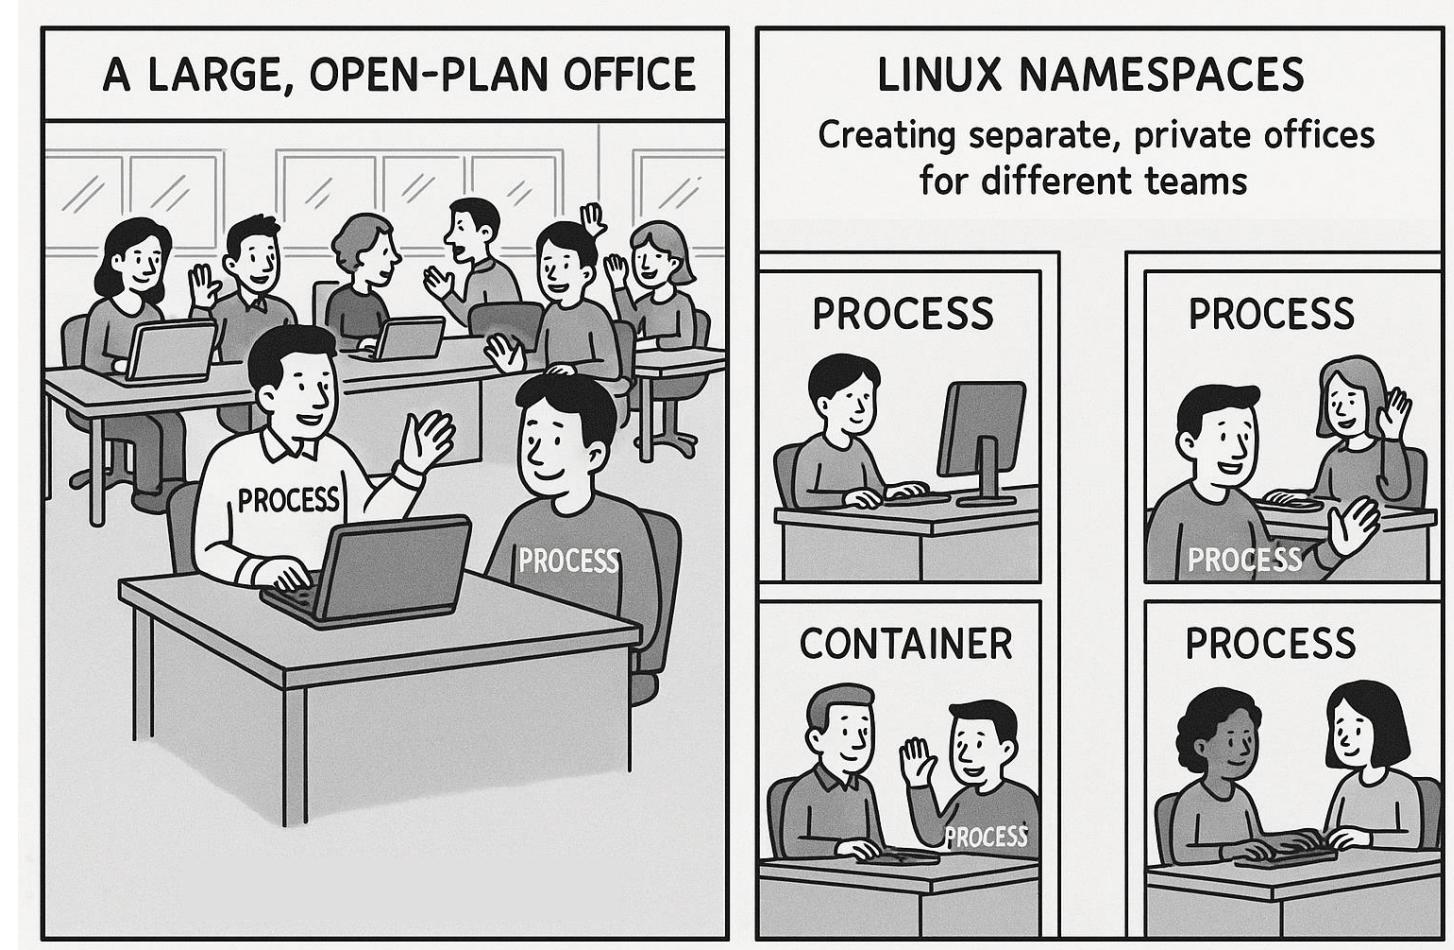
- Before container:
  - Chroot system call in 1979 and FreeBSD jails in 2000.
  - Cgroups (2006, 2007) and namespaces (2002)
  - Linux Containers (LXC) in 2008
- Containers virtualise at the operating system level
- Containerisation provides a clean separation of concerns:
  - developers focus on their application logic and dependencies
  - IT operations teams focus on deployment and management without bothering with application details (e.g. specific software versions and apps configurations).



<https://www.redhat.com/en/blog/history-containers>  
<https://cloud.google.com/containers/>

# Linux Namespaces: Isolation of System Resources

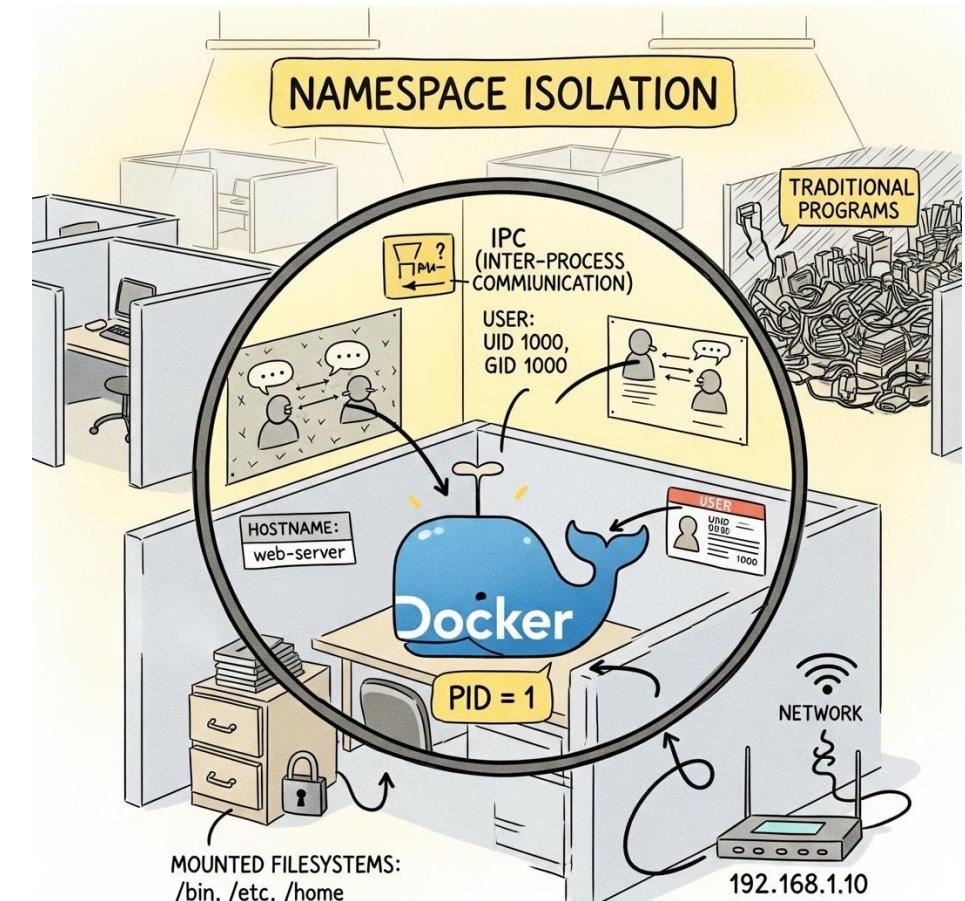
- Namespaces are a feature of the Linux kernel that partition kernel resources such that one set of processes sees one set of resources, while another set of processes sees a different set.
- Namespaces provides a process with its own isolated view of the system.



# Linux Namespaces: Isolation of System Resources

Namespaces are utilized to give a container the illusion that it's running on its own, separate machine. Key namespaces used by containers include:

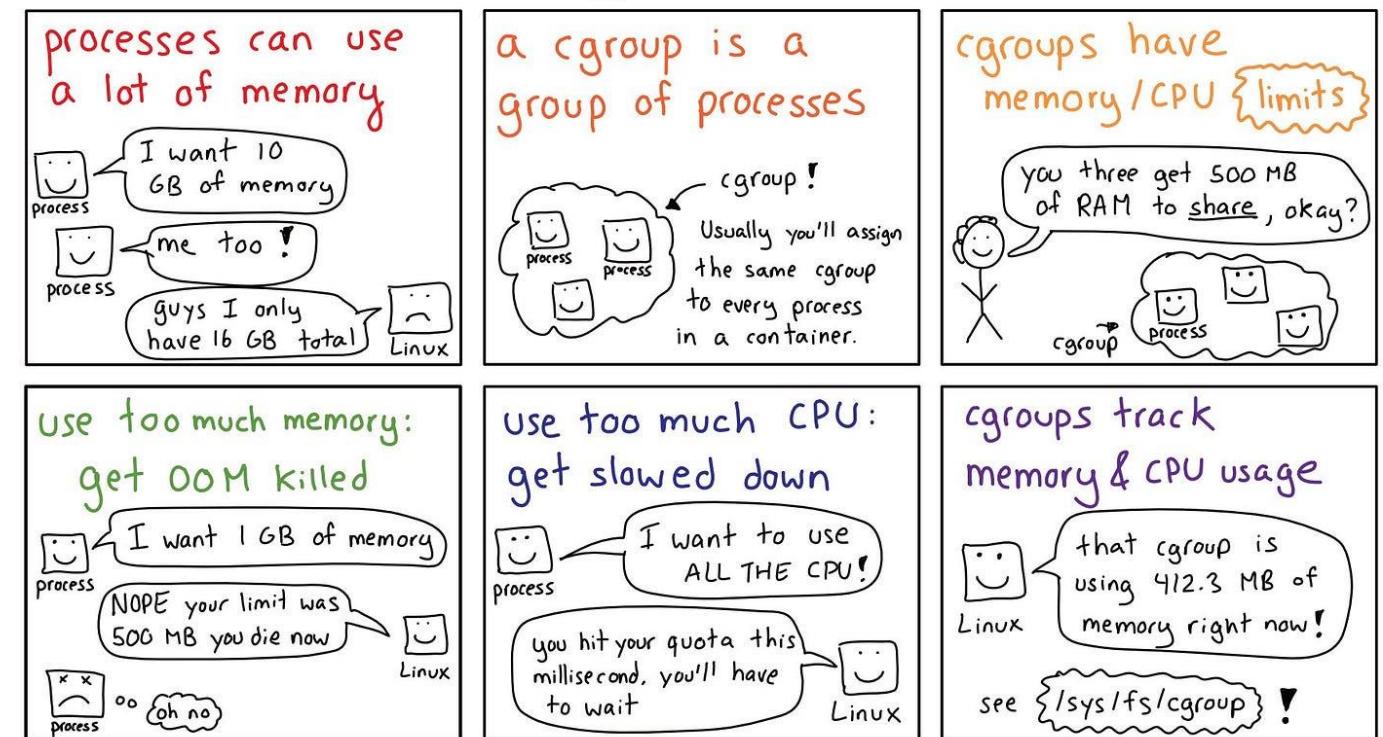
- PID (Process ID): Isolates the process tree. Inside a container, the main process has a PID of 1, but on the host machine, it has a different PID.
- NET (Network): Provides each container with its own network stack (IP address, routing tables, etc.).
- MNT (Mount): Gives each container its own filesystem.
- UTS (UNIX Time-sharing System): Allows each container to have its own hostname.
- IPC (Inter-Process Communication): Isolates IPC resources.
- USER: Isolates user and group IDs.



# Managing Resources with Control Groups (Cgroups)

- Control Groups (cgroups) are a Linux kernel feature that limits, accounts for, and isolates the resource usage (CPU, memory, disk I/O, network, etc.) of a collection of processes.
- While namespaces handle what a container can see, cgroups control what a container can use.
- Cgroups prevent a single container from consuming all of the host's resources and impacting other containers or the host system itself.

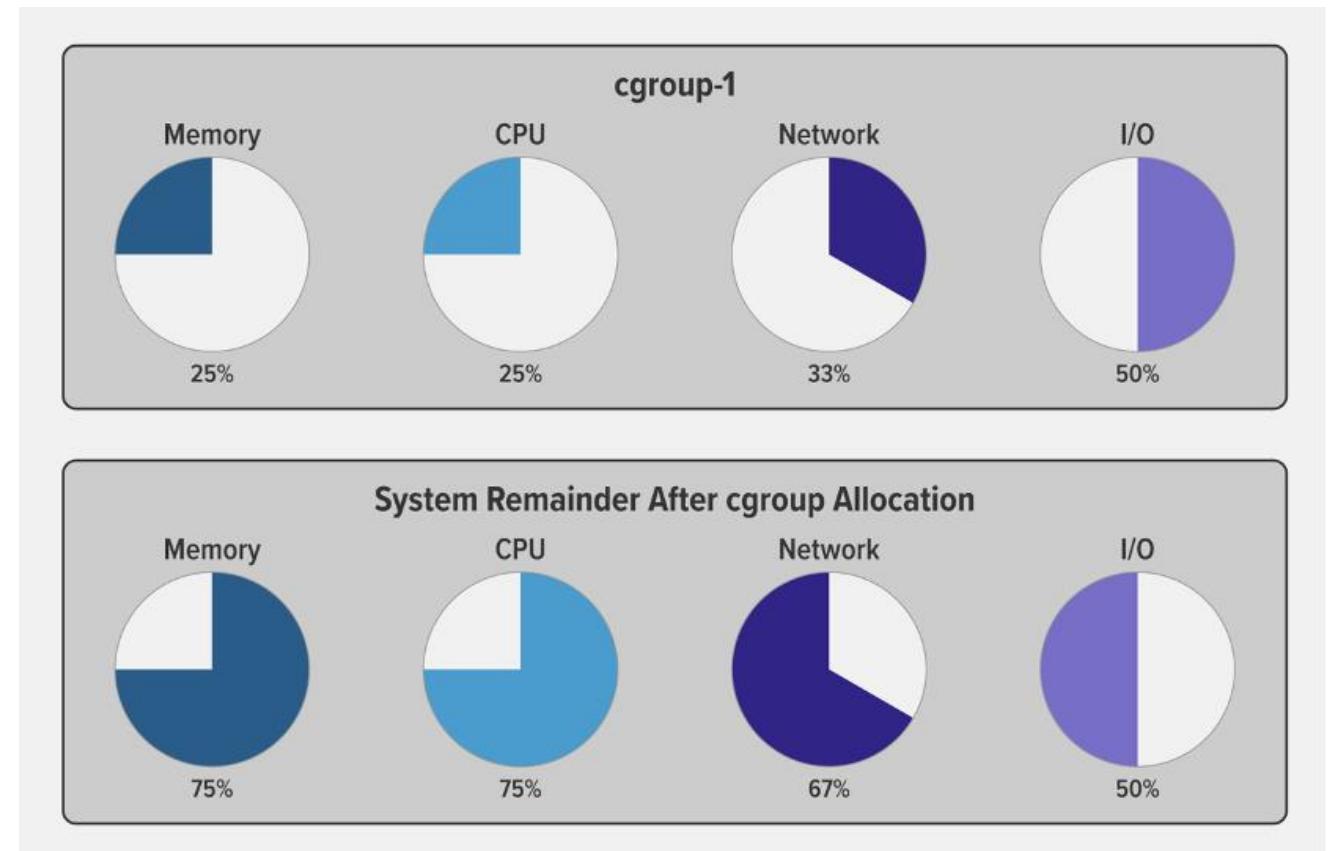
JULIA EVANS  
@b0rk



<https://medium.com/@boutnaru/linux-cgroups-control-groups-part-1-358c636ffde0>

# Managing Resources with Control Groups (Cgroups)

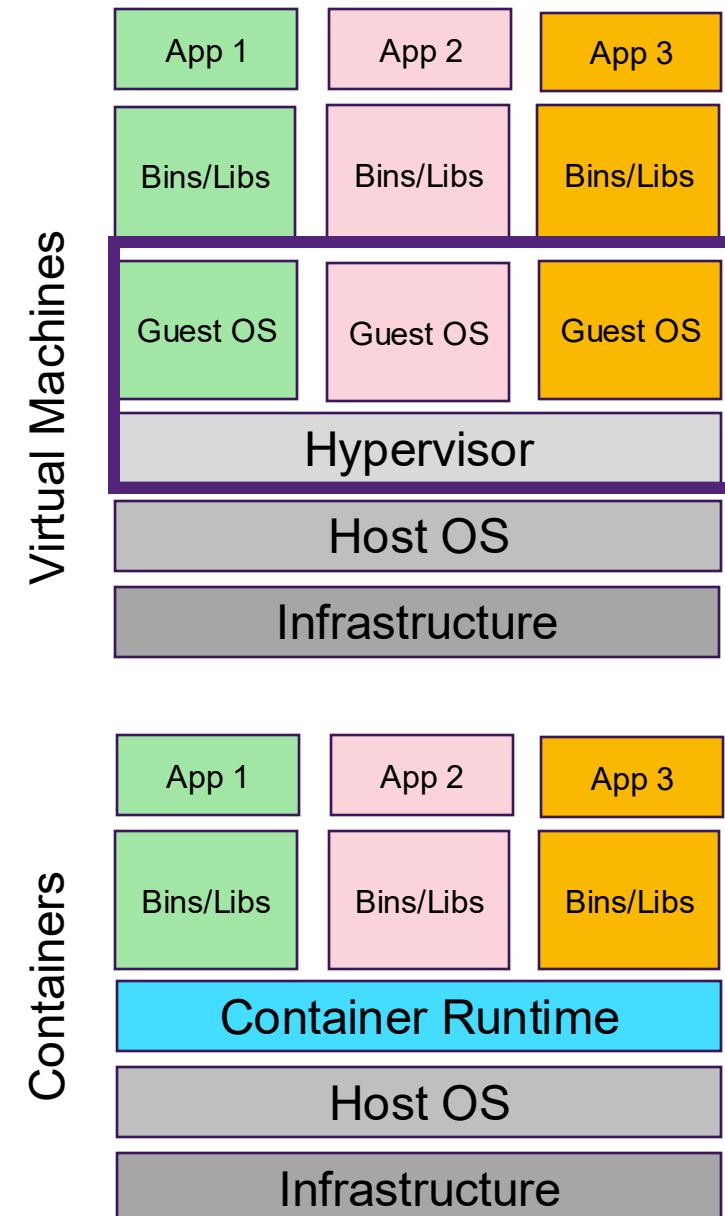
- Resource Limiting: You can set hard limits on the amount of memory or CPU a container can use.
- Prioritization: You can allocate more or fewer resources to certain containers.
- Accounting: Cgroups monitor and report on the resource usage of containers.
- Control: You can freeze and resume all processes within a cgroup.



<https://blog.nginx.org/blog/what-are-namespaces-cgroups-how-do-they-work>

# The History of Containers

- Before container:
  - Chroot system call in 1979 and FreeBSD jails in 2000.
  - Cgroups (2006, 2007) and namespaces (2002)
  - Linux Containers (LXC) in 2008
- Containers virtualise at the operating system level
- Containerisation provides a clean separation of concerns:
  - developers focus on their application logic and dependencies
  - IT operations teams focus on deployment and management without bothering with application details (e.g. specific software versions and apps configurations).



<https://www.redhat.com/en/blog/history-containers>  
<https://cloud.google.com/containers/>

# Why Containers?

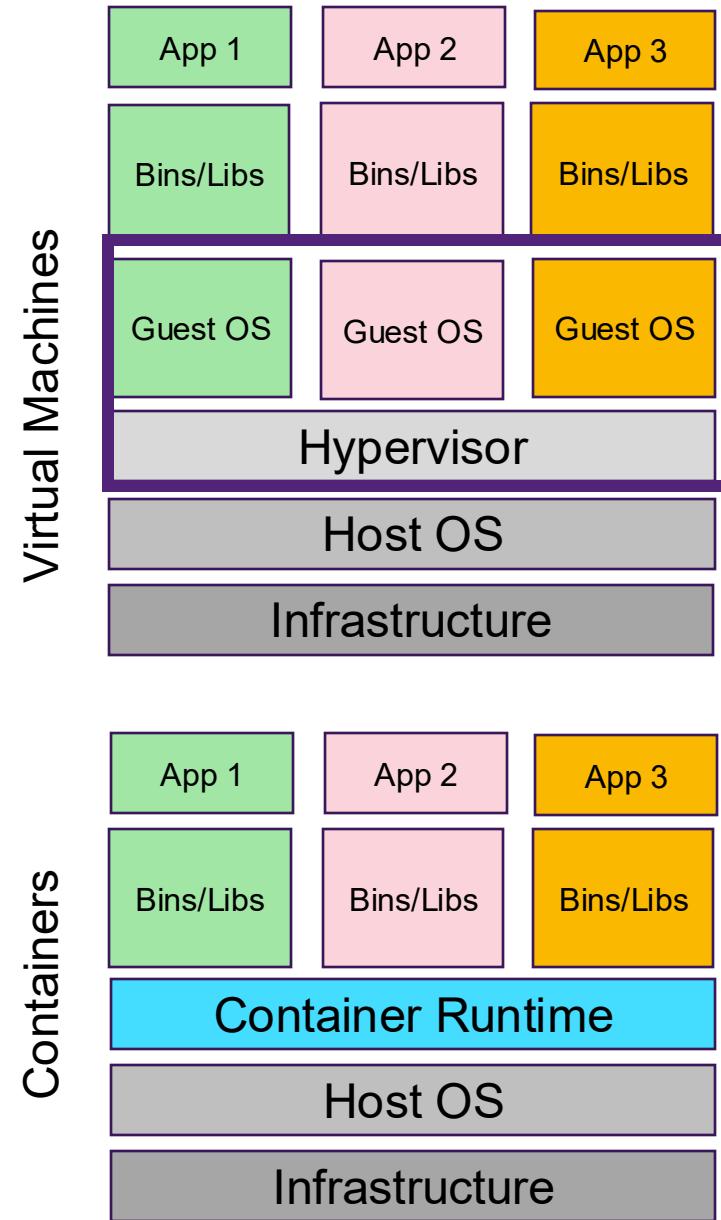
- **Run Anywhere:**

- Containers are able to run virtually anywhere, greatly easing development and deployment:
  - on Linux, Windows, and Mac operating systems;
  - on virtual machines;
  - on a developer's machine;
  - in data centres on-premises;
  - in the public cloud.

- **Isolation:**

- Containers virtualise hardware at the OS-level:
  - CPU,
  - memory, storage,
  - network resources.
- Containers provide developers with a sandboxed view of the OS logically isolated from other applications.

<https://cloud.google.com/containers/>



# Containers vs. Virtual Machines

Features	Containers	Virtual Machines
<b>Isolation</b>	<b>Process-level</b> isolation using namespaces and cgroups	<b>Full OS-level</b> isolation with a separate guest OS
<b>Resource Efficiency</b>	<b>High</b> , due to sharing of the host OS kernel	<b>Lower</b> , as each VM runs a full OS
<b>Startup Time</b>	<b>Very fast</b> (seconds or less)	<b>Slower</b> (minutes)
<b>Performance</b>	<b>Near-native</b> performance	<b>Overhead</b> from hypervisor and full OS
<b>Storage</b>	Uses <b>layered file systems</b> (e.g., overlayfs)	Requires dedicated disk space for each VM
<b>Deployment</b>	<b>Lightweight</b> and quick to deploy	<b>Heavier</b> and slower to deploy
<b>Scalability</b>	<b>Easily scalable</b>	<b>Less flexible</b> , more resource-intensive to scale
<b>Portability</b>	<b>High portability</b> (consistent environments)	<b>Moderate portability</b> (more dependencies)
<b>Security</b>	<b>Less secure</b> than VMs (shared kernel vulnerabilities)	<b>More secure</b> (isolated OS kernel)
<b>Use Cases</b>	Microservices, DevOps, CI/CD, lightweight apps	Legacy applications, full OS environments, high-security isolation
<b>Management Tools</b>	Docker, Kubernetes	VMware, Hyper-V, VirtualBox
<b>Overhead</b>	<b>Minimal</b> overhead (shares host kernel)	<b>Significant</b> overhead (runs separate OS for each VM)

# Outline

- Container
- • What is Docker
- Basic concepts in Docker
  - Images
  - Layer architecture
  - Containers
  - Registry
- Docker Commands
- Containerisation and Dockerfile
  - Dokcerfile instructions

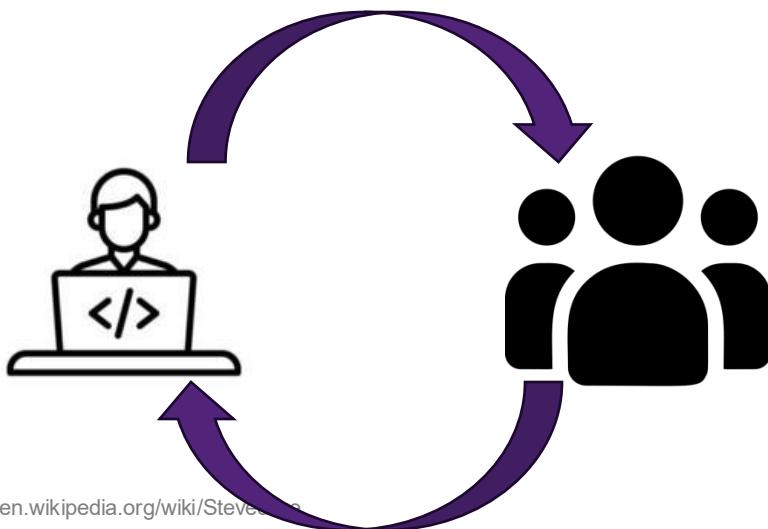
# What is a “Docker”?



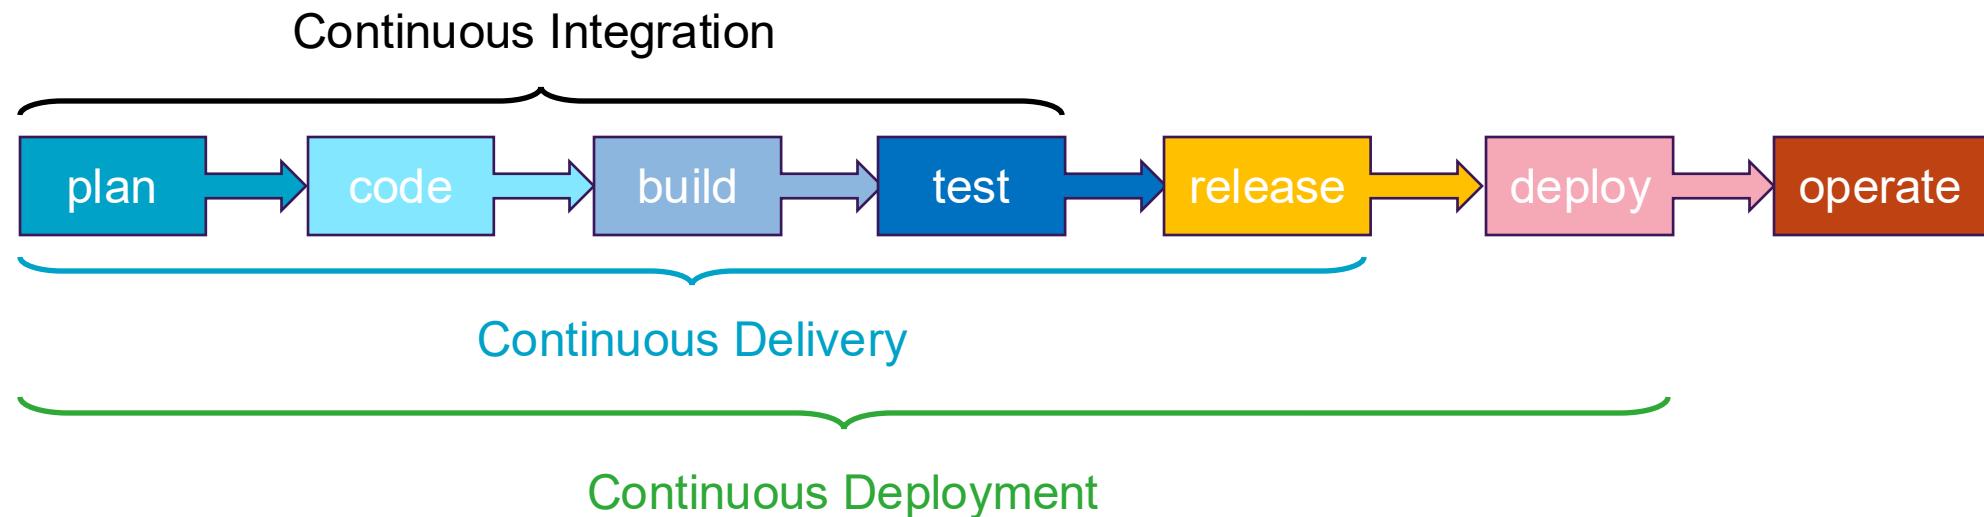
**stevedore** (*/'stiːvə,dɔːr/*), **longshoreman**, **docker**, or **dockworker**

# What is a “Docker”

- Cargo can be quickly packed at source of production
- Cargo can be quickly shipped from A to B
- Cargo can be quickly unpacked at destination
- Cargo is what it is after shipping



# Modern Development Practices - DevOps



- Continuous Integration (CI): automated testing
- Continuous Delivery (CD): automated release
- Continuous Deployment: fully automated

**Needs Isolation & Consistency**  
**Lightweight**

# What is Docker

- Docker is
  - a job (literally),
  - a company providing solutions to containerisation (Google's LXD),
  - a new implementation of container technologies.
- Some definitions:
  - Docker is a container image that is a **lightweight, stand-alone, executable** package of a piece of software that includes everything needed to run it: code, runtime, system tools, system libraries, settings.
  - Docker is an open platform for **developing, shipping, and running applications** (official documentation). Simply, Docker is a tool to deploy applications in a **sandbox** (called containers) to run on the host operating system i.e. Linux, Windows, Mac OSX, etc.

<https://www.docker.com/what-container>

[https://en.wikipedia.org/wiki/Docker\\_\(software\)](https://en.wikipedia.org/wiki/Docker_(software))



<https://thenewstack.io/solomon-hykes-departs-from-docker/>

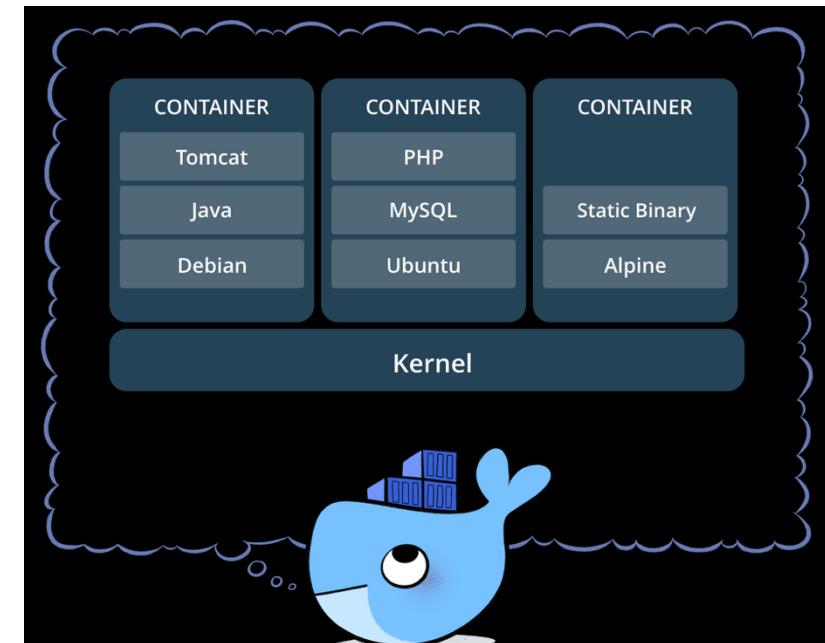
Docker founders  
Kamel Founadi,  
Solomon Hykes,  
Sebastien Pahl

# Outline

- Container
- What is Docker
- Basic concepts in Docker
  - Images
  - Containers
  - Registry
  - Layer architecture
- Docker Commands
- Containerisation and Dockerfile
  - Dockerfile instructions

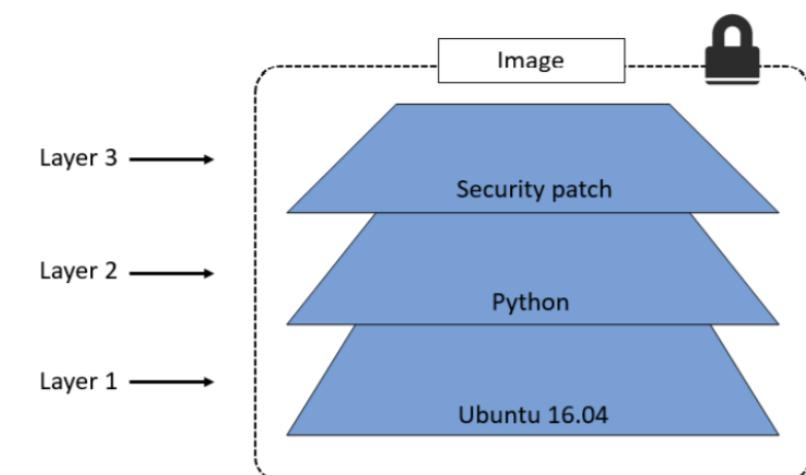
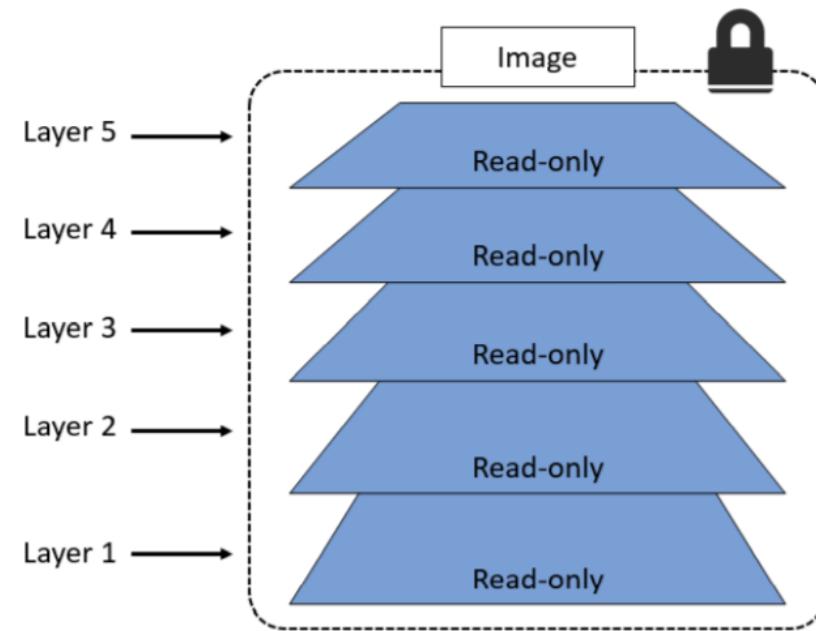
# Basic Concepts - Docker Image

- Docker image is an object that contains
  - an [OS filesystem](#)
  - and one or more [applications](#)
- An image is a template that is [readable only](#) to build a container.
- An image is based on another image ([base image](#)) with some additional customisation
  - Example: an image that contains ubuntu 22.04 (LTS) operating system and apache web server
- Docker provides a very good mechanism to build images and can even pull the existing images.
- There are heaps of resources on Docker Hub.
- Many software is released as Docker images, instead of software packages.



# Images and Layers

- With Union FS (Union Filesystems), Docker is designed as a **layer-wise** architecture.
- Union FS: create a union of (1) a list of lower directories – read only (2) upper directory – reads and writes
- Each image will not be changed (**read-only**) after it has been constructed – lower directories.
- Layer-wise architecture makes **reuse** and **customization** of images much easier.
  - Construct a new image based on the existing images by adding new layers on top of the current layers.
  - E.g. Ubuntu + Python + security patch

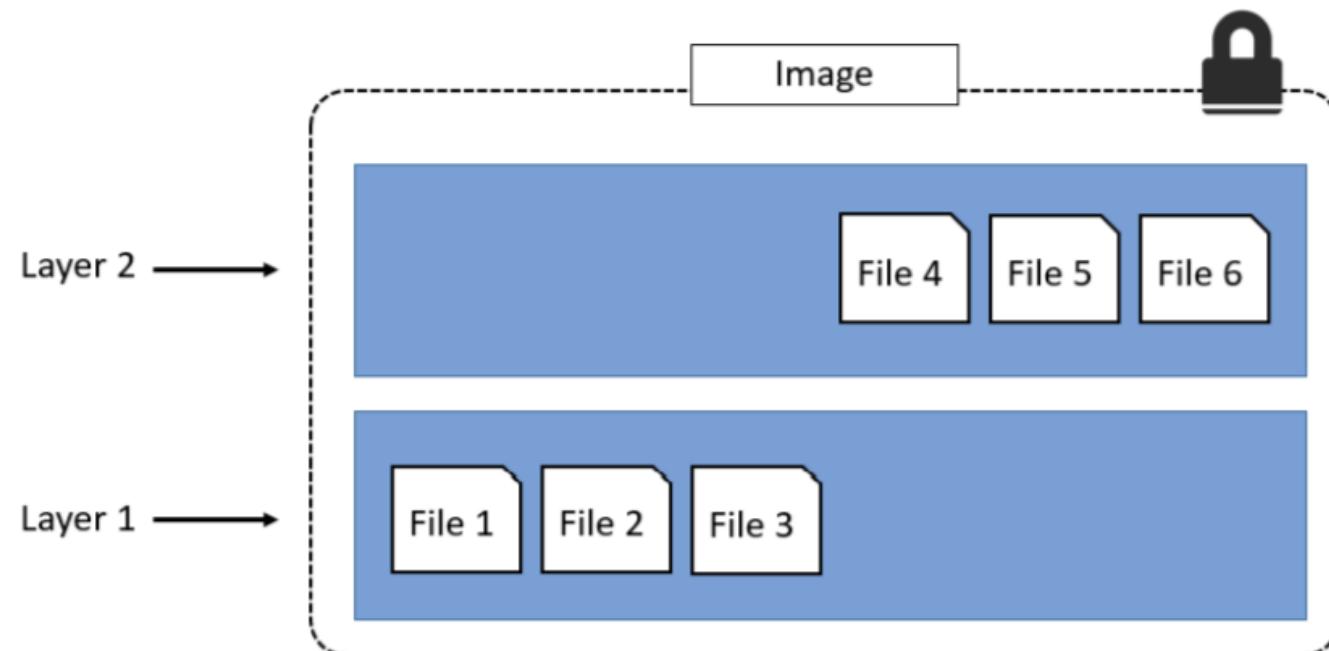


<https://docs.docker.com/storage/storagedriver/>

<https://blog.netapp.com/blogs/containers-vs-vms/>

# Copy-on-write Strategy: Example of Updating Files

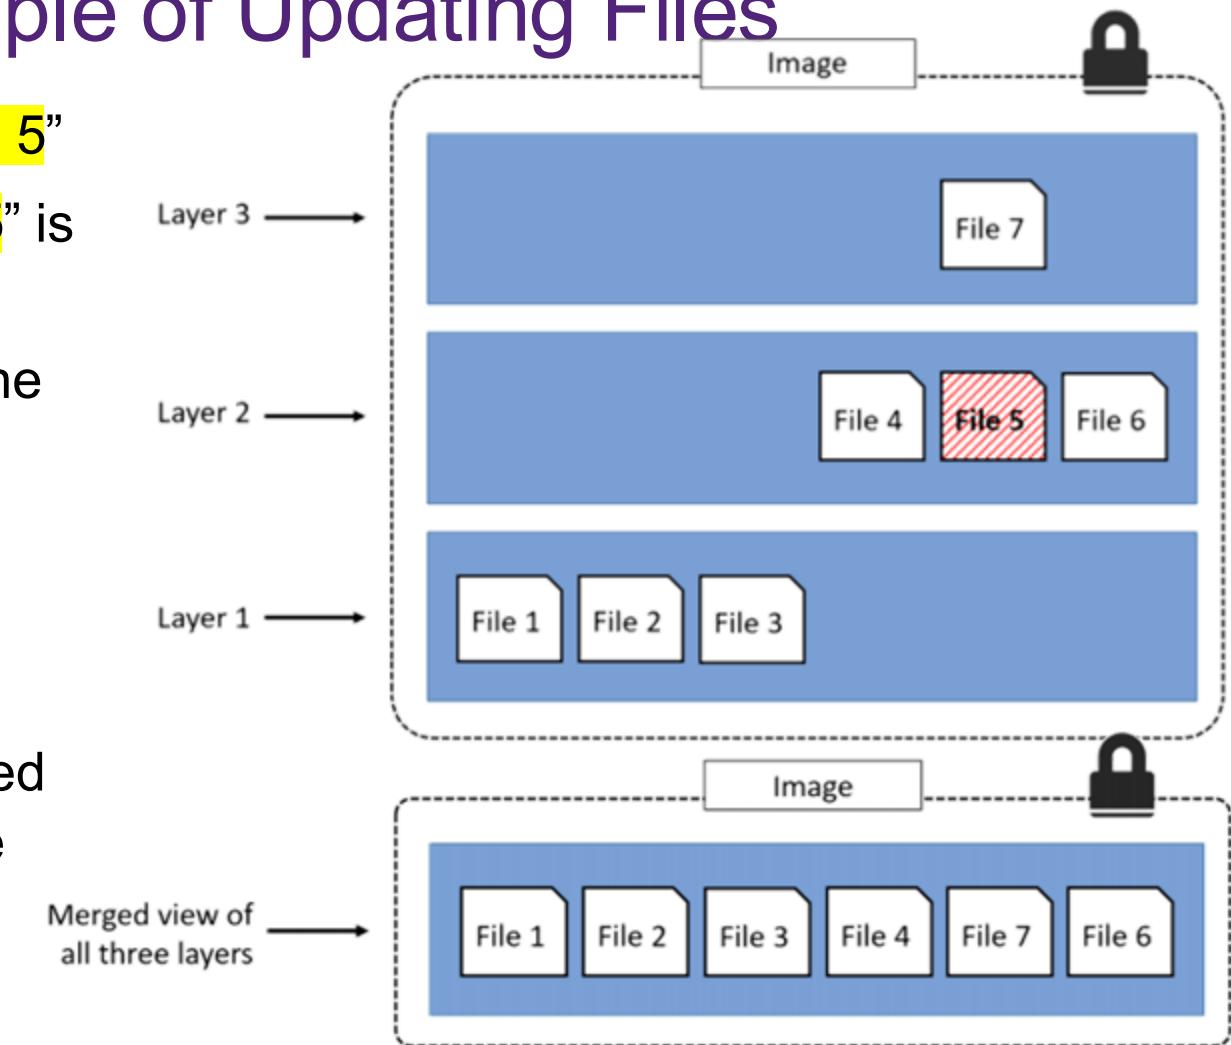
- Copy-on-write is a strategy of sharing and copying files for maximum efficiency.
- Example: after constructing two layers, three files in Layer 1 and three files in Layer 2.
- What will happen if we want to update “File 5” in Layer 2



<https://docs.docker.com/storage/storagedriver/#the-copy-on-write-cow-strategy>

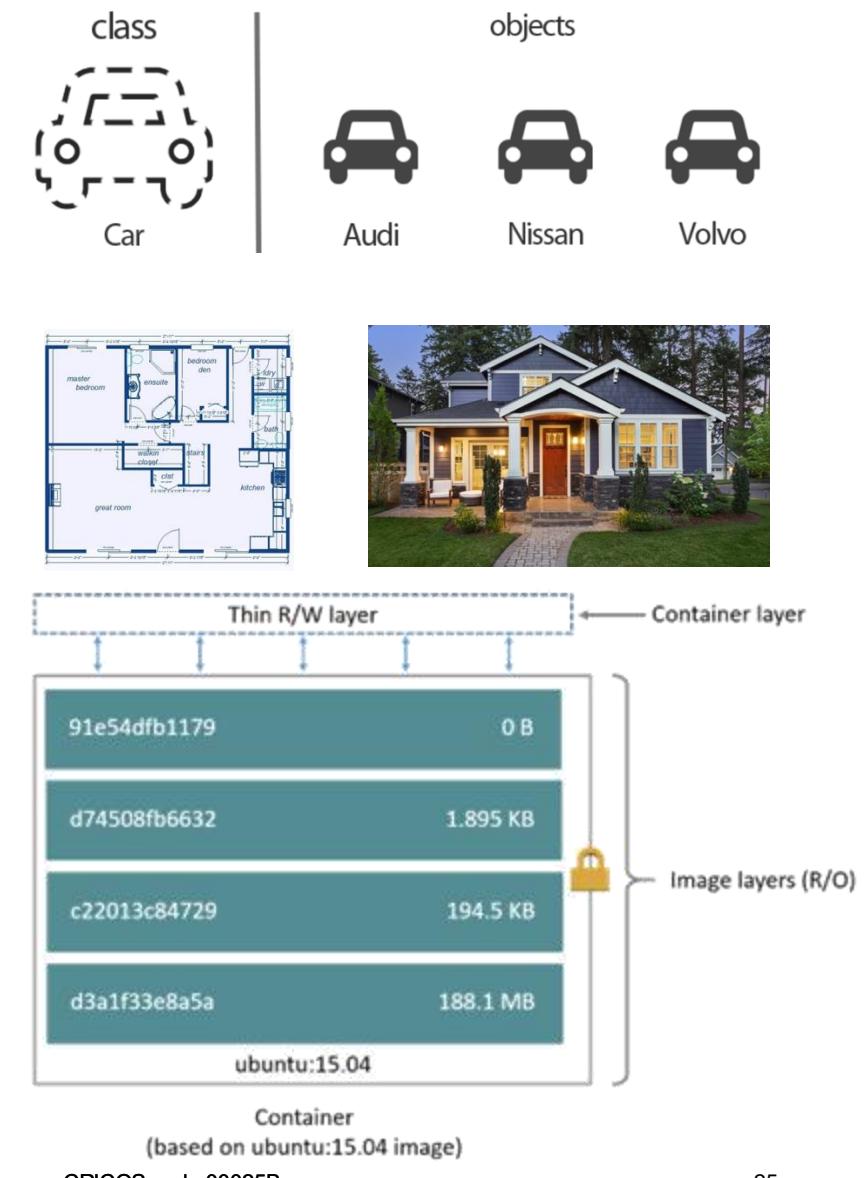
# Copy-on-write Strategy: Example of Updating Files

- “File 7” in layer 3 is an updated version of “File 5”
- Due to read-only property of each layer, “File 5” is still stored in the image.
- “File 5” is updated as a new layer added into the image.
- Good practice:
  - only **add** required files or apps in the current layer
  - Temporary or redundant files must be removed from the current layer **before** constructing the next layer
  - Very important when writing dockerfiles.



# Basic Concepts - Container

- Container: a running instance of an image
  - Class vs Instances (in Objective Oriented Programming)
  - E.g. House Blueprints vs House
- Container is a running process
  - Independent namespace
    - E.g. variable `$JAVA_HOME` in Container A is different from `$JAVA_HOME` in Container B
  - Independent root filesystem
    - E.g. `/usr/local` in Container A is different from `/user/local` in Container B
  - Independent networks
    - E.g. `192.168.0.1` in Container A is different from `192.168.0.1` in Container B
  - Independent user space
    - E.g. user “John” in Container A is different from “John” in Container B, even user root is different.



# Basic Concepts - Docker Registry

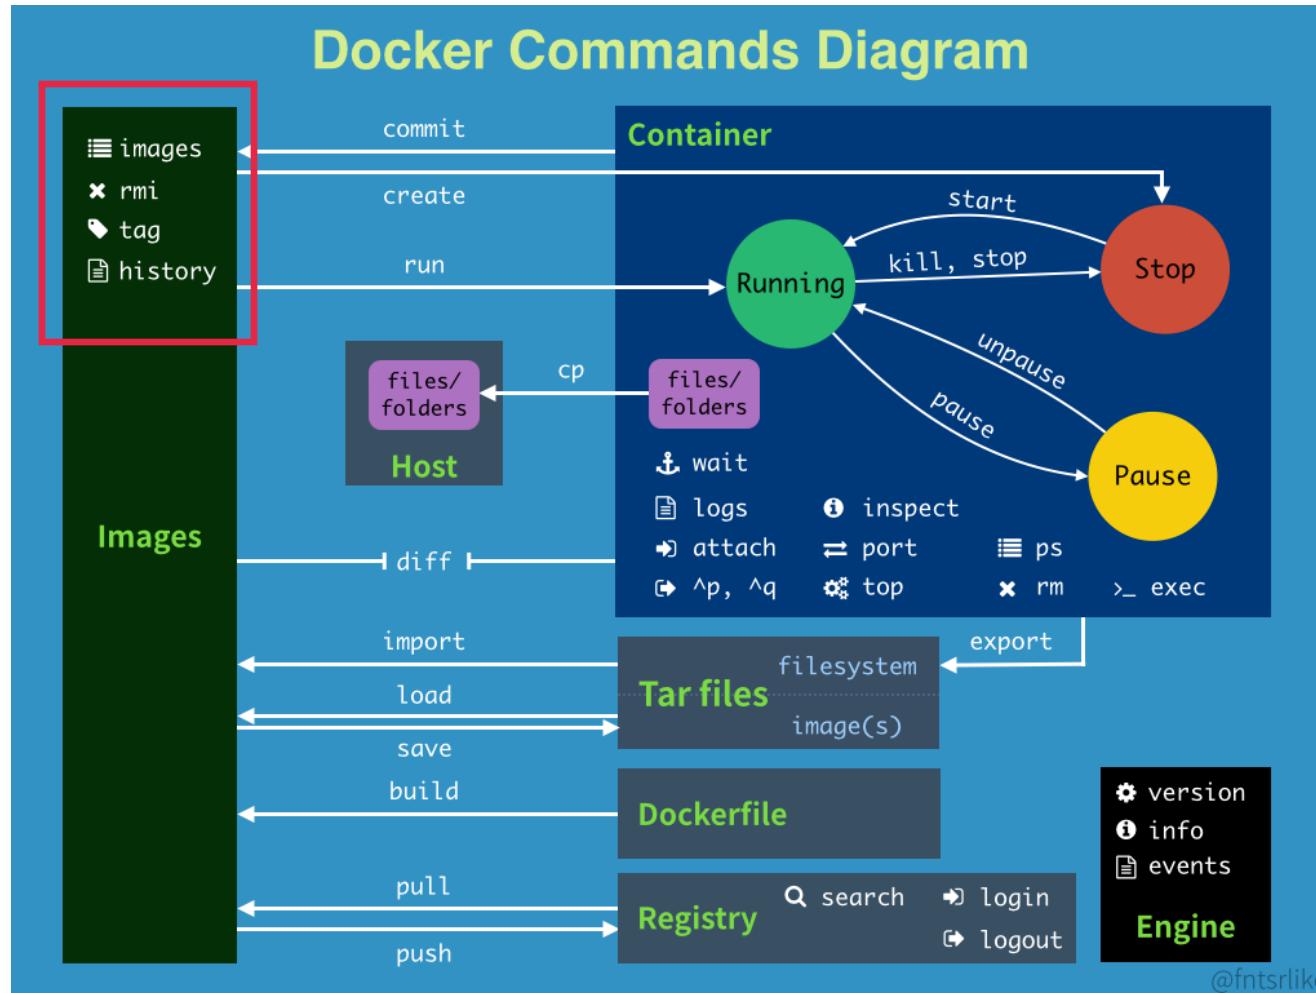
- After releases of docker images, a registry is needed to store and manage images for sharing.
- Docker images are stored in a Docker registry:
  - **Public Docker registry:** Docker Hub (available to everyone)
  - **Private Docker registry:** company's hub (only available to the employees)
- Each Docker **registry** can have multiple **repositories** and each repository can have multiple **tags**. Each tag corresponds to an image:
  - <repository\_name>:<tag>
  - Example: for ubuntu images – **ubuntu:20.04** vs **ubuntu:22.04**
  - Tags can be ignored (**ubuntu** -> **ubuntu:latest**)
- Registry's name is normally represented by a path with two components (**user/software**)
  - Example: **docker pull microsoft/cntk:2.7-cpu-python3.5**



# Outline

- Container
- What is Docker
- Basic concepts in Docker
  - Images
  - Layer architecture
  - Containers
  - Registry
- ➡ • Docker Commands
- Containerisation and Dockerfile
  - Dokcerfile instructions

# Docker Command Diagram – Image management



- List images:

`docker images [OPTIONS] [REPOSITORY[:TAG]]`

- Remove one or more image:

`docker rmi [OPTIONS] IMAGE [IMAGE...]`

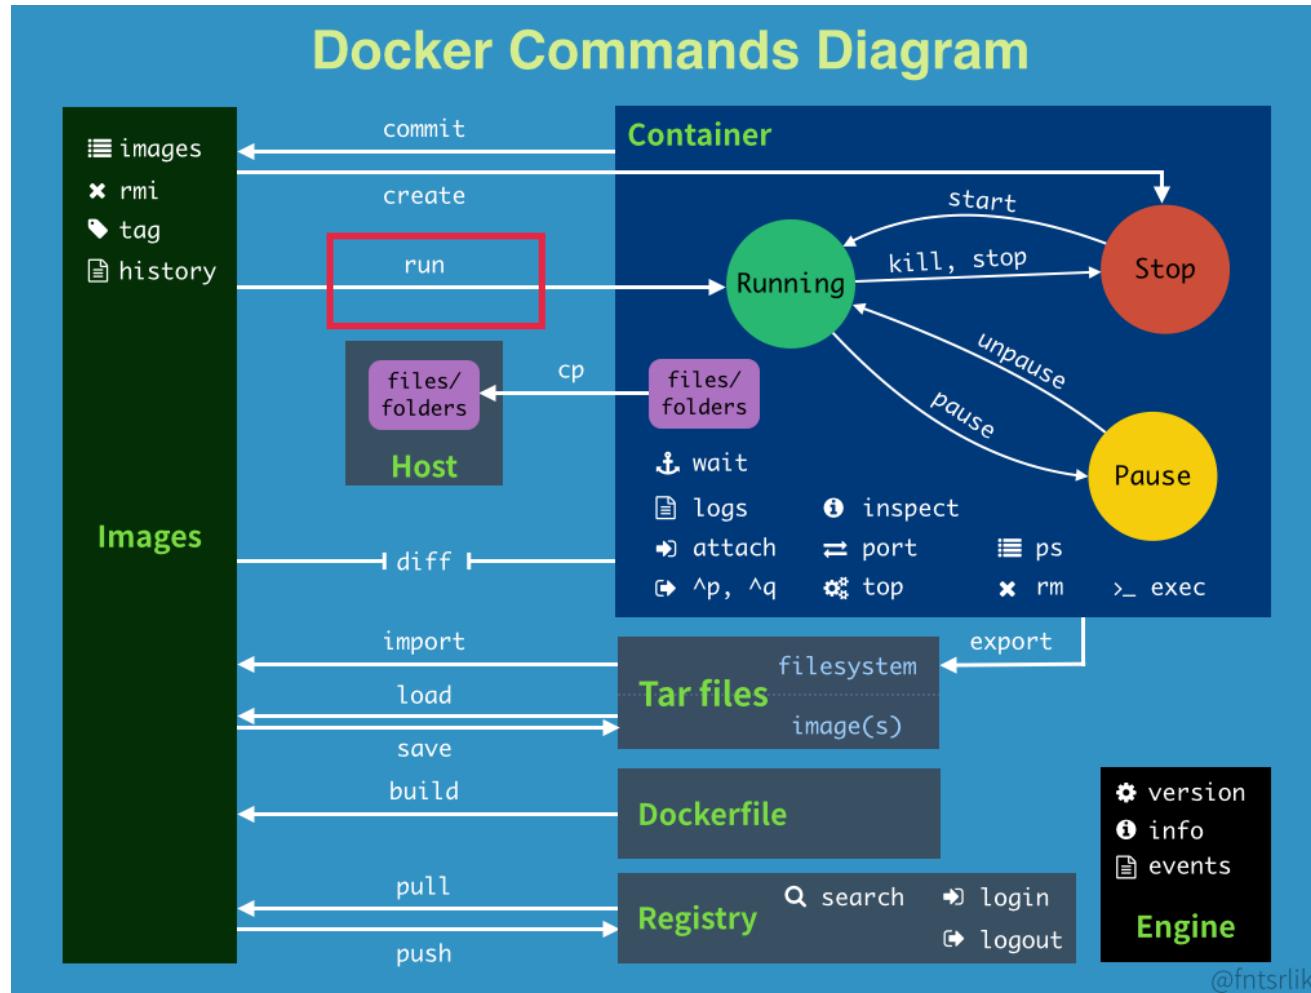
- Create a tag **TARGET\_IMAGE** that refers to **SOURCE\_IMAGE**

`docker tag SOURCE_IMAGE[:TAG] TARGET_IMAGE[:TAG]`

- Show the history of an image

`docker history [OPTIONS] IMAGE`

# Docker Command Diagram – Docker Run Images

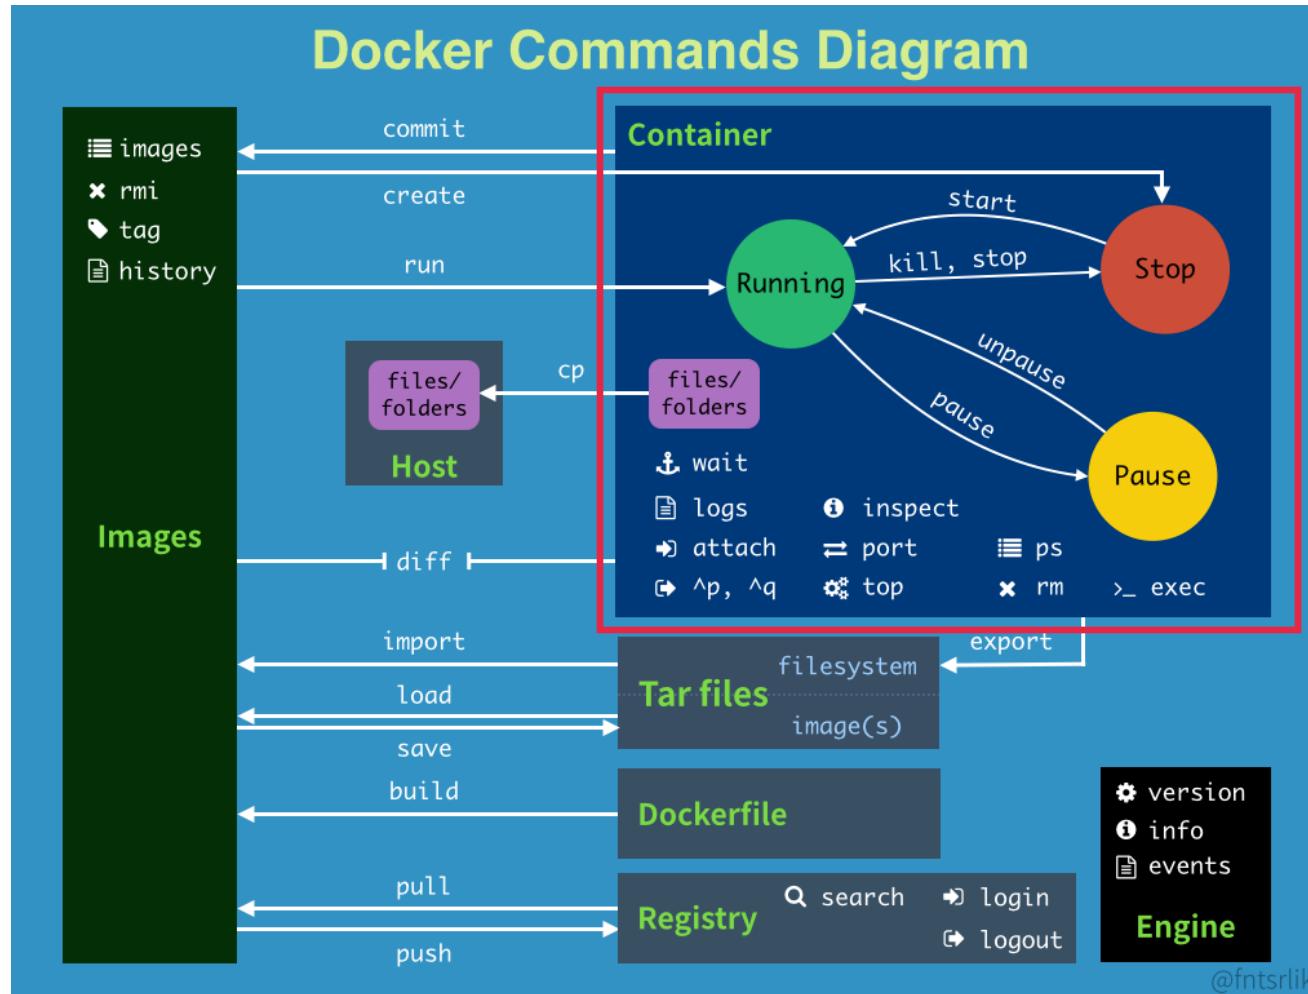


- **Run images:**

`docker run [OPTIONS] IMAGE[:TAG|@DIGEST] [COMMAND] [ARG...]`

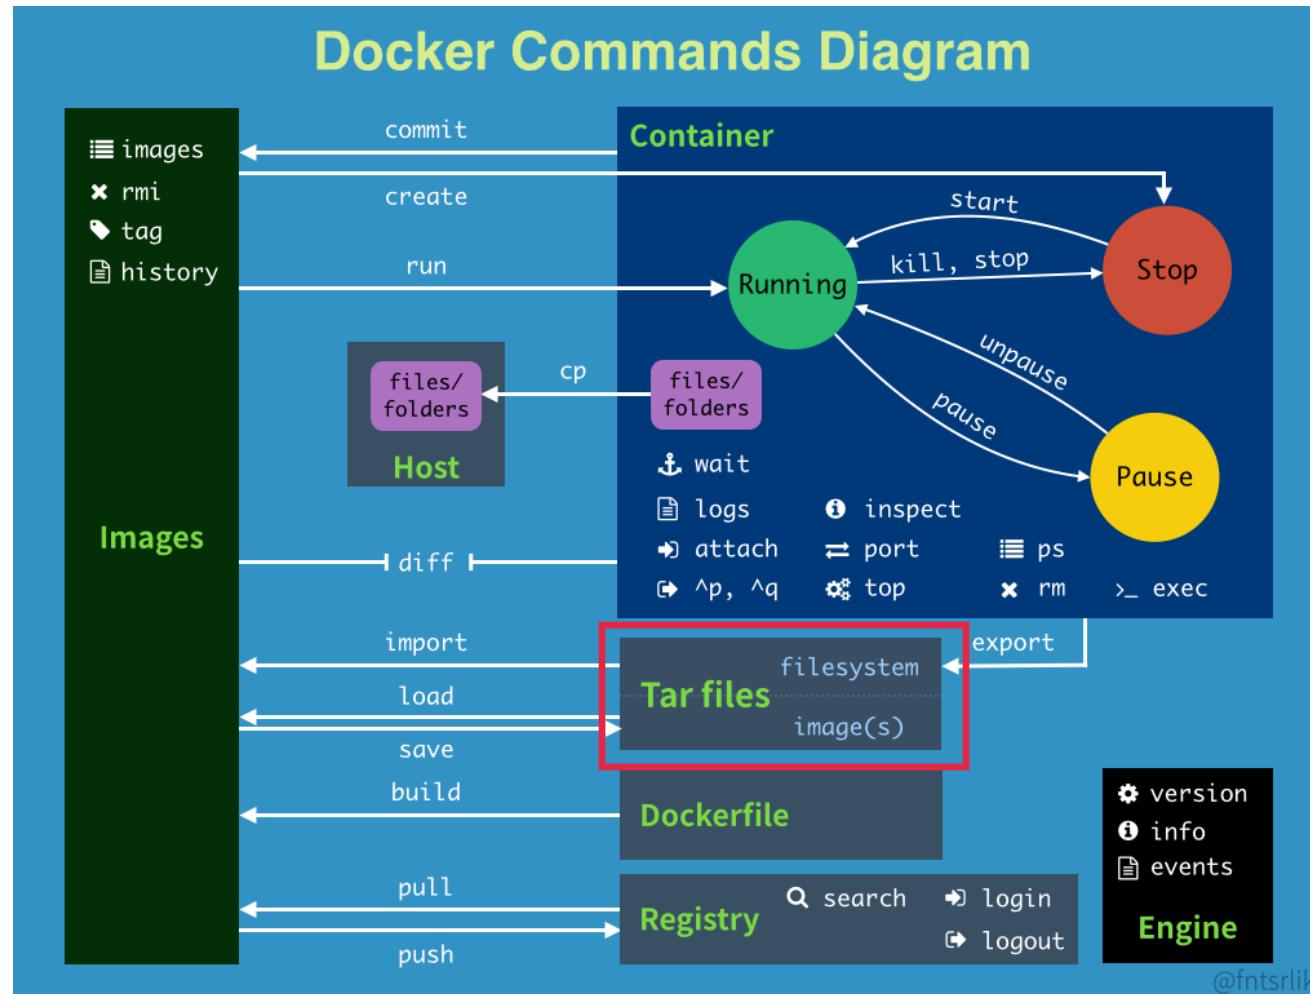
- Must specify an IMAGE to derive the container form
- Foreground (-it) vs Detached (-d)
  - E.g. `docker run -it alpine`
  - E.g. `docker run -d nginx`
- Name (--name)
  - E.g. `docker run --name my-redis -d redis`
- EXPOSE (incoming ports -P or -p)
- VOLUME (shared filesystems)
- and more...

# Docker Command Diagram – Container management



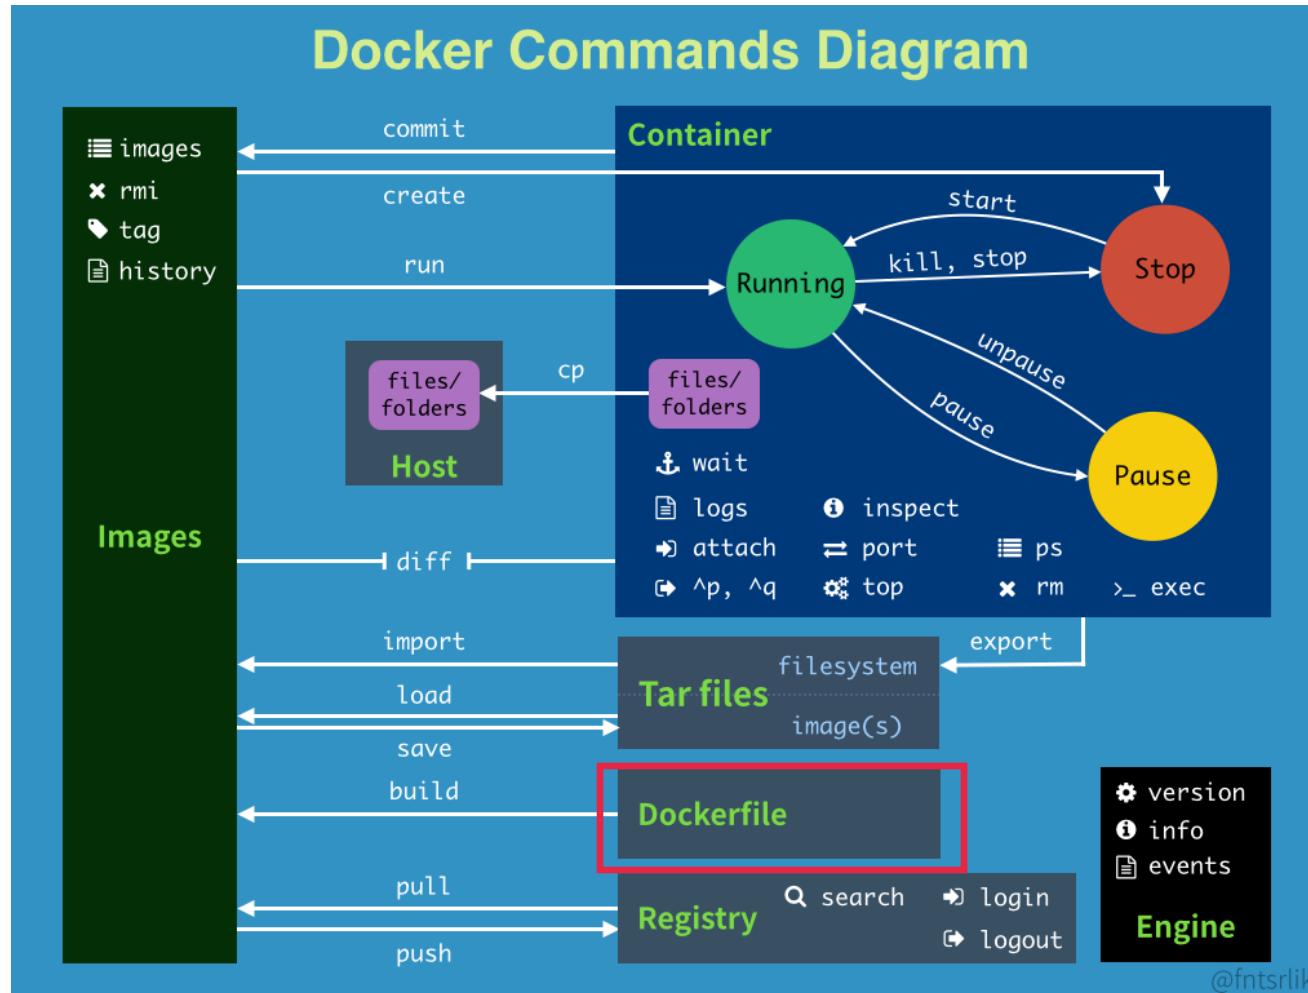
- Kill/Stop one or more running containers:**  
`docker container kill/stop [OPTIONS] CONTAINER [CONTAINER...]`  
Or  
`docker kill/stop [OPTIONS] CONTAINER [CONTAINER...]`
- Start one or more stopped containers:**  
`docker container start [OPTIONS] CONTAINER [CONTAINER...]`
- Pause/Unpause all processes within one or more containers**  
`docker container pause/unpause [OPTIONS] CONTAINER [CONTAINER...]`
- Other container management commands**  
List containers: `docker ps`  
Attach containers: `docker attach [OPTIONS] CONTAINER`  
Run a bash: `docker exec [OPTIONS] CONTAINER COMMAND [ARG...]`  
and more ...

# Docker Command Diagram – Docker sharing



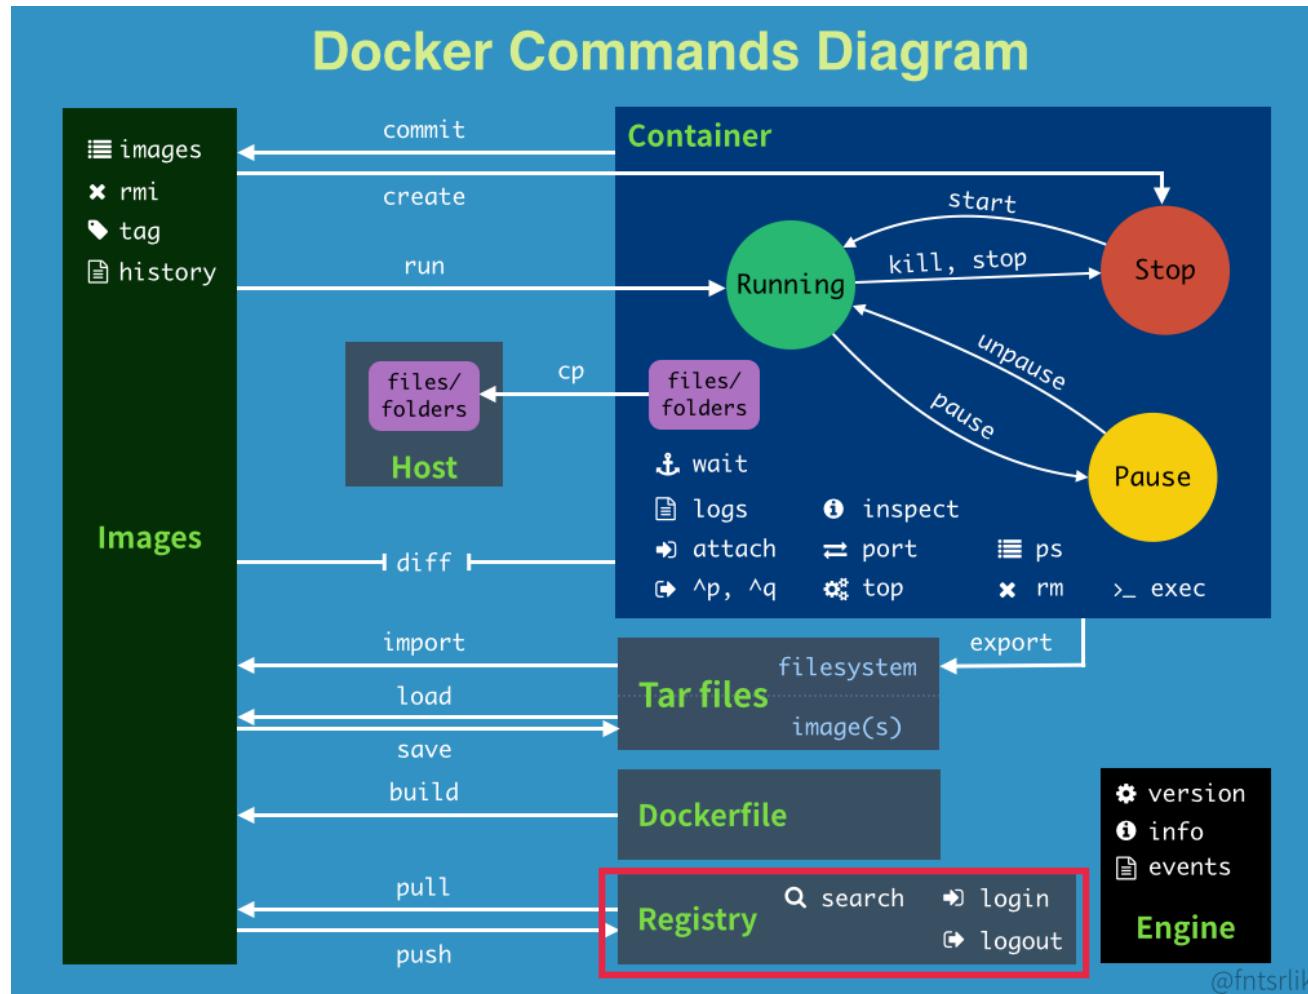
- Export a container's filesystem as a tar archive:  
`docker export [OPTIONS] CONTAINER`
- Import the contents from a tarball to create a filesystem image:  
`docker import [OPTIONS] file|URL|- [REPOSITORY[:TAG]]`
- Load an image from a tar archive or STDIN:  
`docker load [OPTIONS]`
- Save one or more images to a tar archive (streamed to STDOUT by default)  
`docker save [OPTIONS] IMAGE [IMAGE...]`

# Docker Command Diagram – Build image from dockerfile



- Build an image from a Dockerfile:  
`docker build [OPTIONS] PATH | URL | -`

# Docker Command Diagram – Interact with Registry



- Pull an image or a repository from a registry:  
`docker pull [OPTIONS] NAME[:TAG|@DIGEST]`
- Push an image or a repository to a registry:  
`docker push [OPTIONS] NAME[:TAG]`
- Search the Docker Hub for images:  
`docker search [OPTIONS] TERM`

# Outline

- Container
- What is Docker
- Basic concepts in Docker
  - Images
  - Layer architecture
  - Containers
  - Registry
- Docker Commands
- • Containerisation and Dockerfile
  - Dokcerfile instructions

# Containerizing an app

The process of taking an application and configuring it to run as a container is called “[containerizing](#)” or “[Dockerizing](#)”.

Containers are all about apps! In particular, they’re about making apps simple to [build](#), [ship](#), and [run](#).

One sharing example (with docker commands):

- Use docker commands to construct an image layer by layer
- Use docker commit to make changes when necessary
- Use docker pull/push to share with others via public/private registry

Any problems that may concern the other users?

- Size of image/container
- Contents of image/container(malicious software)

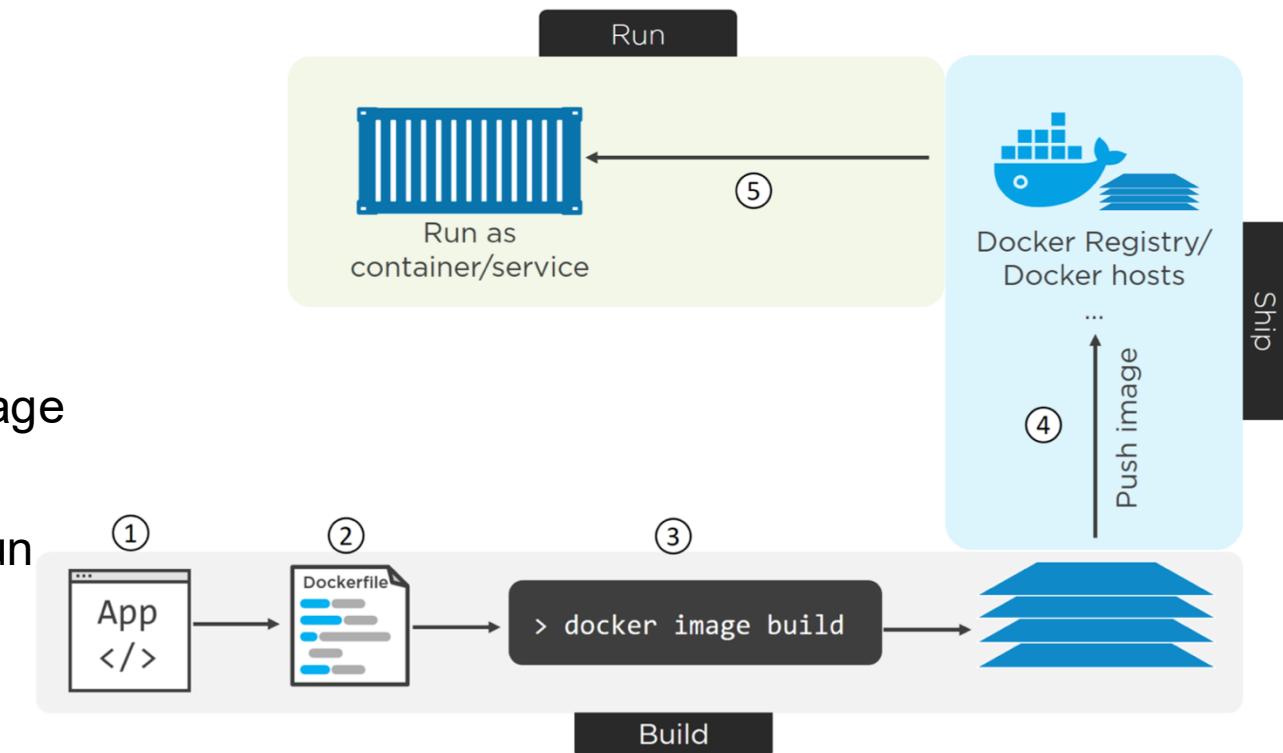
A simple solution: an “[audit](#)” file that lists all contents in this image



# Containerizing Steps

The process of containerizing an app:

1. Start with your application code.
2. Create a Dockerfile that describes your app, its dependencies, and how to run it.
3. Feed this Dockerfile into the docker image build command.
4. Docker builds your application into a Docker image and push image to Docker Registry
5. Your teammates and you can remotely/locally run the image as a container



# Dockerfile

- is a text file that defines the **environment** inside the container
- is a collection of instructions and commands (blueprint)
- clearly tells what are contained in the image.
- Docker can build images automatically by reading the instructions from a Dockerfile
  - Usage example: **docker build <dockerfile>**
- Docker image is in form of dockerfile

```
1  # This dockerfile uses the ubuntu image
2  # VERSION 2 - EDITION 1
3  # Author: docker_user
4  # Command format: Instruction [arguments / command] ..
5
6  # Base image to use, this must be set as the first line
7  FROM ubuntu
8
9  # Maintainer: docker_user <docker_user at email.com>
10 # (@docker_user)
11 MAINTAINER docker_user docker_user@email.com
12
13 # Commands to update the image
14 RUN echo "deb http://archive.ubuntu.com/ubuntu/ raring main
15 universe" >> /etc/apt/sources.list
16 RUN apt-get update && apt-get install -y nginx
17 RUN echo "\ndaemon off;" >> /etc/nginx/nginx.conf
18
19 # Commands when creating a new container
20 CMD /usr/sbin/nginx
```

# Launch a Nginx container with Docker CLI commands

```
uqteaching@instance-1:~$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
my_ubuntu	1.1	499e9fad3b9c	3 days ago	64.2MB
my_ubuntu	latest	499e9fad3b9c	3 days ago	64.2MB
wordpress	latest	ee2256095234	8 days ago	543MB
mariadb	latest	7c1c380b0dd8	12 days ago	407MB
ubuntu	latest	adafef2e596e	12 days ago	73.9MB
ubuntu	18.04	d27b9ffc5667	12 days ago	64.2MB
hello-world	latest	bf756fb1ae65	6 months ago	13.3kB

```
uqteaching@instance-1:~$ docker run -it -d -p 80:80 --name ubuntu_nginx ubuntu
```

```
1d28fafcac6b5820e7be1d294b047b40866b454b2c6b6a4048fd35ef2be29a18
```

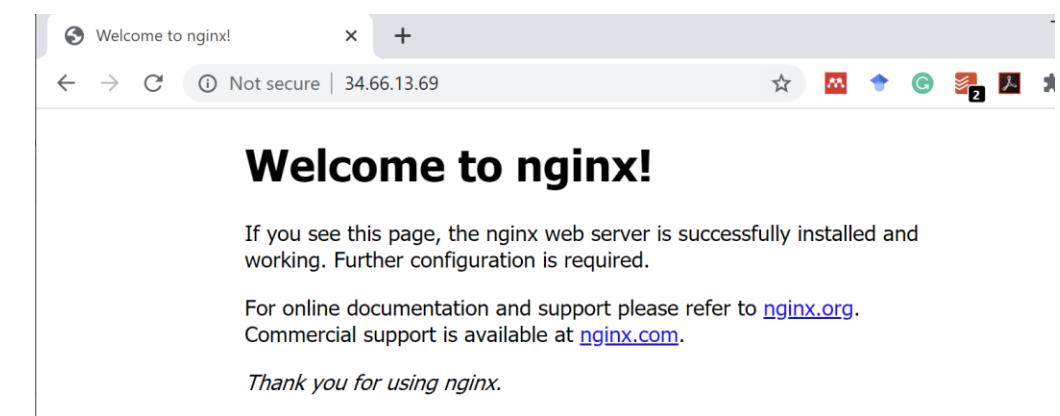
```
uqteaching@instance-1:~$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
	NAMES				
1d28fafcac6b	ubuntu	"/bin/bash"	2 seconds ago	Up 1 second	0.0.0.
0:80->80/tcp	ubuntu_nginx				

```
uqteaching@instance-1:~$ docker exec -it 1d28 bash
```

```
root@1d28fafcac6b:/# apt-get update && apt-get install -y nginx
```

```
root@1d28fafcac6b:/# service nginx start
```



# Launch a Nginx container with a Dockerfile

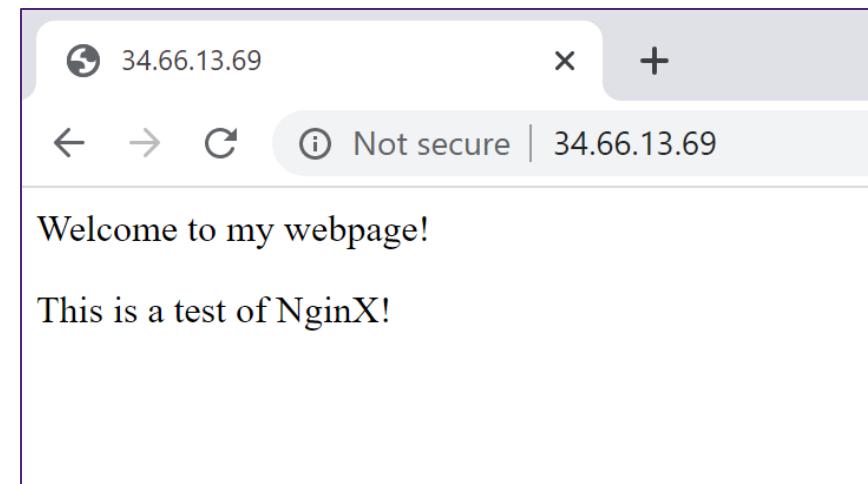
```
1 FROM ubuntu
2 LABEL maintainer="sen.wang@uq.edu.au"
3
4 # Install Nginx.
5 RUN apt-get -yqq update
6 RUN apt-get -yqq install nginx
7 COPY website /var/www/html
8 WORKDIR /var/www/html
9 # Expose ports.
10 EXPOSE 80
11 CMD ["nginx", "-g", "daemon off;"]
```

dockerfile

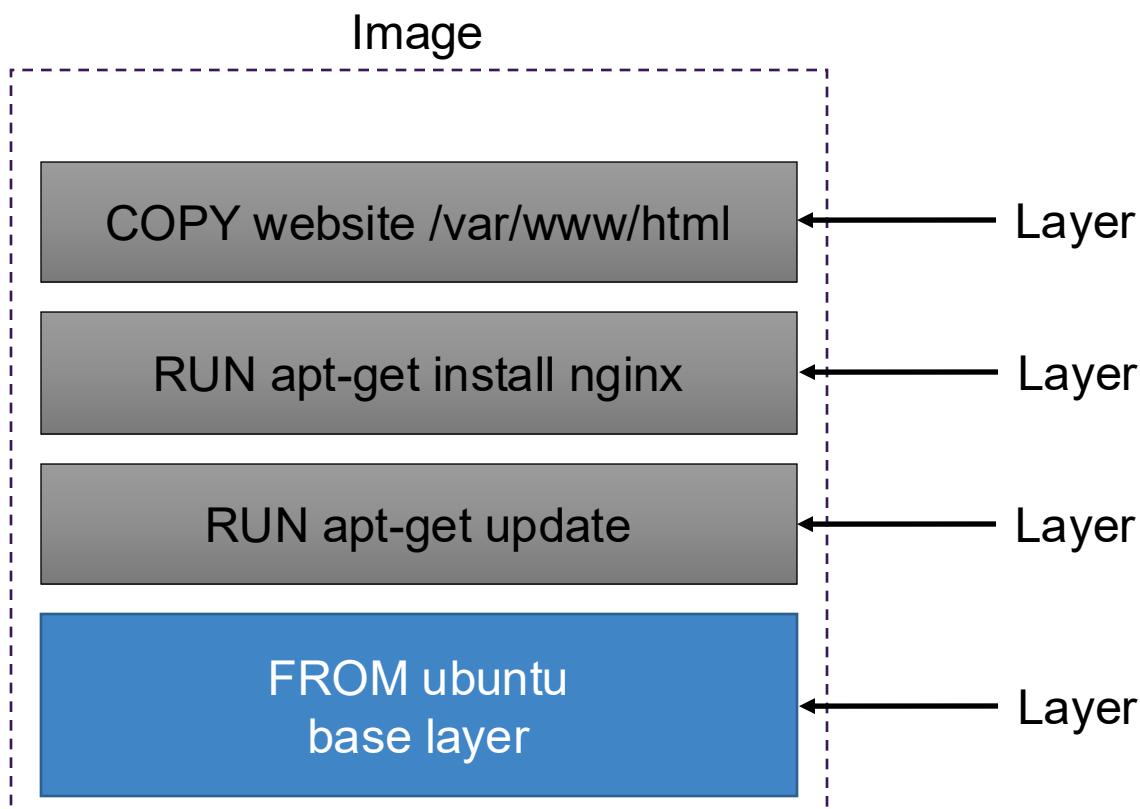
```
<h>Welcome to my webpage!</h>
<p>This is a test of NginX!</p>
~
```

Index.html

```
uqteaching@instance-1:~/lecture_demo/13/d1$ docker build -t df_ubuntu_nginx .
uqteaching@instance-1:~/lecture_demo/13/d1$ docker images
REPOSITORY          TAG      IMAGE ID      CREATED        SIZE
df_ubuntu_nginx     latest   ad1a43b49204  9 seconds ago  155MB
ubuntu_nginx         latest   23924d7573f2  About an hour ago  155MB
my_ubuntu            1.1     499e9fad3b9c  3 days ago    64.2MB
my_ubuntu            latest   499e9fad3b9c  3 days ago    64.2MB
nginx               latest   0901fa9da894  8 days ago    132MB
wordpress           latest   ee2256095234  8 days ago    543MB
mariadb              latest   7c1c380b0dd8  12 days ago   407MB
ubuntu               latest   adafef2e596e  12 days ago   73.9MB
ubuntu               18.04   d27b9ffc5667  12 days ago   64.2MB
debian               jessie   a3590c0e9ff9  5 weeks ago   129MB
hello-world          latest   bf756fb1ae65  6 months ago  13.3kB
uqteaching@instance-1:~/lecture_demo/13/d1$ docker run -it -d -p 80:80 df_ubuntu_nginx
212a95bf535cb528e4c5a8c2d63ab96d6efe95006f9516fca42234255dde1c17
```



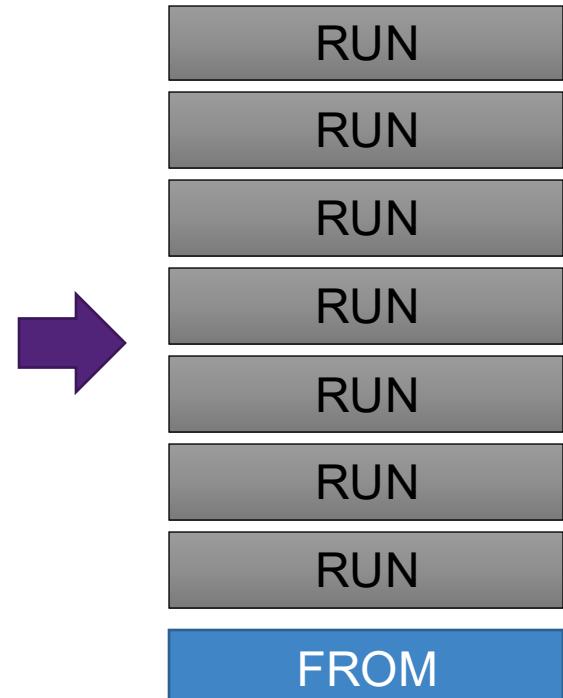
# Union FS in Image Build



```
1 FROM ubuntu
2 LABEL maintainer="sen.wang@uq.edu.au"
3
4 # Install Nginx.
5 RUN apt-get -yqq update
6 RUN apt-get -yqq install nginx
7 COPY website /var/www/html
8 WORKDIR /var/www/html
9 # Expose ports.
10 EXPOSE 80
11 CMD ["nginx", "-g", "daemon off;"]
```

# A Dockerfile Example

```
1 FROM debian
2
3 RUN apt-get update
4 RUN apt-get install -y gcc Libc6-dev make wget
5 RUN wget -O redis.tar.gz "http://download.redis.io/releases/redis-5.0.3.tar.gz"
6 RUN mkdir -p /usr/src/redis
7 RUN tar -xzf redis.tar.gz -C /usr/src/redis --strip-components=1
8 RUN make -C /usr/src/redis
9 RUN make -C /usr/src/redis install
```



- We have a base image as the first layer, “RUN” will add SEVEN different layers:
- Union FS has maximum layer number, so you cannot write as many lines as you write a shell script (e.g. > 128 lines/layers)
- There are many libs/apps that are used as temporary tools, resulting in a big and redundant image.

# A Dockerfile Example

```
1 FROM debian
2
3 RUN buildDeps='gcc libc6-dev make wget' \
4     && apt-get update \
5     && apt-get install -y $buildDeps \
6     && wget -O redis.tar.gz "http://download.redis.io/releases/redis-5.0.3.tar.gz" \
7     && mkdir -p /usr/src/redis \
8     && tar -xzf redis.tar.gz -C /usr/src/redis --strip-components=1 \
9     && make -C /usr/src/redis \
10    && make -C /usr/src/redis install \
11    && rm -rf /var/lib/apt/lists/* \
12    && rm redis.tar.gz \
13    && rm -r /usr/src/redis \
14    && apt-get purge -y --auto-remove $buildDeps
```

RUN

FROM

# Dockerfile: FROM

Syntax:

```
FROM [--platform=<platform>] <image> [AS <name>]
```

Example:

```
FROM ubuntu:22.04
```

- FROM instruction initializes a new build stage and sets the Base Image for subsequent instructions.
- As such, a valid Dockerfile must start with a FROM instruction. ARG is the only instruction that may precede FROM in the Dockerfile.
  - ARG CODE\_VERSION=latest
  - FROM base:\${CODE\_VERSION}
- The image can be any valid image
  - e.g. ubuntu/busybox on Docker Hub.
  - scratch: an explicitly empty image, especially for building images "FROM scratch"

# Dockerfile: RUN

RUN has two forms:

- **RUN <command>**
  - Shell commands (by default is /bin/sh -c on Linux)
  - E.g. **RUN apt-get update**
- **RUN ["executable", "param1", "param2"] (exec form)**
  - The exec form is parsed as a JSON array, which means that you must use double-quotes ("") around words not single-quotes ('').
  - E.g. **RUN ["/bin/bash", "-c", "echo hello"]**

To avoid redundant layers, it's recommend to concatenate all commands with “**&&**” or “**;**” and separate them with “**\**”, which is a line continuation character.

```
1 FROM debian
2
3 RUN buildDeps='gcc libc6-dev make wget' \
4   && apt-get update \
5   && apt-get install -y $buildDeps \
6   && wget -O redis.tar.gz "http://download.redis.io/releases/redis-5.0.3.tar.gz" \
7   && mkdir -p /usr/src/redis \
8   && tar -xzf redis.tar.gz -C /usr/src/redis --strip-components=1 \
9   && make -C /usr/src/redis \
10  && make -C /usr/src/redis install \
11  && rm -rf /var/lib/apt/lists/* \
12  && rm redis.tar.gz \
13  && rm -r /usr/src/redis \
14  && apt-get purge -y --auto-remove $buildDeps|
```

# Dockerfile: COPY

COPY has two forms:

- **COPY** [--chown=<user>:<group>] <src>... <dest>
  - **COPY** [--chown=<user>:<group>] [<src>, ... <dest>]
- <src> may contain **wildcards** and matching will be done using Go's filepath.Match rules.
- **COPY** website/\*.**html** /var/www/html
    - Copy all html files in website to the destination folder.
  - **COPY** website/in?**ex.html** /var/www/html
    - Copy any file named as in?ex.html to the destination folder. ? Is replaced with any single character, e.g. index.html or in\_ex.html.

Option “**--chown=<user>:<group>**” can change user and group information in destination folder (Not working on Windows-based containers).

```
1 FROM ubuntu
2 LABEL maintainer="sen.wang@uq.edu.au"
3
4 # Install Nginx.
5 RUN apt-get -yqq update
6 RUN apt-get -yqq install nginx
7 COPY website /var/www/html
8 WORKDIR /var/www/html
9 # Expose ports.
10 EXPOSE 80
11 CMD ["nginx", "-g", "daemon off;"]
```

NOTE:

“\*” matches any sequence of non-Separator characters  
‘?’ matches any single non-Separator character

# Dockerfile: ADD

ADD has two forms:

- `ADD [--chown=<user>:<group>] <src>... <dest>`
- `ADD [--chown=<user>:<group>] [<src>, ... <dest>]`

ADD is similar with COPY, but more powerful.

- `ADD <src_URL> <dest>`
  - If source is a remote file URL, ADD can download it into <dest> folder and set its permissions of 600 by default.
  - 600 (rw-,---,---): Owner(6-110)/Group(0)/Others(0)
  - If source is compressed file (tar/gzip/bzip2/gz), ADD can download it and uncompress it into destination folder.

In Best Practices for Writing Dockerfiles [1], `COPY` is preferred than `ADD` due to its semantic transparency.

[1] [https://docs.docker.com/develop/develop-images/dockerfile\\_best-practices/](https://docs.docker.com/develop/develop-images/dockerfile_best-practices/)

# Dockerfile: VOLUME

**VOLUME** ["PATH"] or **VOLUME** <PATH>

- creates a mount point with the specified name and marks it as holding externally mounted volumes from native host or other containers.
- the value can be
  - either **a JSON array**, e.g. VOLUME ["/var/log/"],
  - or **a plain string** with multiple arguments, e.g. VOLUME /var/log.
- The docker run command initializes the newly created volume with any data that exists at the specified location within the base image.

```
FROM ubuntu
RUN mkdir /myvol
RUN echo "hello world" > /myvol/greeting
VOLUME /myvol
```

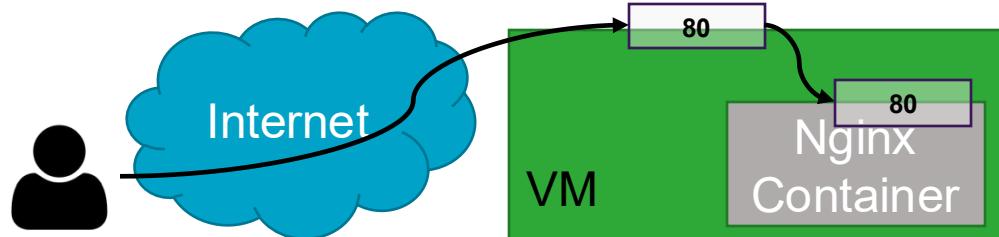
- docker run -it -d -v \$PWD/mydata:/myvol [image\_name]

# Dockerfile: EXPOSE

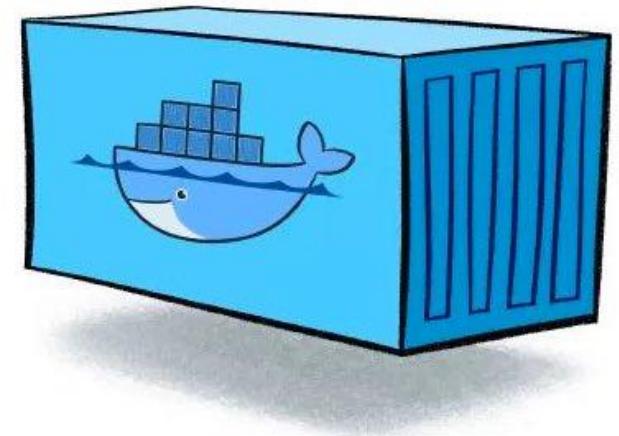
`EXPOSE <port> [<port>/<protocol>...]`

- informs Docker that the container listens on the specified network ports at runtime.
- can specify whether the port listens on TCP or UDP, and the default is TCP if the protocol is not specified.
  - `EXPOSE 80/tcp`
  - `EXPOSE 80/udp`
- In dockerfile, `EXPOSE 80` is declaring the container will be using PORT 80, but it won't actually open the port.
- Use `docker run -p <host_port>:<container_port>` to map ports.

```
uqteaching@instance-1:~/lecture_demo/13/d1$ docker run -it -d -p 80:80 df_ubuntu_nginx  
212a95bf535cb528e4c5a8c2d63ab96d6efe95006f9516fca42234255dde1c17
```



TCP: 80, 443, etc.



```
1 FROM ubuntu  
2 LABEL maintainer="sen.wang@uq.edu.au"  
3  
4 # Install Nginx.  
5 RUN apt-get -yqq update  
6 RUN apt-get -yqq install nginx  
7 COPY website /var/www/html  
8 WORKDIR /var/www/html  
9 # Expose ports.  
10 EXPOSE 80  
11 CMD ["nginx", "-g", "daemon off;"]
```

# Dockerfile: CMD

The CMD instruction has three forms:

- **CMD** ["executable","param1","param2"] (exec form, this is the preferred form)
- **CMD** ["param1","param2"] (as default parameters to ENTRYPPOINT)
- **CMD** command param1 param2 (shell form)

The main purpose of a **CMD** is to provide defaults for an executing container.

- For example, the default **CMD** of ubuntu container is "/bin/bash"
- **CMD** ["sh", "-c", "echo /etc/os-release"]

There can only be **ONE CMD** instruction in a Dockerfile.

- If you list more than one CMD then only **the last CMD** will take effect.

```
1 FROM ubuntu
2 LABEL maintainer="sen.wang@uq.edu.au"
3
4 # Install Nginx.
5 RUN apt-get -yqq update
6 RUN apt-get -yqq install nginx
7 COPY website /var/www/html
8 WORKDIR /var/www/html
9 # Expose ports.
10 EXPOSE 80
11 CMD ["nginx", "-g", "daemon off;"]
```

```
FROM ubuntu
RUN mkdir /myvol
RUN echo "hello world" > /myvol/greeting
VOLUME /myvol

CMD echo $USER
CMD [ "sh", "-c", "echo $PWD" ]
~
uqteaching@instance-1:~/lecture_demo/13/d3$ docker run -it cmd_test
/
uqteaching@instance-1:~/lecture_demo/13/d3$ █
```

# Dockerfile: CMD

```
1  FROM ubuntu
2  LABEL maintainer="sen.wang@uq.edu.au"
3
4  # Install Nginx.
5  RUN apt-get -yqq update
6  RUN apt-get -yqq install nginx
7  COPY website /var/www/html
8  WORKDIR /var/www/html
9  # Expose ports.
10 EXPOSE 80
11 CMD ["nginx", "-g", "daemon off;"]
```

- The base image is ubuntu and its default executable is “/bin/bash” (main process)
- **CMD** [“nginx”,“-g”,“daemon off;”] set the default executable of this container is nginx and set it running in the frontend
- Why not to use **“CMD service nginx start”?**
  - Improper understanding the essence of container (it’s a process not a VM)
  - Container does not exist after finishing its main process.

# Dockerfile: ENTRYPOINT

An ENTRYPOINT allows you to configure a container that will run as an executable.

ENTRYPOINT has two forms:

- **ENTRYPOINT** ["executable", "param1", "param2"] (exec form - preferred)
- **ENTRYPOINT** command param1 param2 (The shell form)

```
FROM ubuntu
ENTRYPOINT ["echo", "Hello"]
```

```
uqteaching@instance-1:~/lecture_demo/13/d4$ docker run -it entry_test
Hello
uqteaching@instance-1:~/lecture_demo/13/d4$ docker run -it entry_test Docker
Hello Docker
uqteaching@instance-1:~/lecture_demo/13/d4$ docker run -it entry_test -setting Docker
Hello -setting Docker
```

When using both ENTRYPOINT and CMD in one dockerfile,

- CMD will become the additional parameter -> **ENTRYPOINT “CMD”**
- **CMD can be overridden** by command line arguments provided when docker container runs.

```
FROM ubuntu
ENTRYPOINT ["echo", "Hello"]
CMD ["Docker World!"]
```

```
uqteaching@instance-1:~/lecture_demo/13/d4$ docker run -it entry_cmd_test
Hello Docker World!
uqteaching@instance-1:~/lecture_demo/13/d4$ docker run -it entry_cmd_test Cloud Computing
Hello Cloud Computing
```

# Dockerfile: ENV

The ENV instruction sets the environment variable <key> to the value <value>.

This value will be in the environment for all subsequent instructions in the build stage and can be replaced inline in many as well.

The ENV instruction has two forms:

- **ENV <key> <value>**
  - **ENV myName Sen Wang**
  - **ENV mySchool EECS**
  - **ENV myOffice GPS649**
- **ENV <key>=<value> ...**
  - **ENV myName="Sen Wang" mySchool="EECS"**

```
Step 2/12 : ENV myName Sen Wang
--> Running in 3ef532f853a0
Removing intermediate container 3ef532f853a0
--> 6e7542d71332
Step 3/12 : ENV mySchool Information Technology and Electrical Engineering
--> Running in 008fa6299f61
Removing intermediate container 008fa6299f61
--> ddd877615546
Step 4/12 : ENV myOffice GPS649
--> Running in 83e2d97e3c29
Removing intermediate container 83e2d97e3c29
--> 02387c5603ab
Step 5/12 : RUN mkdir /myvol
--> Running in d14ec795226e
Removing intermediate container d14ec795226e
--> 9e2da88d4447
Step 6/12 : RUN echo "hello world" > /myvol/greeting
--> Running in a72f35e44235
Removing intermediate container a72f35e44235
--> b822378936c8
Step 7/12 : VOLUME /myvol
--> Running in f41402964340
Removing intermediate container f41402964340
--> e1a9e4ac9ff8
Step 8/12 : RUN echo $myName
--> Running in 313f775754f4
Sen Wang
Removing intermediate container 313f775754f4
--> 91619327228b
Step 9/12 : RUN echo $mySchool
--> Running in a8d2d8edbc85
Information Technology and Electrical Engineering
Removing intermediate container a8d2d8edbc85
--> 508ae1225bb8
Step 10/12 : RUN echo $myOffice
--> Running in 10721e938553
GPS649
```

# Dockerfile: ARG

The ARG instruction defines a variable that users can pass at **build-time** to the builder with the docker build command using the **--build-arg <varname>=<value>** flag.

**ARG <name>[=<default value>]**

```
1  FROM ubuntu
2
3  ARG user1=sen
4  ARG ver=1
```

Run “**docker build --build-arg user1=wang**” to build the dockerfile

- **user1** will be overridden as “wang” is passed to the first argument
- **ver** will be 1 as default.

Note: Do not use ARG to save your credentials (e.g. github keys, password) because **build-time** variable values are visible to any user of the image with the **docker history** command.

# Dockerfile: ENV vs ARG

Both ENV and ARG can be used to define variables.

- Their life cycles are different.
  - ARG defines variables during **build-time**;
  - ENV defines variables during **running container**.
- ARG values can be changed when building the image
  - **--build-arg** can be used to over the ARG values
- ENV cannot be directly changed when building the image

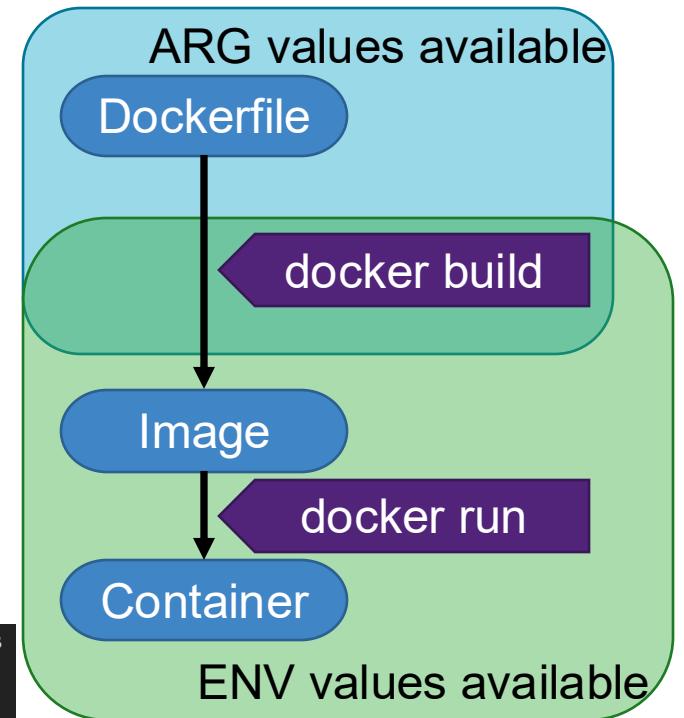
```
FROM ubuntu

ARG VAR_A=5
ENV VAR_B $VAR_A

RUN echo $VAR_A
RUN echo $VAR_B
```

docker build -t env\_arg\_test --build-arg VAR\_A=10 .■

```
Sending build context to Docker daemon 2.048kB
Step 1/5 : FROM ubuntu
--> adafef2e596e
Step 2/5 : ARG VAR_A=5
--> Running in c46b303f6564
Removing intermediate container c46b303f6564
--> 42525016b12b
Step 3/5 : ENV VAR_B $VAR_A
--> Running in 59b797e3a0bd
Removing intermediate container 59b797e3a0bd
--> 1a3c85c8cf63
Step 4/5 : RUN echo $VAR_A
--> Running in 56589fa2dd52
10
Removing intermediate container 56589fa2dd52
--> ae645e07b22d
Step 5/5 : RUN echo $VAR_B
--> Running in bd9ea9da1811
10
```



<https://vsupalov.com/docker-arg-vs-env/>

# Dockerfile: WORKDIR

The WORKDIR instruction sets the working directory (or current working directory)

`WORKDIR /path/to/workdir`

- If the WORKDIR doesn't exist, it will be created even if it's not used in any subsequent Dockerfile instruction.
- The WORKDIR instruction can be used multiple times in a Dockerfile.
- If a relative path is provided, it will be relative to the path of the previous WORKDIR instruction.

```
WORKDIR /a  
WORKDIR b  
WORKDIR c  
RUN pwd
```

`/a/b/c`.

Why do we need WORKDIR?

```
RUN cd /app  
RUN echo "hello" > world.txt
```

```
RUN echo "hello" > world.txt  
RUN cd /app
```

# Dockerfile: USER

The USER instruction sets the user name (or UID) and optionally the user group (or GID) to use when running the image and for any RUN, CMD and ENTRYPOINT instructions that follow it in the Dockerfile.

- **USER <user>[:<group>]** or **USER <UID>[:<GID>]**
- Similarly, USER just change the current user for the next layer.
- The <user> must be pre-defined, otherwise you cannot change.

```
RUN groupadd -r redis && useradd -r -g redis redis
USER redis
RUN [ "redis-server" ]
```

# Building Solutions

Before Docker 17.05, there are two ways to build docker images:

- Pack everything into **ONE** Dockerfile:
  - Dockerfile could be lengthy and not easy to maintain
  - Too many layers result in fatty images and long time of deployment.
  - Code leaks risks
- Distribute to **multiple** Dockerfiles:
  - One dockerfile is to build the image
  - One dockerfile is to copy all the required libs/programs
  - One script to run dockerfiles
  - Can significantly reduce image size, but hard to maintain two dockerfiles synchronised.

## Dockerfile.build

```
1  FROM golang:1.9-alpine
2  RUN apk --no-cache add git
3  WORKDIR /go/src/github.com/go/helloworld COPY app.go .
4  RUN go get -d -v github.com/go-sql-driver/mysql \
5    && CGO_ENABLED=0 GOOS=linux go build -a -installsuffix cgo -o app .
```

## Dockerfile.copy

```
1  FROM alpine:latest
2  RUN apk --no-cache add ca-certificates WORKDIR /root/
3  COPY app .
4  CMD ["./app"]
```

## Build.sh

```
1 #!/bin/sh
2  echo Building go/helloworld:build
3  docker build -t go/helloworld:build . -f Dockerfile.build
4  docker create --name extract go/helloworld:build
5  docker cp extract:/go/src/github.com/go/helloworld/app ./app docker rm -f extract
6  echo Building go/helloworld:2
7  docker build --no-cache -t go/helloworld:2 . -f Dockerfile.copy rm ./app
```

# Multistage Builds

- Use multiple FROM statements in one Dockerfile
- Each FROM can use a different base image
- Each FROM indicates a new stage of the build
- Selectively copy artifacts from one stage to another
  - COPY --from=0 /go/src/xxx .

```
$ docker build -t go/helloworld:3 .
```

```
$ docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
go/helloworld	3	d6911ed9c846	7 seconds ago	6.47MB
go/helloworld	2	f7cf3465432c	22 seconds ago	6.47MB
go/helloworld	1	f55d3e16affc	2 minutes ago	295MB

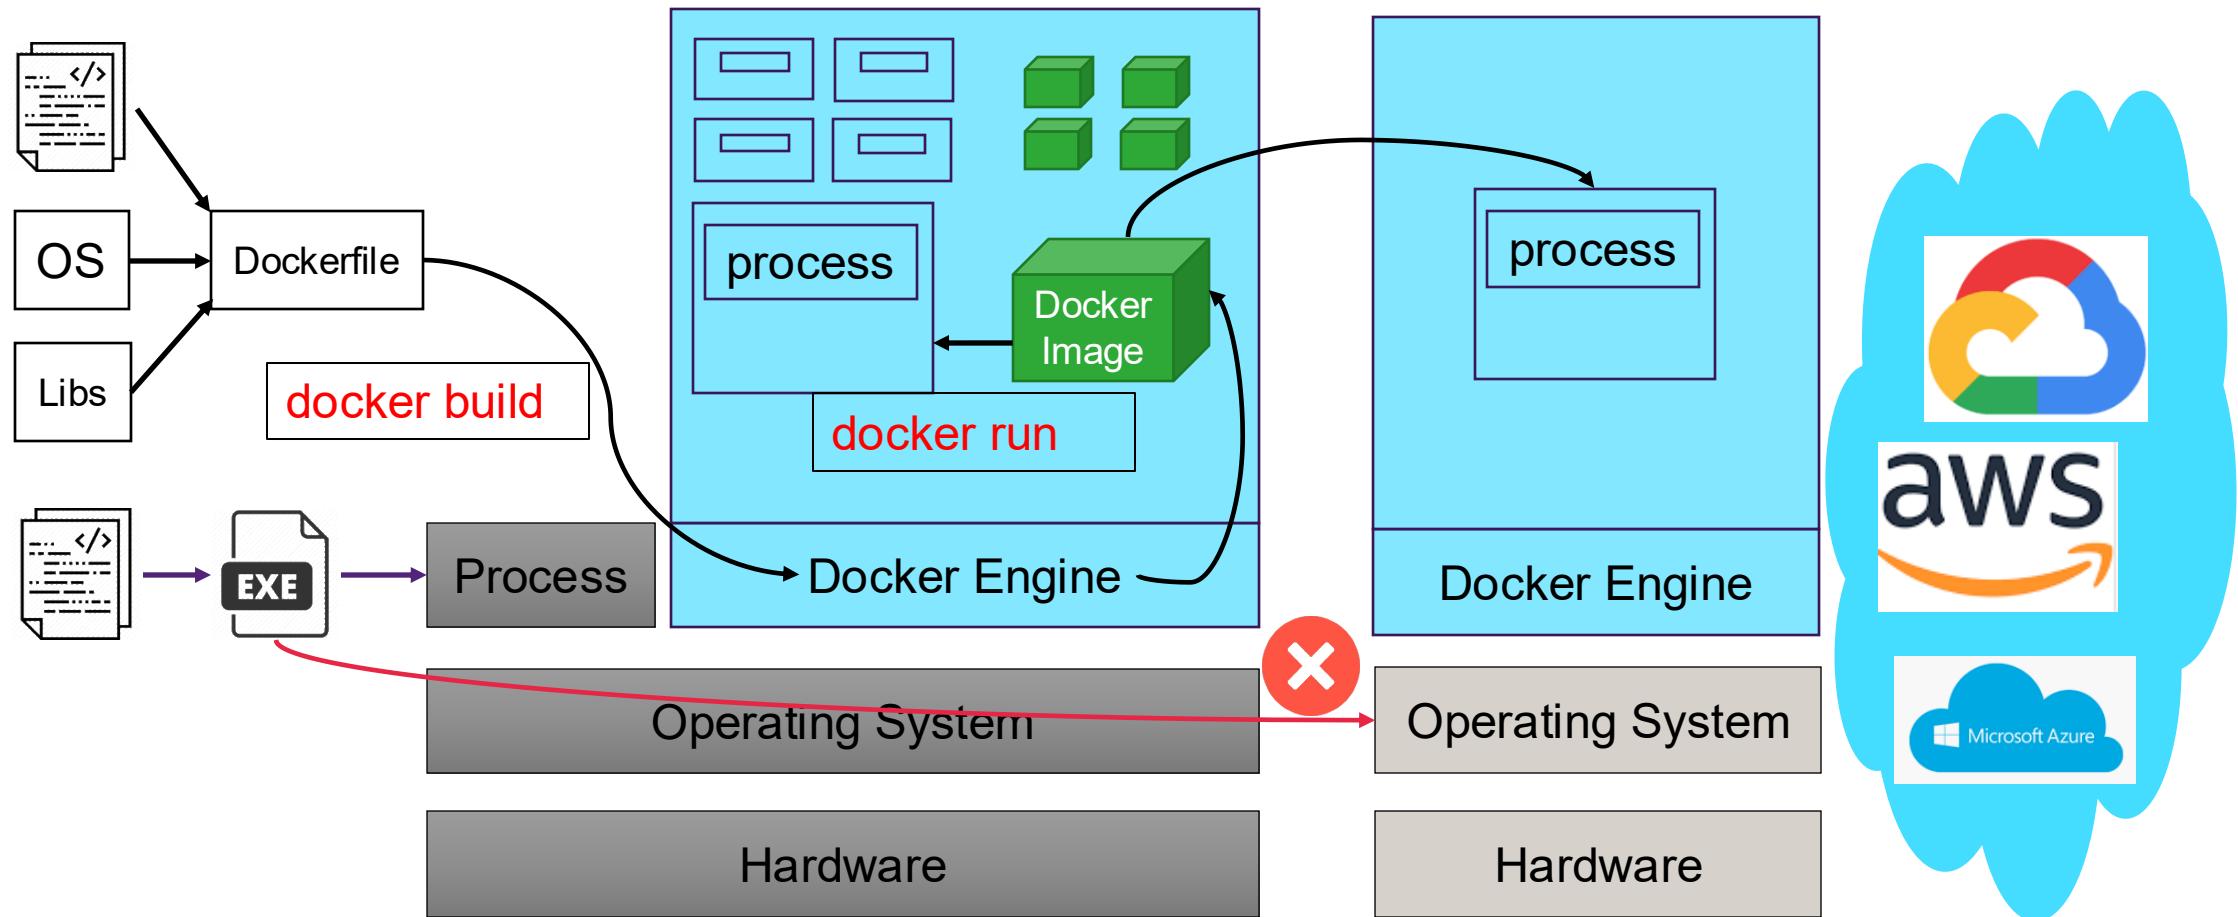
```
1  FROM golang:1.9-alpine
2  RUN apk --no-cache add git
3  WORKDIR /go/src/github.com/go/helloworld/
4  RUN go get -d -v github.com/go-sql-driver/mysql
5  COPY app.go .
6  RUN CGO_ENABLED=0 GOOS=linux go build -a
   -installsuffix cgo -o app .
7
8  FROM alpine:latest
9  RUN apk --no-cache add ca-certificates
10 WORKDIR /root/
11 COPY --from=0 /go/src/github.com/go/helloworld/app .
12 CMD ["../app"]
```

<https://docs.docker.com/develop/develop-images/multistage-build/>

# Best practices for writing Dockerfiles

- Exclude with `.dockerignore`
  - like `.gitignore`
- Use multi-stage builds
- Don't install unnecessary packages
- Decouple applications
- Minimize the number of layers
- Sort multi-line arguments
- ...

# Revisit: How does Dockerfile work in the Cloud?



# Summary

- Container
- What is Docker
- Basic concepts in Docker
  - Images
  - Layer architecture
  - Containers
  - Registry
- Docker Commands
- Containerisation and Dockerfile
  - Dokckerfile instructions

# Tutorial & Practical for Week 3

## Tutorial 2 (Week 3)

1. What is a Virtual Private Cloud (VPC) network?
2. Please discuss the differences between VPCs in AWS and GCP?
3. What is Load Balancing? Please discuss the learned LB algorithms and compare them.
4. Please discuss the Layer 4 and Layer 7 load balancers.

## Practical 2 (Week 3)

1. Based on the existing VM created on GCP, install and set up Docker environment
2. Docker commands for beginners
3. Set up a personal blog website with WordPress and MariaDB dockers on GCP.
4. Write a dockerfile to build your own Nginx image

Next (Week 4) Topic:

Docker II: Docker Compose and Docker Swarm