

Cloud Computing (INFS3208)

Lecture 5: Kubernetes

Lecturer: Dr Heming Du

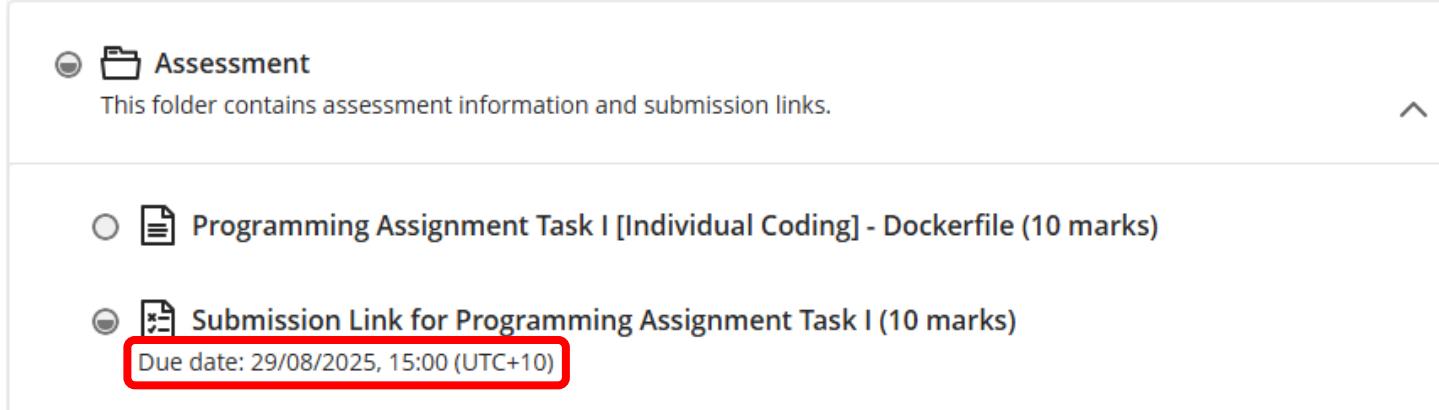
School of Information Technology and Electrical Engineering

Faculty of Engineering, Architecture and Information Technology

The University of Queensland

The remainder of Programming Task I Submission

- Due:
 - **3 p.m. on Friday, 29/08/2025 AEST.**
- Extension application
 - **Must have a medical certificate**
 - **by 5 p.m. Wednesday 27 Aug 2025**
 - **Late submission policy.**
- Ed Discussion (more tutors).
- The submission link is available under “Assessment” on the course website



Assessment

This folder contains assessment information and submission links.

Programming Assignment Task I [Individual Coding] - Dockerfile (10 marks)

Submission Link for Programming Assignment Task I (10 marks)
Due date: 29/08/2025, 15:00 (UTC+10)

Re-cap

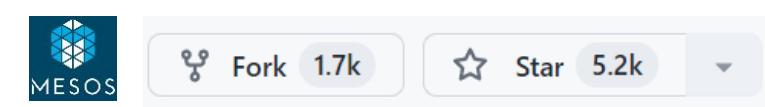
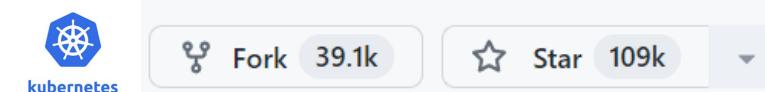
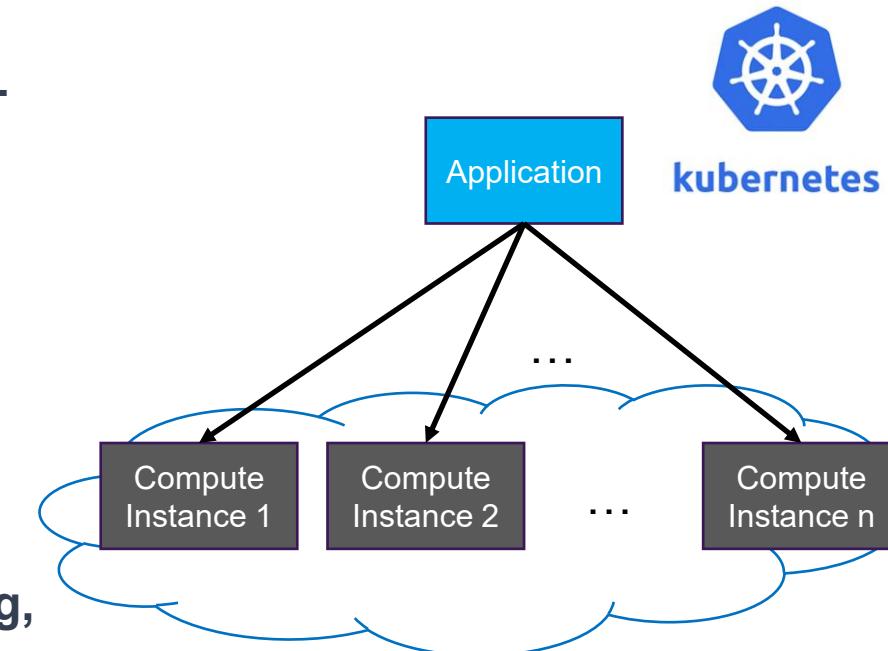
- Microservices
- Docker Compose
- Docker Swarm
 - Docker Machine
 - Create a Swarm
 - Deploy Services to a Swarm
 - Deploy a Stack to a Swarm

Outline

- • Introduction to K8s
- K8s Architecture and its Components
 - Control Plane Node and Work Node
- Declarative Model in K8s
- Reconciliation Loop
- Minikube & Kubectl
- Demo of GKE - Cluster creation
- K8s Objects:
 - Pod & ReplicaSet & Deployment
 - & Service & Ingress & StatefulSet
 - & ConfigMap & Secrets & Volumes
- A K8s deployment example - Mysql & phpMyAdmin

What is Kubernetes?

- Kubernetes (koo-ber-net-eez, commonly stylised as K8s) is an **open-source container-orchestration system** for managing (deploying and scaling) containerised applications across multiple hosts
- K8s provides automating deployment, scaling and management of containerised applications.
- It was originally designed by Google and is now maintained by the Cloud Native Computing Foundation.
- It aims to provide a "**platform for automating deployment, scaling, and operations of application containers across clusters of hosts.**"
- It works with a range of container tools, including **Docker**.



<https://en.wikipedia.org/wiki/Kubernetes>
<https://kubernetes.io/case-studies/>

Why Kubernetes?



kubernetes

- **Scaling & High Availability:**
 - automatically scales applications based on demand without manual intervention.
 - ensures that there are always a specified number of pods running (across multiple nodes), which leads to high availability.
- **Self-healing:**
 - automatically restarts containers that fail, replaces them, and even reschedules containers when nodes die.
 - If a container doesn't respond to user-defined health checks, it's automatically terminated and replaced.
- **Automated Rollouts and Rollbacks:** ensures that updates to applications roll out progressively while monitoring application health to ensure it doesn't kill all instances simultaneously. If something goes wrong, Kubernetes can roll back the change.
- **Integrated Developer Tools:**
 - There's a broad selection of tools that have been built around the Kubernetes ecosystem, providing enhanced capabilities such as CI/CD, monitoring, logging, and more.
 - **More Features:** Automated Service Discovery, Load Balancing, and Storage Orchestration, Secret and Configuration Management, Multi-cloud and Hybrid-cloud Capability (Platform Agnostic), Large Community and Strong Ecosystem, etc.

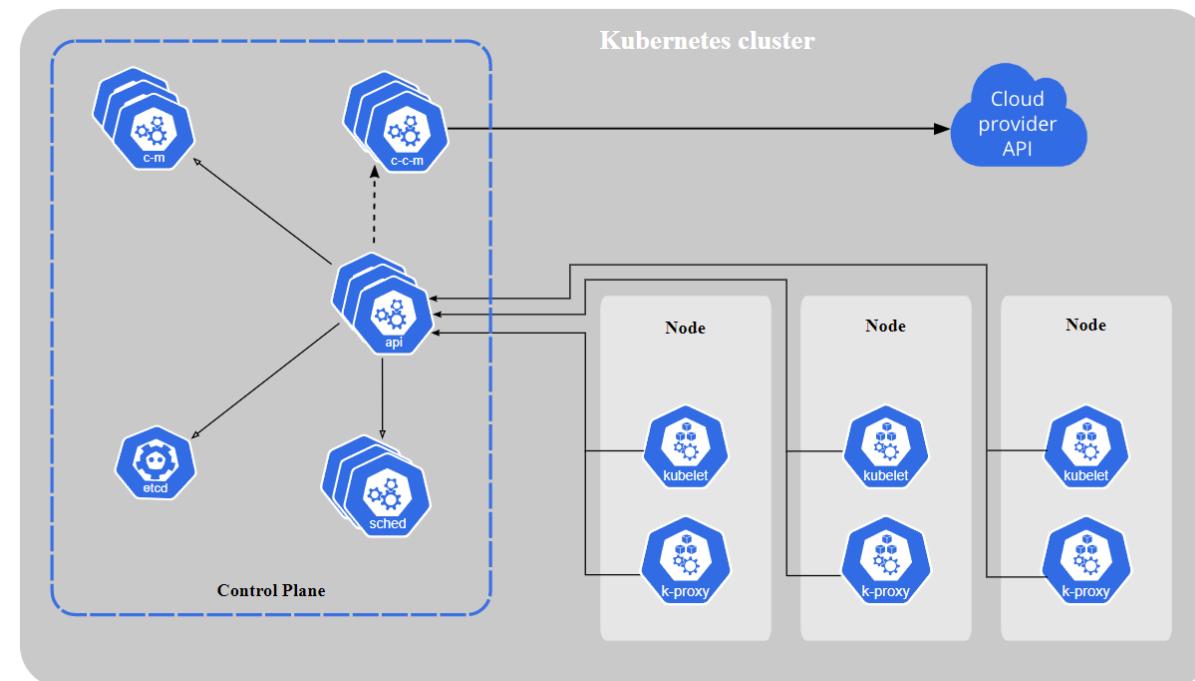
Outline

- Introduction to K8s
- • K8s Architecture and its Components
 - Control Plane Node and Work Node
- Declarative Models in K8s
- Reconciliation Loop
- Minikube & Kubectl
- Demo of GKE - Cluster creation
- K8s Objects:
 - Pod & ReplicaSet & Deployment
 - & Service & Ingress & StatefulSet
 - & ConfigMap & Secrets & Volumes
- A K8s deployment example - Mysql & phpMyAdmin

Kubernetes Architecture

A cluster of machines:

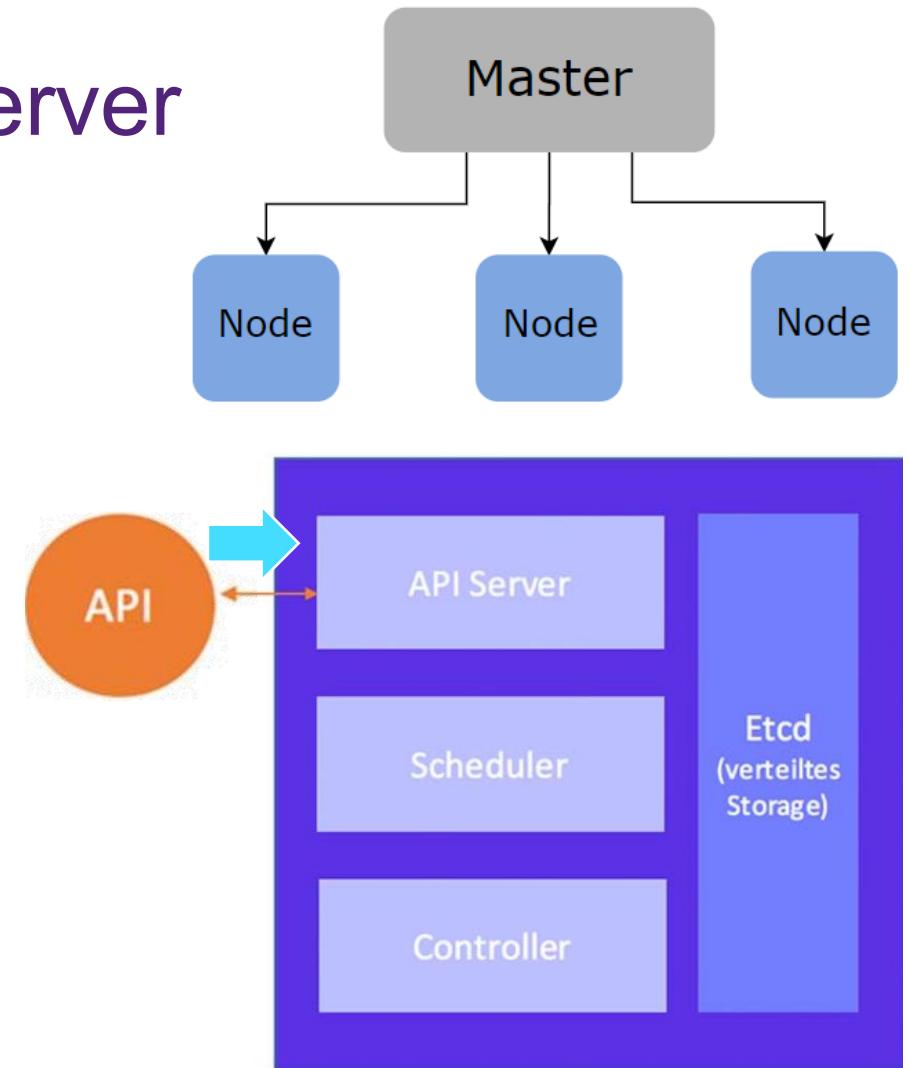
- **Control Plane (Master Node)**
 - Manages the (worker) nodes
 - Schedule (pod),
 - detect or respond to cluster events
 - Provides User Interface
 - **FOUR** processes need to run in Control Plane:
 - API Server, Scheduler, Controller Manager, Etcd
 - Multi-master high availability (**HA**) is a must-have
- **Worker Nodes (Slave Nodes)**
 - Conduct actual works with multiple pods
 - **THREE** processes need to run on every node
 - Container runtime, Kubelet, Kube proxy



Control Plane Components – API Server

API Server

- Central component in Control Plane
 - an **interactive interface** (cluster gateway) between the **cluster** and the **client** (e.g. cmd/UI)
- Updates and queries between the cluster and the client (e.g. kubectl)
 - **Kubectl**: a command line tool for communicating with a Kubernetes cluster's control plane, using the Kubernetes APIs.
- Provides k8s APIs and authentication (i.e. gatekeeper)
- Implementation: **kube-apiserver** (horizontally scalable)



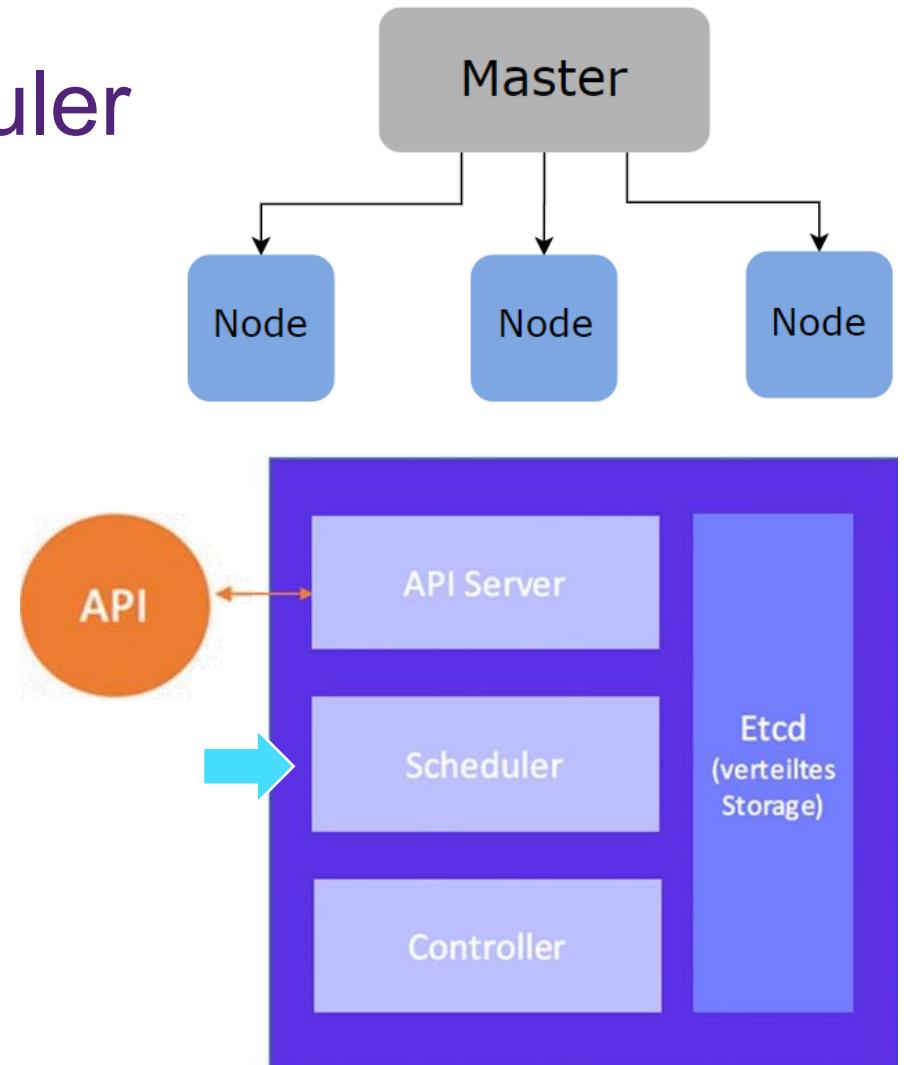
Source:

<https://knowitinfo.com/what-is-kubernetes-control-plane/>
<https://kubernetes.io/docs/concepts/overview/components/>

Control Plane Components - Scheduler

Scheduler

- Watch for the unassigned ‘task’
- Assign it to a node that has available resources (CPU & memory) matching the requirement
 - Kubelet (on Node) actually executes the scheduled pods.
 - Perform predicate checks and rank nodes
- Many factors can be considered when scheduling a task:
 1. Individual and collective resource requirements
 2. Hardware/software/policy constraints
 3. Data locality
 4. Etc.
- Implementation: kube-scheduler



Source:

<https://knowitinfo.com/what-is-kubernetes-control-plane/>
<https://kubernetes.io/docs/concepts/overview/components/>

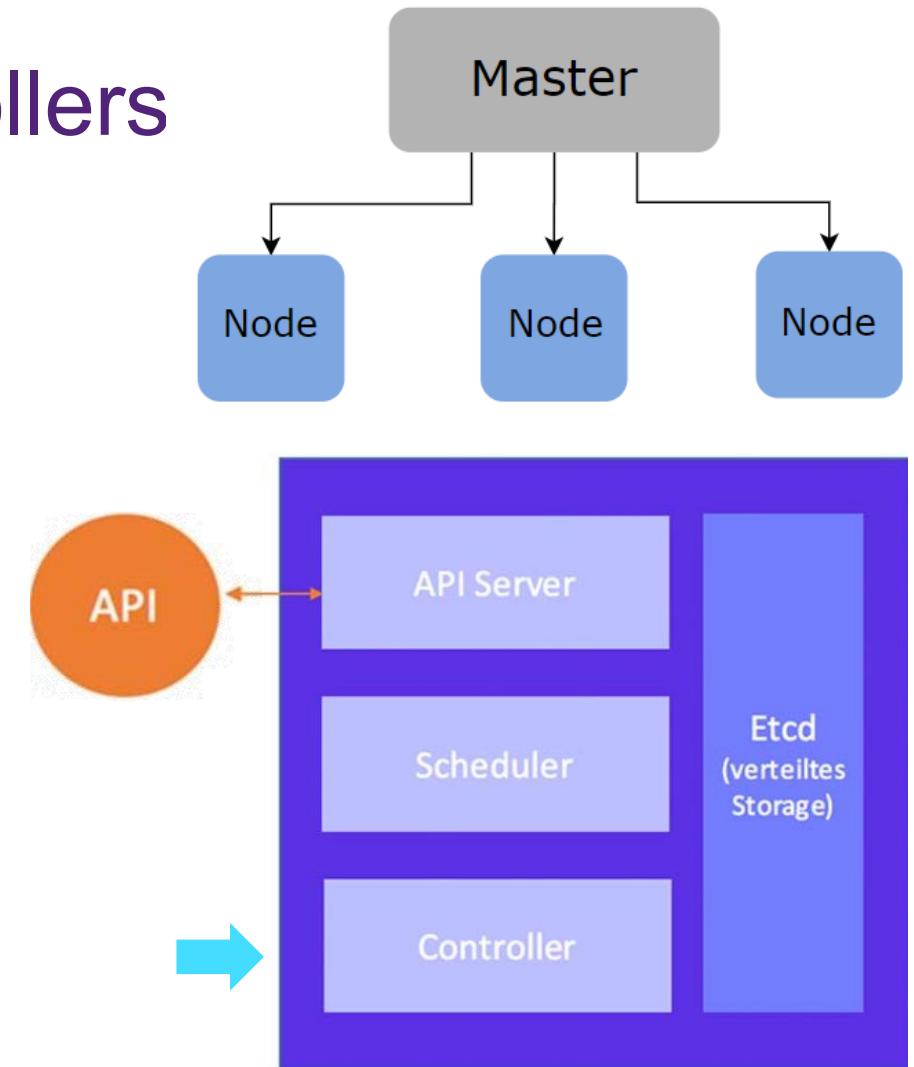
Control Plane Components - Controllers

Controllers

- **Node controller** - Responsible for noticing and responding when nodes go down.
- **Job controller** - Watches for Job objects that represent one-off tasks, then creates Pods to run those tasks to completion.
- **Endpoint controller** - Populates the Endpoints object (e.g. pod-ip:service-port).
- **Namespace controller & Service Account controller & Token controller**, etc.

Controller Manager – controller of controllers

- **Detects** cluster state changes (pods dead, etc)
- **Informs** Scheduler for re-scheduling
- Implementation: **kube-controller-manager**



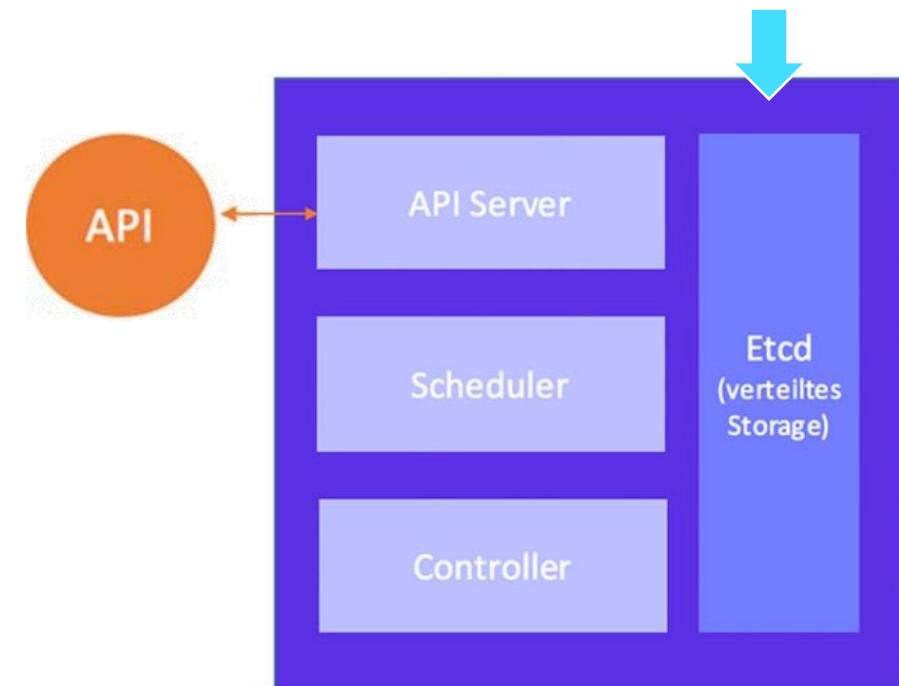
Source:

<https://knowitinfo.com/what-is-kubernetes-control-plane/>
<https://kubernetes.io/docs/concepts/overview/components/>

Control Plane Components - Etcd

Etcd

- a distributed **key-value** (JSON) store **database**
 - Any information about the cluster
 - e.g. **Cluster configuration**
 - All state changes (cluster changes)
 - e.g. Cluster resource condition
 - e.g. Actual state and desired state of the system
 - Secure
 - automatic Transport Layer Security (TLS)
 - Fast and highly-available
 - Over **30,000 requests per second** under heavy load
 - **Raft** consensus algorithm
 - Application data is **NOT** stored in Etcd.



Source:

<https://knowitinfo.com/what-is-kubernetes-control-plane/>

<https://kubernetes.io/docs/concepts/overview/components/>

<https://etcd.io/docs/v3.4/op-guide/performance/#:~:text=In%20common%20cloud%20environments%2C%20like,per%20second%20under%20heavy%20load.>

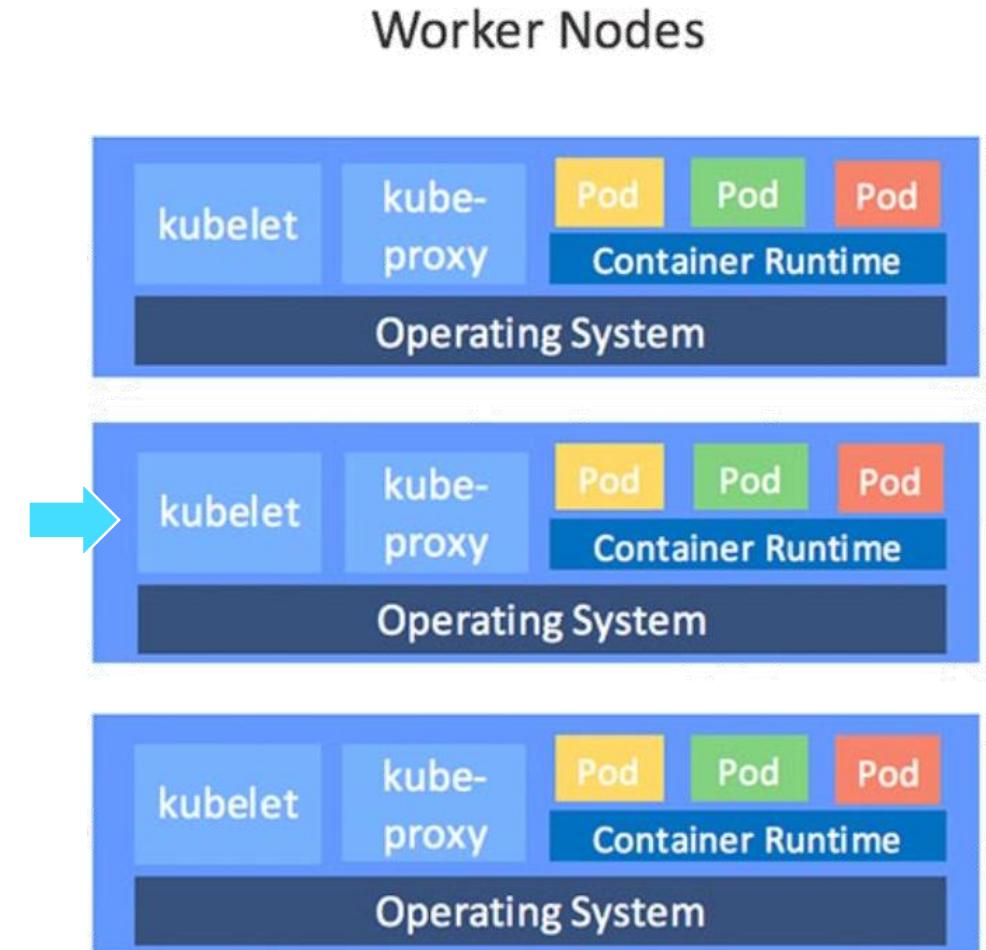
Worker Node Components - Kubelet

Revisited - Worker Nodes (Slave Nodes)

1. Watch the API server for new work assignments
2. Execute new work assignments
3. Report back to the control plane (via the API server)

Kubelet

- **Core agent** - runs on **every worker node** in the cluster
- **Interacts** with both the **container** and **node**
 - **Registers** the node(s) within the cluster
 - **Watches** the API server,
 - **Executes** the task - **starts** the pod with a container inside
 - **Maintains** the reporting channel and **Reports** task failure



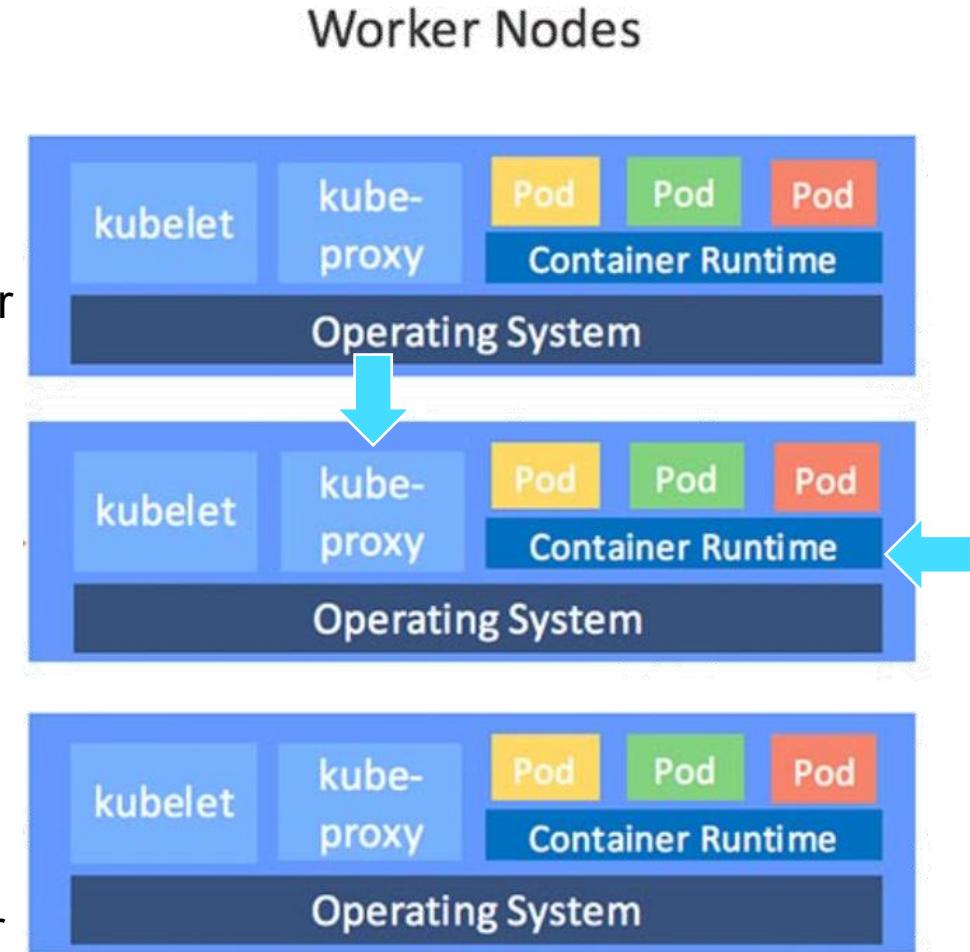
Worker Node Components – CR and K-Proxy

Container Runtime

- Software that is responsible for running containers.
- Perform container-related tasks
- Implementations: e.g. [containerd](#) and Kubernetes [CRI](#) (Container Runtime Interface)

Kube-proxy

- Network proxy that is responsible for local cluster networking
- Makes sure each node gets its unique IP address
- Maintains network rules on nodes
 - Handle routing and load-balancing (IPTABLES / IPVS)
- Implements parts of [K8s Service](#) concept
 - forwards the requests for communication, like a smart balancer

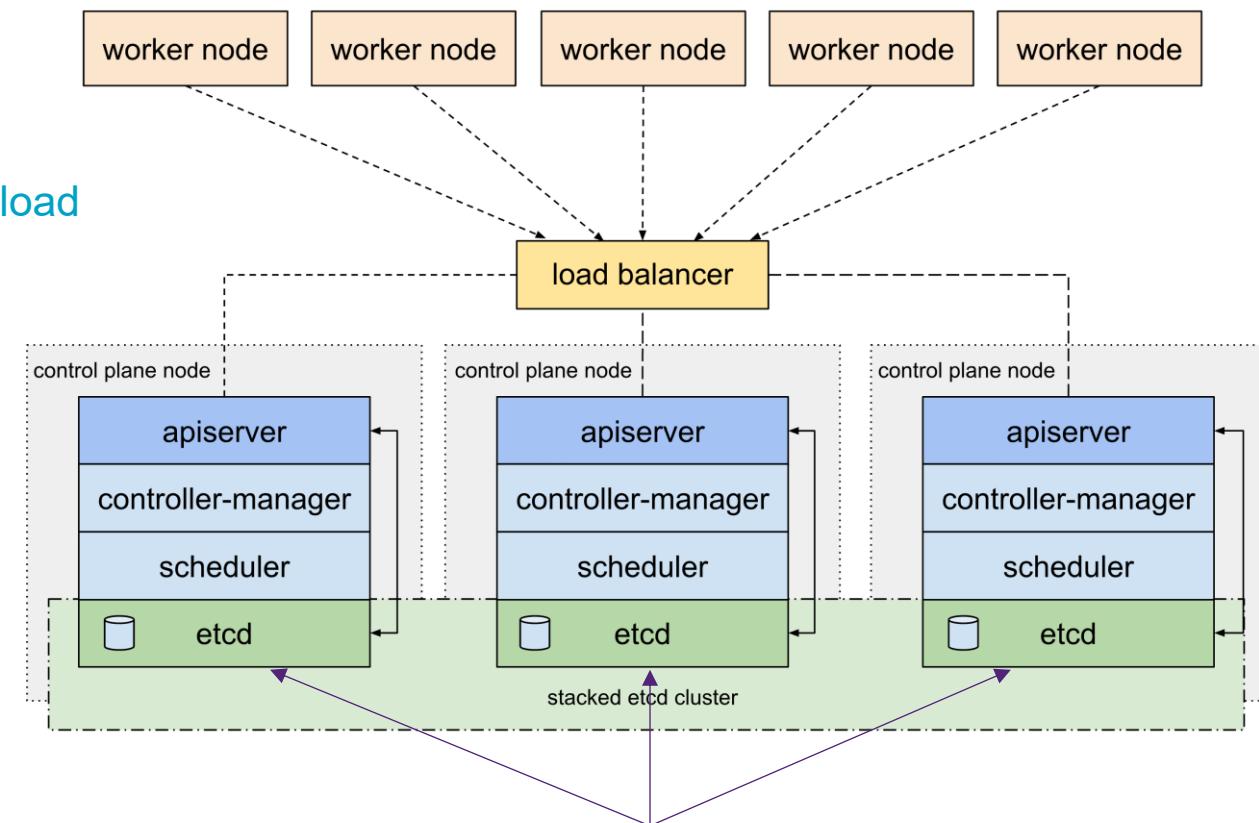


High-Availability for Control Plane (Master Node)

1. With stacked control plane nodes:

- Each CP runs an instance of the [kube-apiserver](#), [kube-scheduler](#), and [kube-controller-manager](#).
- The [kube-apiserver](#) is exposed to worker nodes using a [load balancer](#).
- Each CP node creates a [local etcd member](#) and it communicates only with the [kube-apiserver](#) of this node.
- CPs and etcd members on the same nodes.
 - [simpler to set up than an external etcd cluster](#)
 - [simpler to manage](#) for replication
 - [requires less infrastructure](#) (etcd members and control plane nodes are co-located)
- a stacked cluster runs the [risk of failed coupling](#).
- [a minimum of three](#) stacked CP nodes for an HA cluster.

kubeadm HA topology - stacked etcd



Etcd is stacked on top of the cluster formed by the nodes.

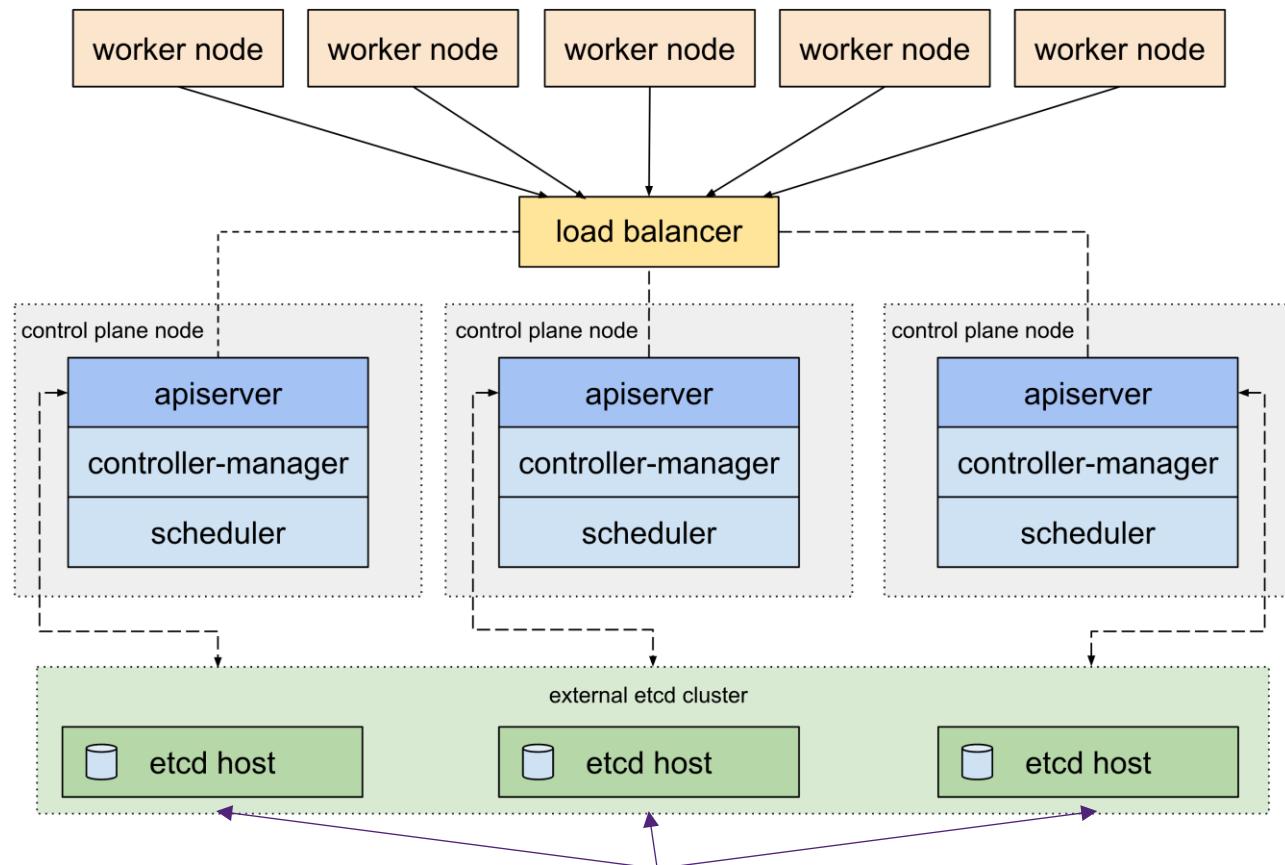
Multiple Control Planes (Master Nodes)

2. With an external etcd cluster:

- etcd is an external cluster.
- Like the stacked etcd topology, each CP node runs an instance of the kube-apiserver, kube-scheduler, and kube-controller-manager.
- The kube-apiserver is exposed to worker nodes using a load balancer.
- CP and its etcd member are decoupled:
 - etcd members run on separate hosts,
 - Each etcd host communicates with the kube-apiserver of each CP node.
- The CP's failure has less impact and does not affect the cluster redundancy compared to stacked etcd topology
- But it requires more infrastructure (2x more hosts)

<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/high-availability/>
<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/ha-topology/>

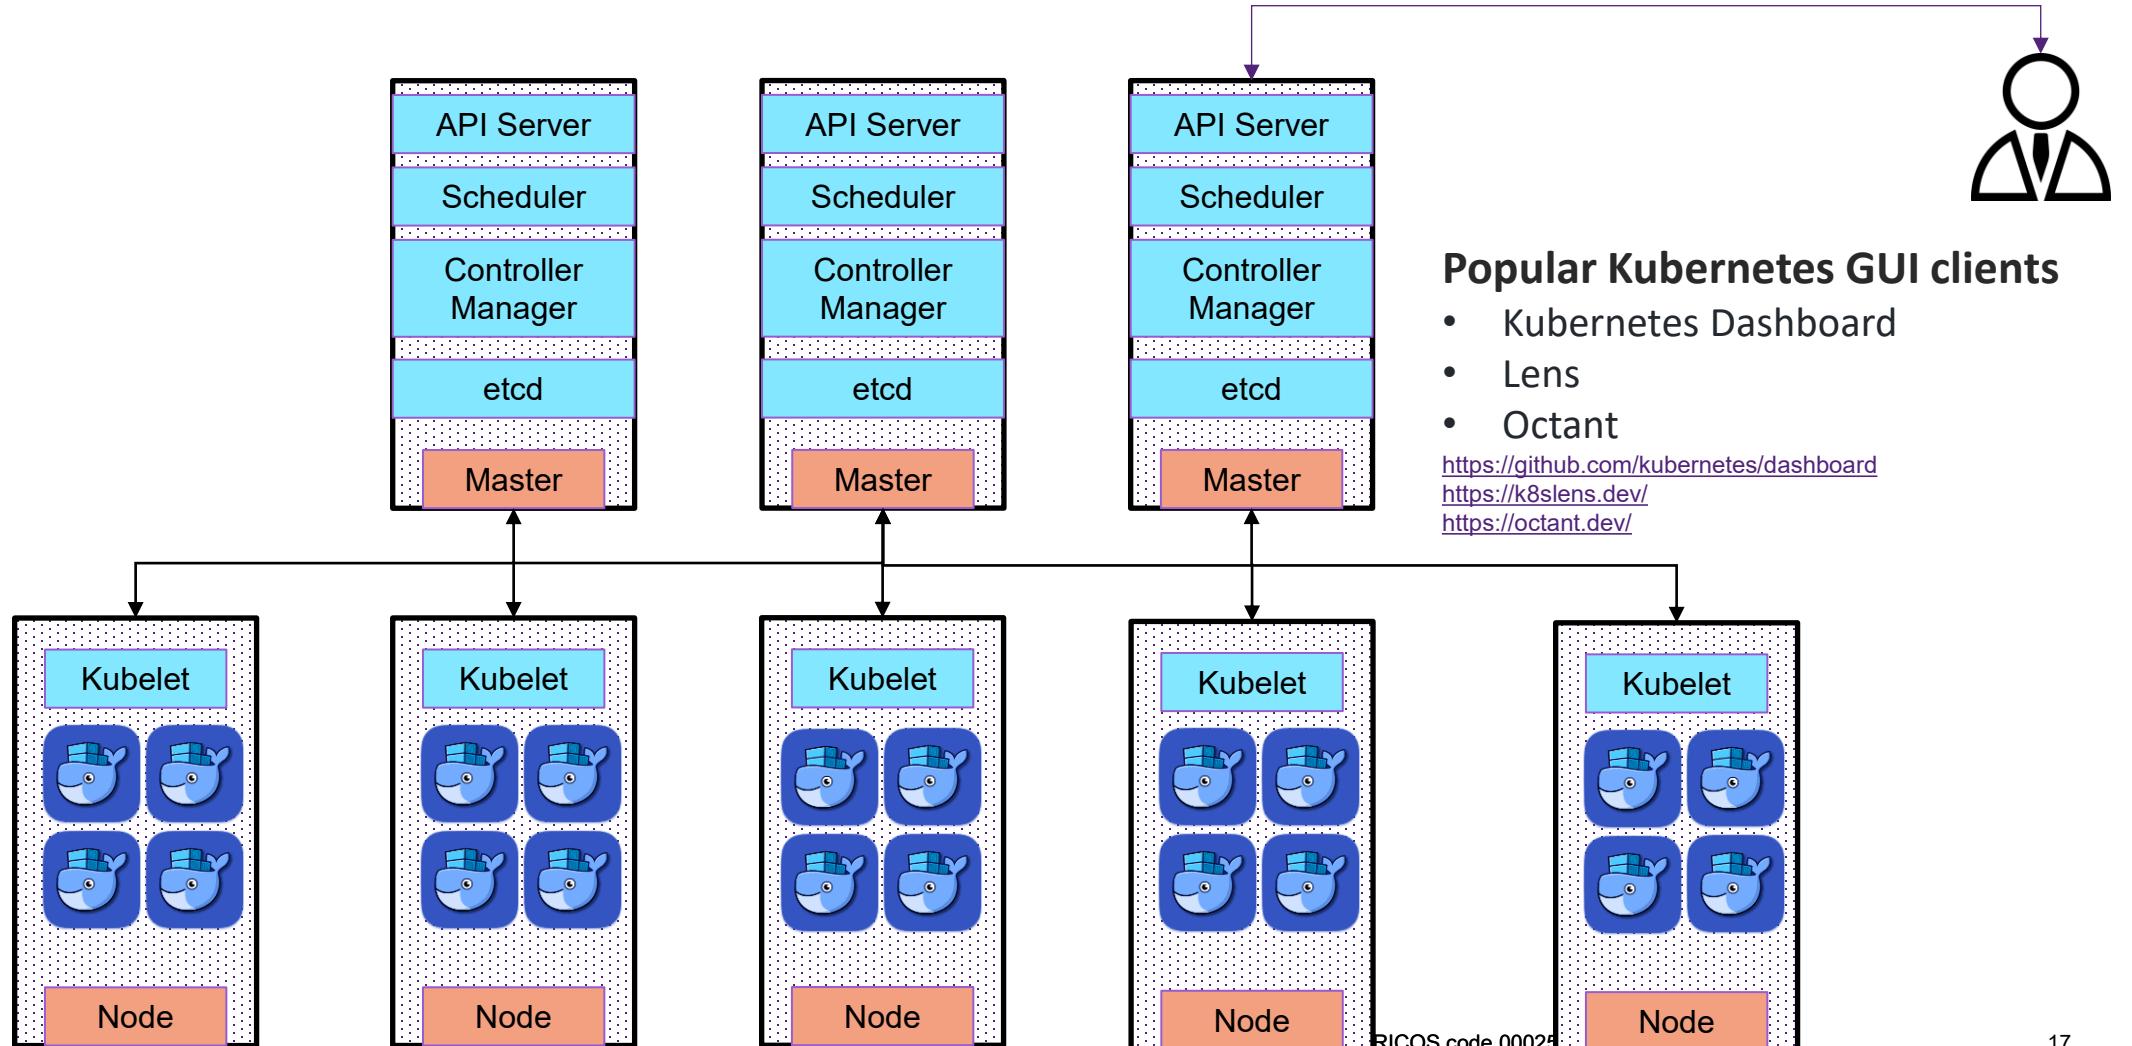
kubeadm HA topology - external etcd



Three extra hosts for the external etcd cluster

Example of K8s

Client apps
e.g. kubectl (CLI)

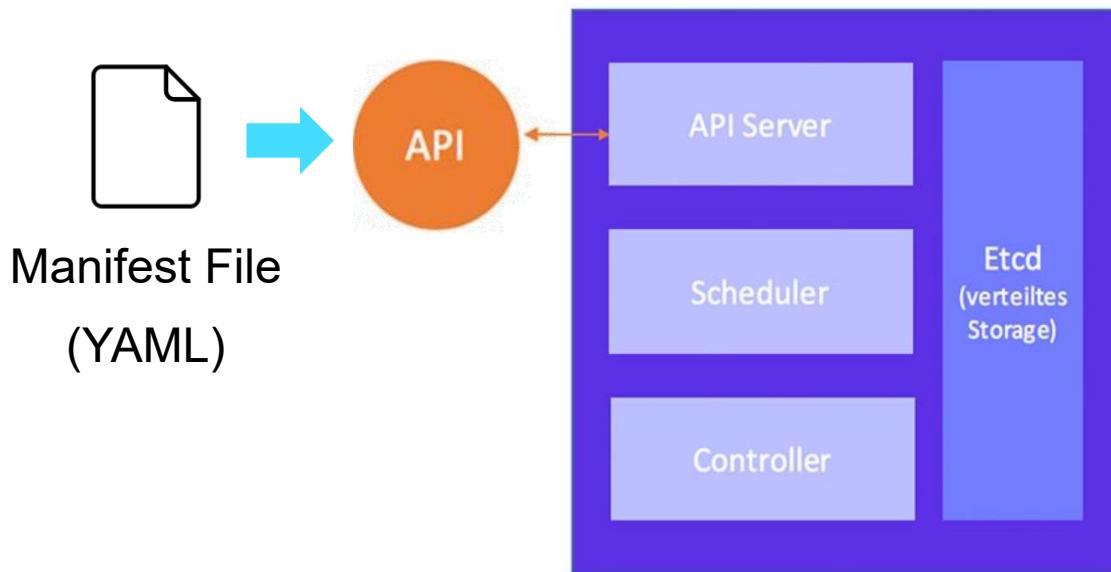


Outline

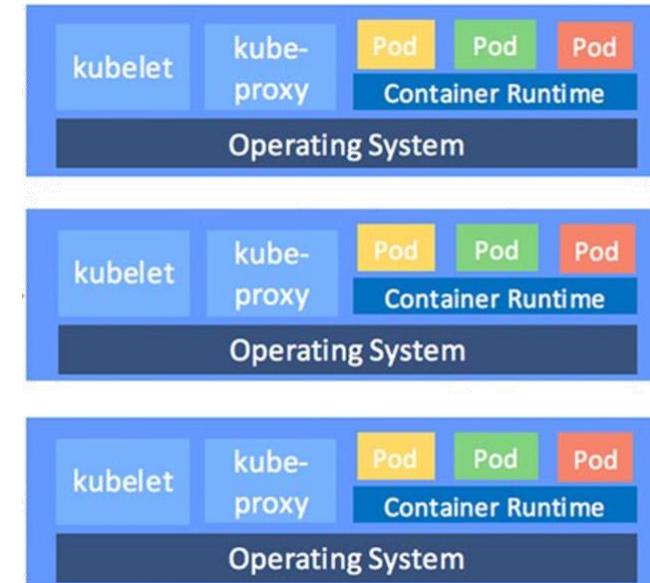
- Introduction to K8s
- K8s Architecture and its Components
 - Control Plane Node and Work Node
- • Declarative Model in K8s
 - Reconciliation Loop
 - Minikube & Kubectl
 - Demo of GKE - Cluster creation
 - K8s Objects:
 - Pod & ReplicaSet & Deployment
 - & Service & Ingress & StatefulSet
 - & ConfigMap & Secrets & Volumes
 - A K8s deployment example - Mysql & phpMyAdmin

Declarative Model

- The declarative model and the concept of desired state are at the very heart of K8s.
- In K8s, the declarative model works like:



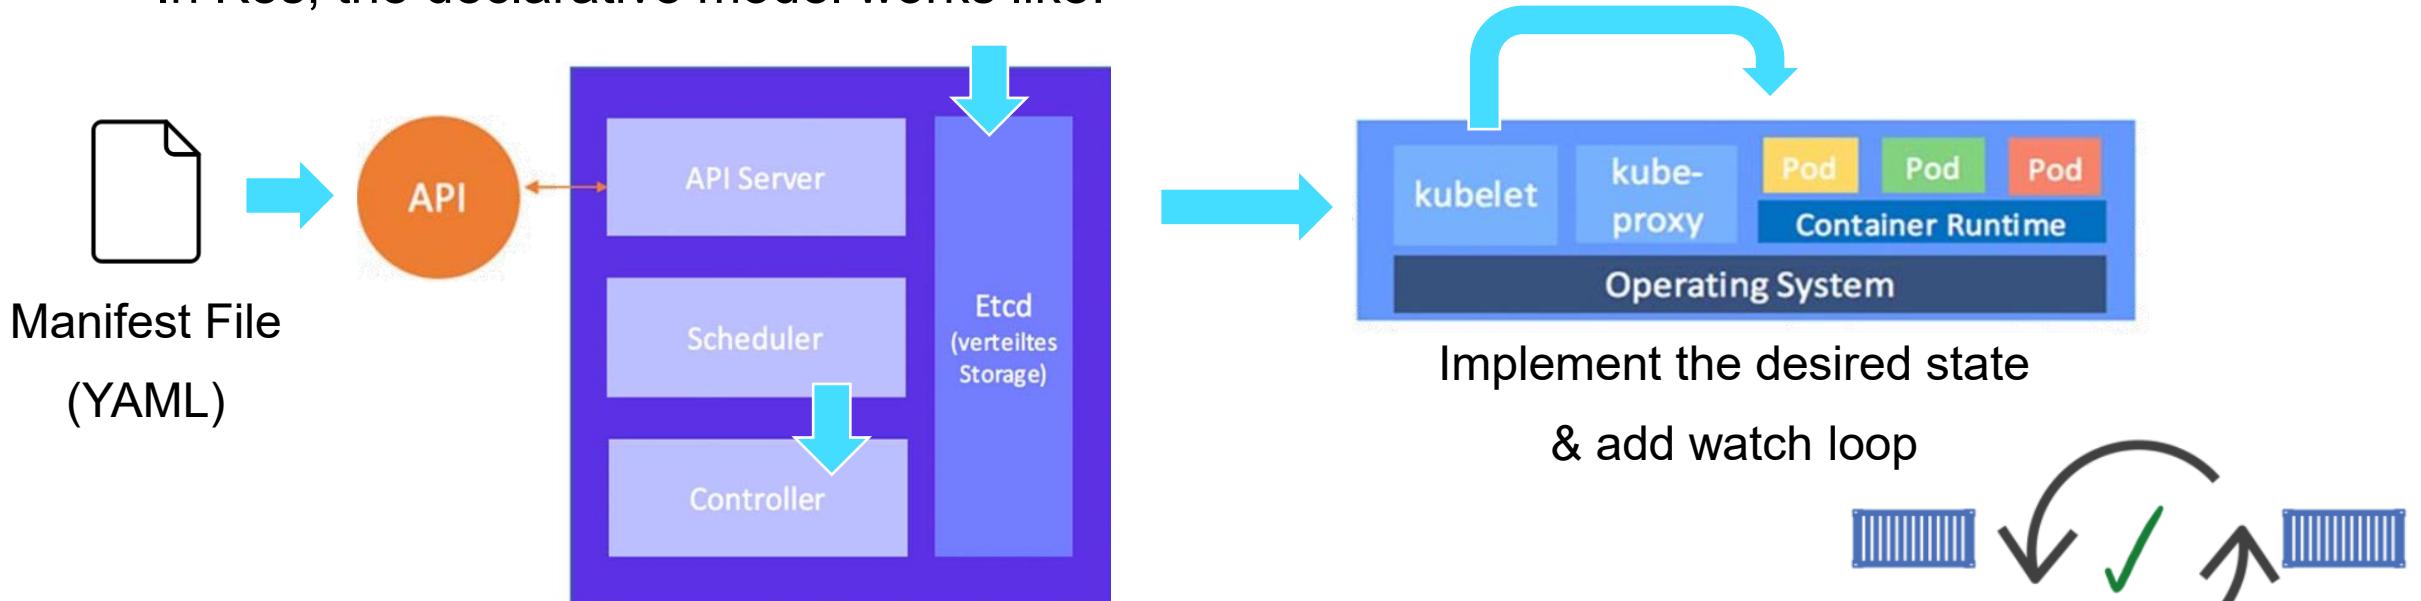
Manifest File
(YAML)



1. Declare the desired state of an application (microservice) in a manifest file
2. POST it to the API server

Declarative Model

- The declarative model and the concept of desired state are at the very heart of K8s.
- In K8s, the declarative model works like:



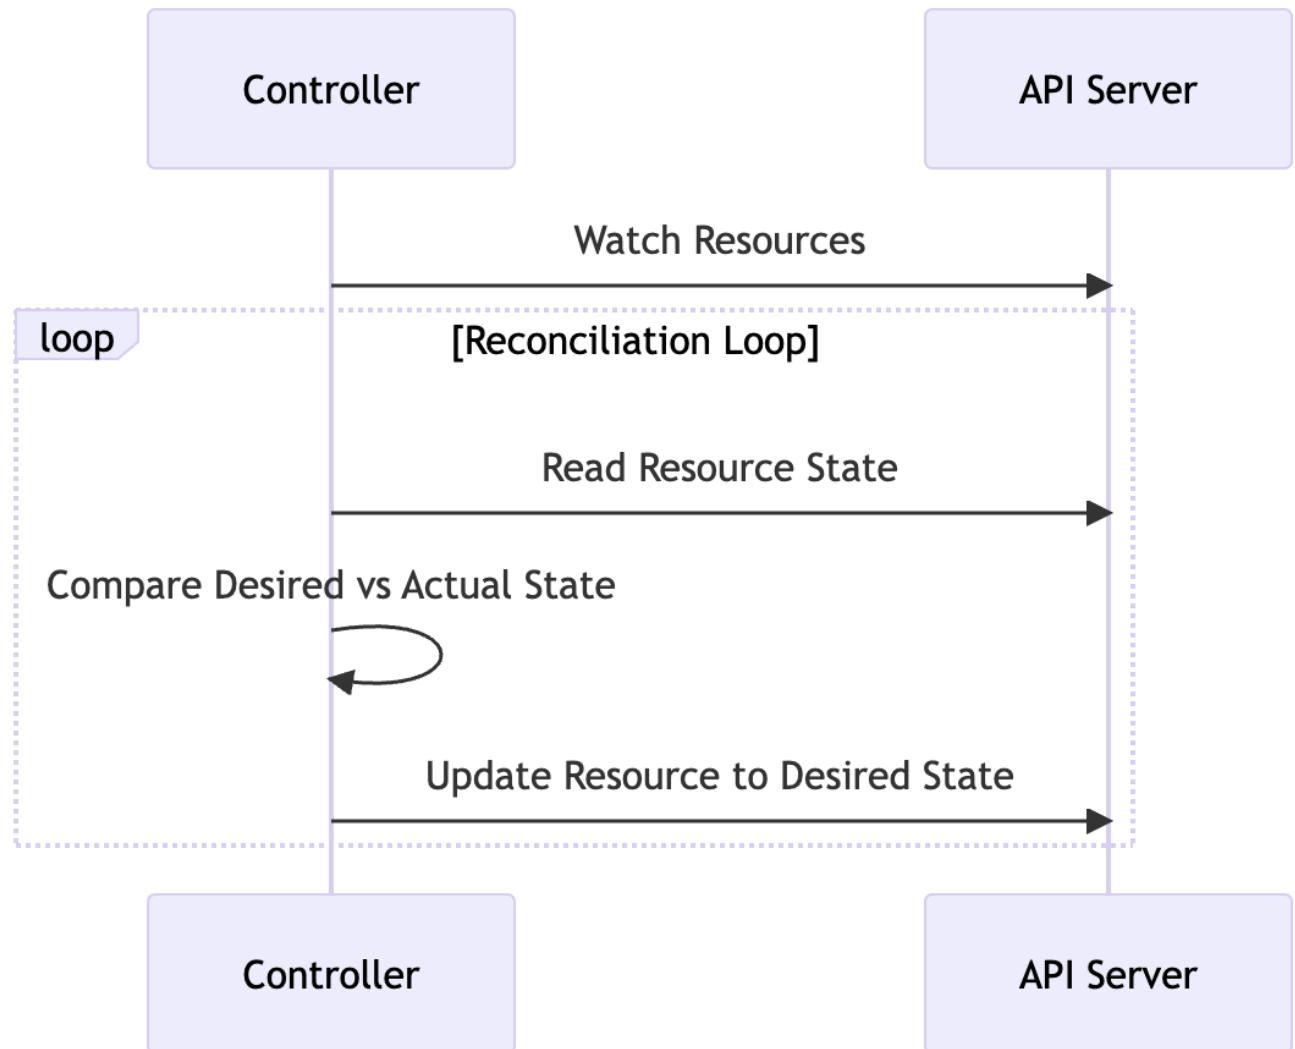
3. Kubernetes stores it in the Etcd as the application's desired state
4. Kubernetes implements the desired state on the cluster
5. Kubernetes implements watch loops -> current state of the application doesn't vary from the desired state

Outline

- Introduction to K8s
- K8s Architecture and its Components
 - Control Plane Node and Work Node
- Declarative Model in K8s
- • Reconciliation Loop
- Minikube & Kubectl
- Demo of GKE - Cluster creation
- K8s Objects:
 - Pod & ReplicaSet & Deployment
 - & Service & Ingress & StatefulSet
 - & ConfigMap & Secrets & Volumes
- A K8s deployment example - Mysql & phpMyAdmin

Reconciliation Loop

In Kubernetes, the Reconciliation Loop ensures that the system's actual state always matches the user's desired state. By continuously detecting and correcting differences, it provides self-healing and consistency in a dynamic, unreliable environment.



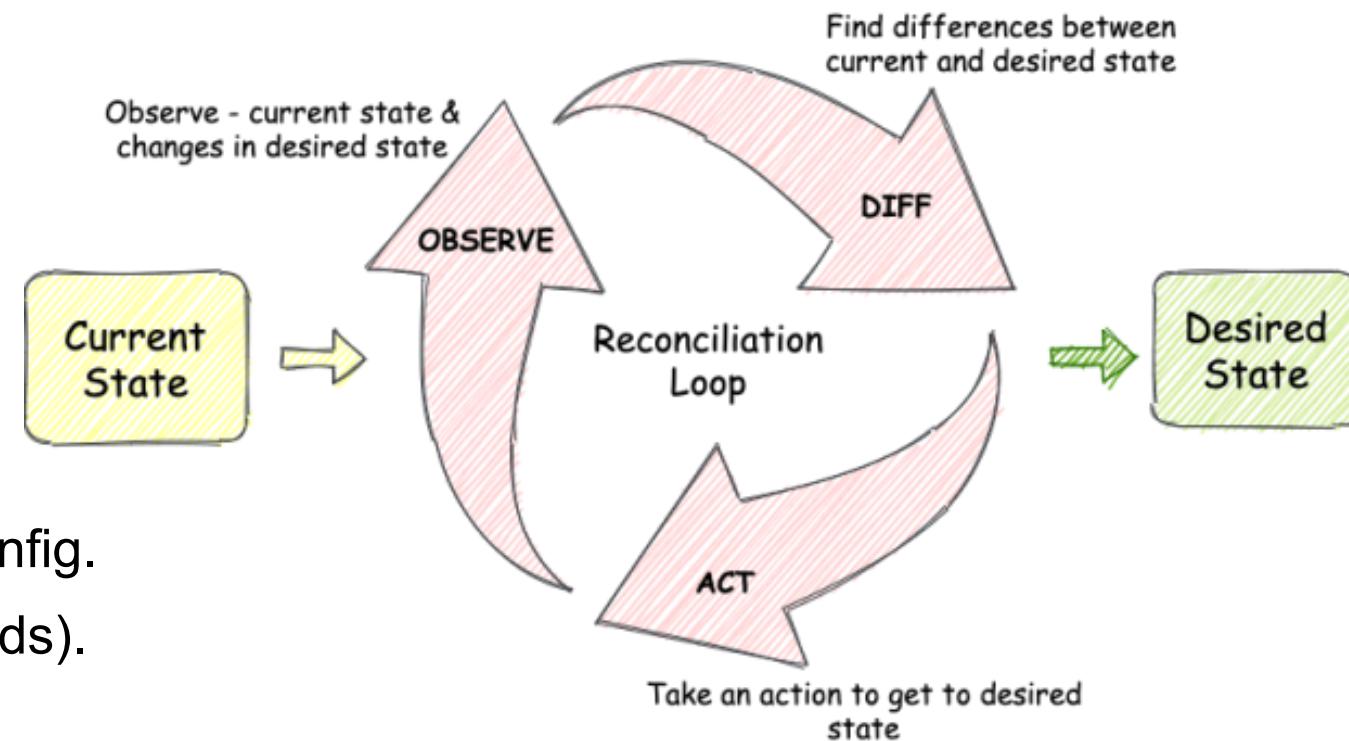
Reconciliation Loop

Goal:

- Make the cluster's current state match the desired state.

How it works:

- Observe – Check what's running now.
- Compare – See if it matches the YAML config.
- Act – Fix differences (e.g., start or stop pods).



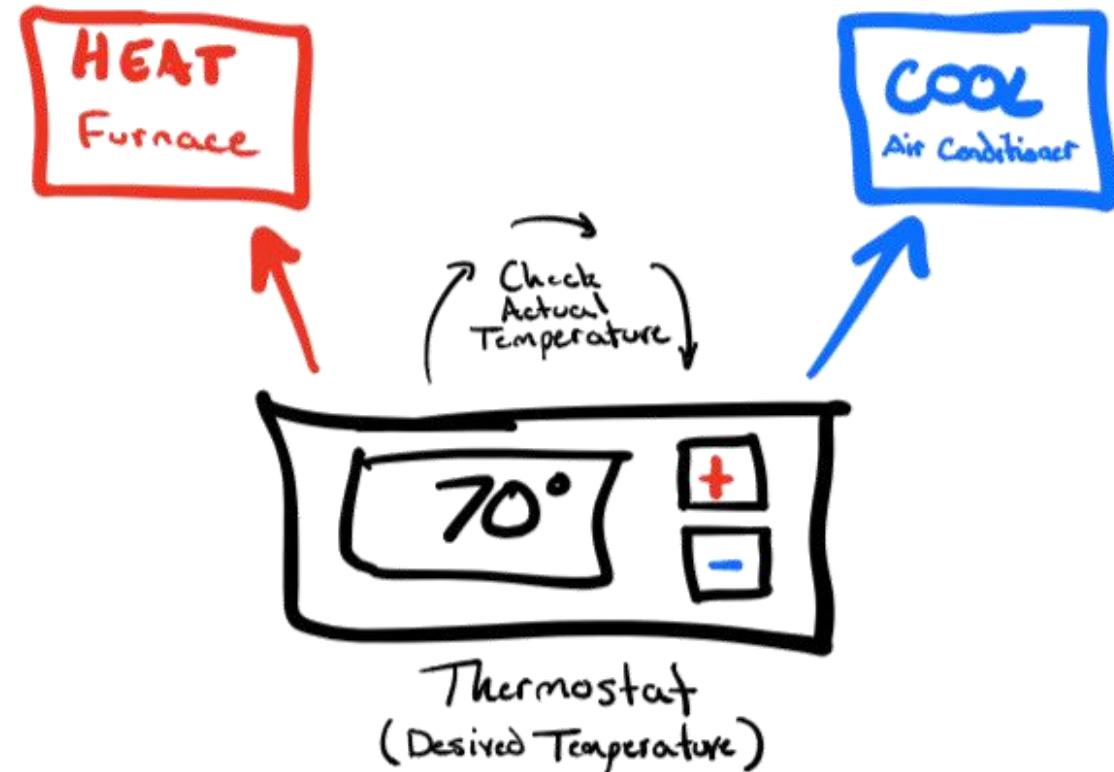
An Intuitive Example: The Thermostat

Desired State: You set the thermostat to 70°C. This is the state you want.

Current State: The thermostat's sensor constantly measures the actual room temperature, let's say it's 70°C.

The Reconciliation Loop in Action:

- Observe: The thermostat reads the current temperature (60°C).
- Compare: It compares 60°C to your desired temperature of 70°C. They don't match.
- Act: Because the current temperature is too low, the thermostat sends a signal to turn on the heater.

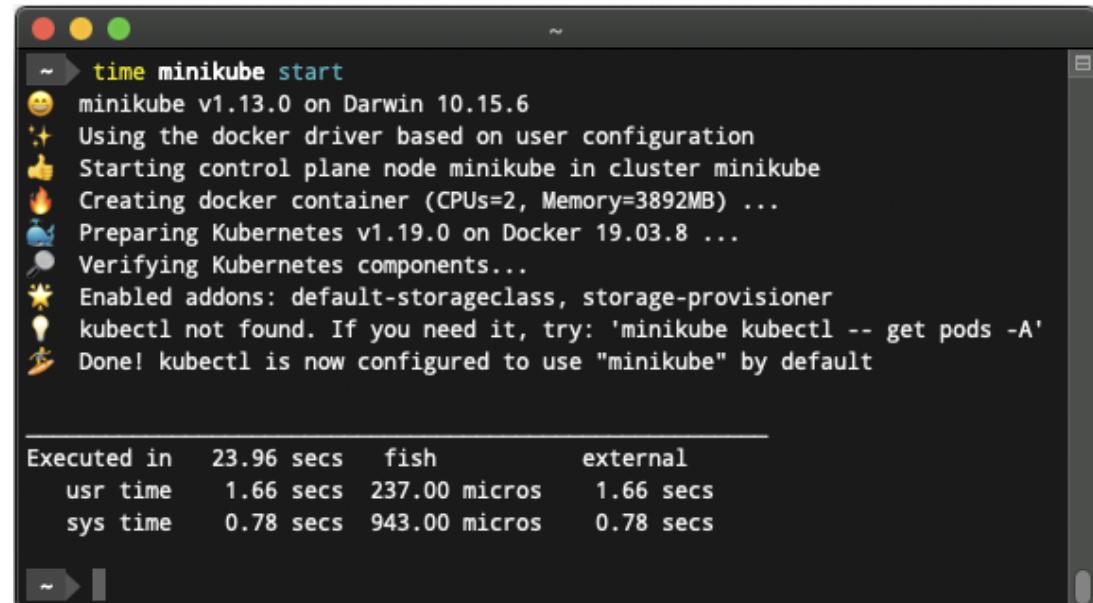


Outline

- Introduction to K8s
- K8s Architecture and its Components
 - Control Plane Node and Work Node
- Declarative Model in K8s
- Reconciliation Loop
- • Minikube & Kubectl
- Demo of GKE - Cluster creation
- K8s Objects:
 - Pod & ReplicaSet & Deployment
 - & Service & Ingress & StatefulSet
 - & ConfigMap & Secrets & Volumes
- A K8s deployment example - Mysql & phpMyAdmin

Minikube

- [minikube](#) quickly sets up a local Kubernetes cluster on
 - macOS, Linux, and Windows.
- Supports the [latest](#) Kubernetes release
- Supports [multiple container runtimes](#) (CRI-O, containerd, docker)
- Provides [advanced](#) features:
 - Load Balancing,
 - Filesystem mounts
 - Network policy and etc.
- Support [add-ons](#) (extensions of minikube for added functionality)
- Useful for application developers and helpful for new Kubernetes beginners.
- Alternatives to Minikube:
 - Kind, micro-k8s, etc.



```
~ ➤ time minikube start
😊 minikube v1.13.0 on Darwin 10.15.6
✨ Using the docker driver based on user configuration
👍 Starting control plane node minikube in cluster minikube
🔥 Creating docker container (CPUs=2, Memory=3892MB) ...
🌐 Preparing Kubernetes v1.19.0 on Docker 19.03.8 ...
🌐 Verifying Kubernetes components...
⭐ Enabled addons: default-storageclass, storage-provisioner
💡 kubectl not found. If you need it, try: 'minikube kubectl -- get pods -A'
💡 Done! kubectl is now configured to use "minikube" by default

Executed in 23.96 secs fish external
    usr time 1.66 secs 237.00 micros 1.66 secs
    sys time 0.78 secs 943.00 micros 0.78 secs
```

<https://minikube.sigs.k8s.io/docs/>

Kubernetes command-line tool - kubectl

- Run commands against Kubernetes clusters.
 - deploy applications (`create/apply pod, deployment, service, pv/pvc`),
 - inspect and manage cluster resources (`get/delete`),
 - view logs (`describe`)
- `kubectl` command syntax:
 - `kubectl [command] [TYPE] [NAME] [flags]`
 - `[command]`: specify the operation do you want to perform on resources,
 - E.g.: `create, get, describe, delete, apply`.
 - `[TYPE]`: specify resource type.
 - E.g.: `pods, nodes, services, etc.`
 - `[NAME]`: specify the name of the resource (case-sensitive).
 - `[flags]`: specify optional flags.
 - E.g.: use “`-s`” or “`--server`” flags to specify the address and port of the Kubernetes API server.

Outline

- Introduction to K8s
- K8s Architecture and its Components
 - Control Plane Node and Work Node
- Declarative Model in K8s
- Reconciliation Loop
- Minikube & Kubectl
- • Demo of GKE - Cluster creation
- K8s Objects:
 - Pod & ReplicaSet & Deployment
 - & Service & Ingress & StatefulSet
 - & ConfigMap & Secrets & Volumes
- A K8s deployment example - Mysql & phpMyAdmin

Demo: Google Kubernetes Engine (GKE) – Create a Cluster

The screenshot shows the Google Cloud Platform navigation bar on the left and a detailed view of the Kubernetes Engine interface on the right.

Navigation Bar:

- App Engine
- Compute Engine
- Kubernetes Engine** (selected, indicated by a red box around the icon)
- Cloud Functions
- Cloud Run
- VMware Engine

Storage (category)

Bigtable

Kubernetes Engine Submenu (opened from the navigation bar):

- Clusters** (selected, indicated by a red box around the text)
- Workloads
- Services & Ingress
- Applications
- Configuration
- Storage
- Object Browser
- Migrate to containers

Kubernetes Engine Main View (right side):

Kubernetes Engine

Kubernetes clusters

Containers package an application so that it can be easily deployed to run in its own isolated environment. Containers are managed in clusters that automate VM creation and maintenance. [Learn more](#)

[Create cluster](#) [Deploy container](#) [Take the quickstart](#)

Demo: GKE – Create a Cluster

[+ ADD NODE POOL](#) [REMOVE NODE POOL](#)

SUMMARY Cluster summary

Cluster basics

The new cluster will be created with the name, version, and in the location you specify here. After the cluster is created, name and location can't be changed.

To experiment with an affordable cluster, try My first cluster in the Cluster set-up guides

Name cluster-demo

Location type Zonal

Zone australia-southeast1-a

Specify default node locations [?](#)
Current default: australia-southeast1-a

Control plane version

Choose a release channel for automatic management of your cluster's version and upgrade cadence. Choose a static version for more direct management of your cluster's version. [Learn more](#).

Static version

Release channel

Release channel Regular channel (default)

Cluster set-up guides

My first cluster An affordable cluster to experiment with **Cost-optimized cluster**

Immutable properties

⚠ These properties can't be changed after creation:

Name	cluster-demo	Status	Name	Location	Number of nodes	Total vCPUs	Total memory	Notifications	Labels
Location type	Zonal	⚠	cluster-1	australia-southeast1-a	3	6	12 GB	–	⋮
Zone	australia-southeast1-a	✓	cluster-demo	australia-southeast1-a	3	6	12 GB	–	⋮
Private cluster	Disabled								
Network	default								

[▼ SHOW MORE](#)

Cost

You will be billed for 3 nodes.

[▼ SHOW DETAILS](#)

[Learn more](#)

Beta properties

You haven't selected any properties in beta.

Limitations

⚠ Pod address range limits the maximum number of nodes.

[Learn more](#)

Node pools

default-pool

Connect to the cluster

You can connect to your cluster via command-line or using a dashboard.

Command-line access

Configure [kubectl](#) command line access by running the following command:

```
$ gcloud container clusters get-credentials cluster-demo --zone australia-southeast1-a --project mythic-dynamo-300704
```

[RUN IN CLOUD SHELL](#)

Demo: GKE – Log in using Cloud Shell

```
uqyluo@cloudshell:~ (mythic-dynamo-300704)$ kubectl cluster-info
Kubernetes control plane is running at https://35.197.191.106
GLBCDefaultBackend is running at https://35.197.191.106/api/v1/namespaces/kube-system/services/default-http-backend:http/proxy
KubeDNS is running at https://35.197.191.106/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy
Metrics-server is running at https://35.197.191.106/api/v1/namespaces/kube-system/services/https:metrics-server:/proxy
```

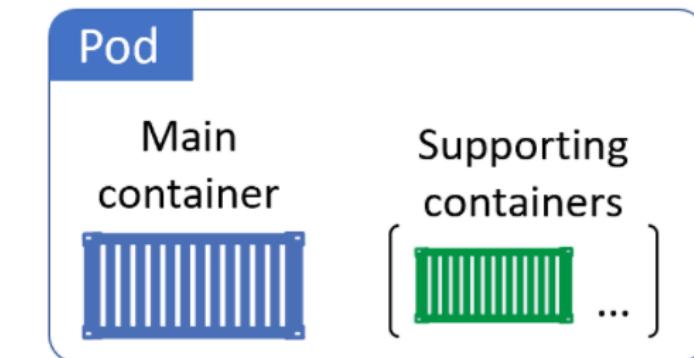
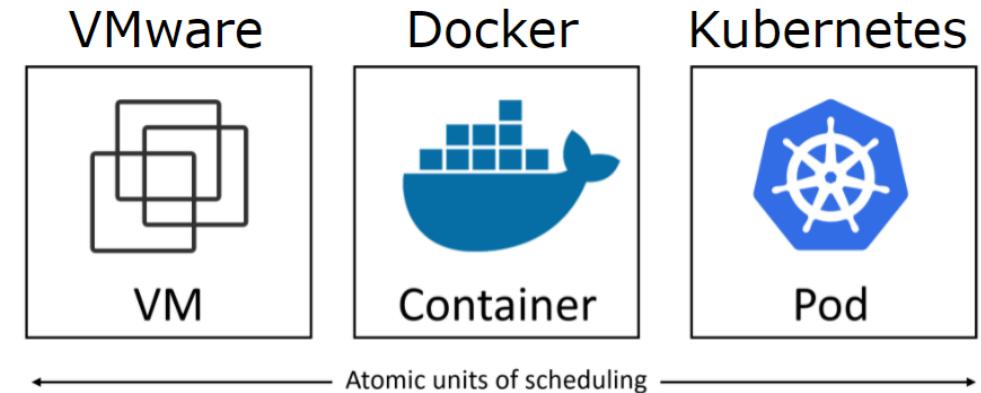
```
uqyluo@cloudshell:~ (mythic-dynamo-300704)$ kubectl get nodes
NAME                  STATUS   ROLES      AGE     VERSION
gke-cluster-demo-default-pool-368af1f8-6drq   Ready    <none>    8m30s   v1.20.8-gke.900
gke-cluster-demo-default-pool-368af1f8-19d6   Ready    <none>    8m32s   v1.20.8-gke.900
gke-cluster-demo-default-pool-368af1f8-xh3q   Ready    <none>    8m32s   v1.20.8-gke.900
```

Outline

- Introduction to K8s
- K8s Architecture and its Components
 - Control Plane Node and Work Node
- Declarative Models in K8s
- Reconciliation Loop
- Minikube & Kubectl
- Demo of GKE - Cluster creation
- • K8s Objects:
 - Pod & ReplicaSet & Deployment
 - & Service & Ingress & StatefulSet
 - & ConfigMap & Secrets & Volumes
 - A K8s deployment example - Mysql & phpMyAdmin

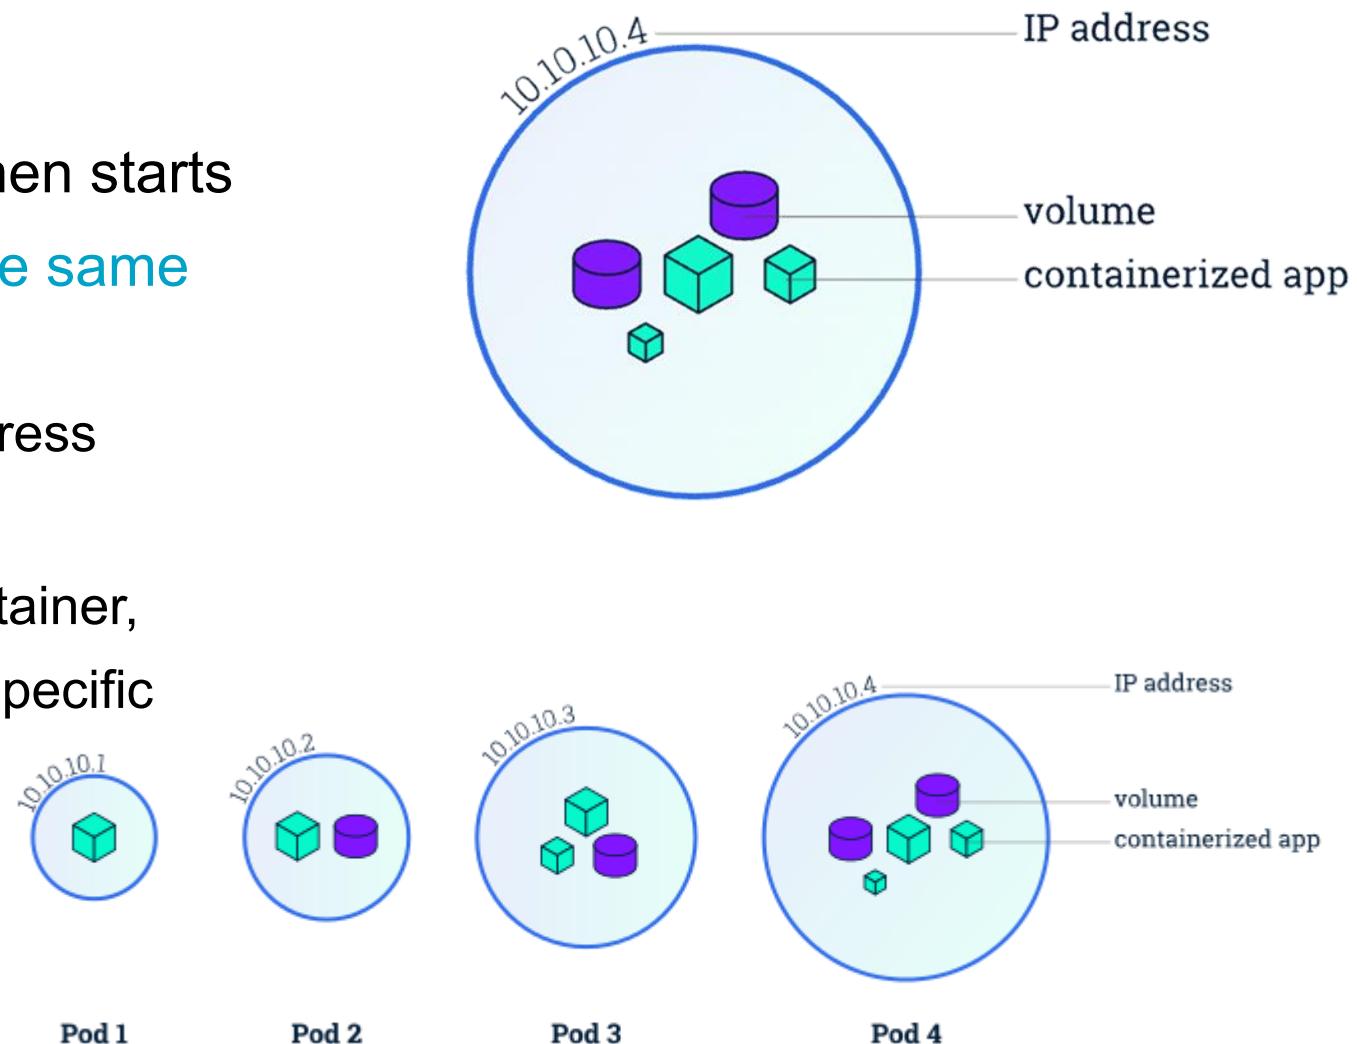
What is a Pod?

- Pod - The **atomic (smallest) unit** of scheduling in K8s
- Pod – The abstraction over a container(s)
 - Containers must **always** run inside of Pods
- Simplest model: **one container per Pod**; Advanced use cases that run multiple containers inside a single Pod
- How do isolated containers communicate?



What is a Pod?

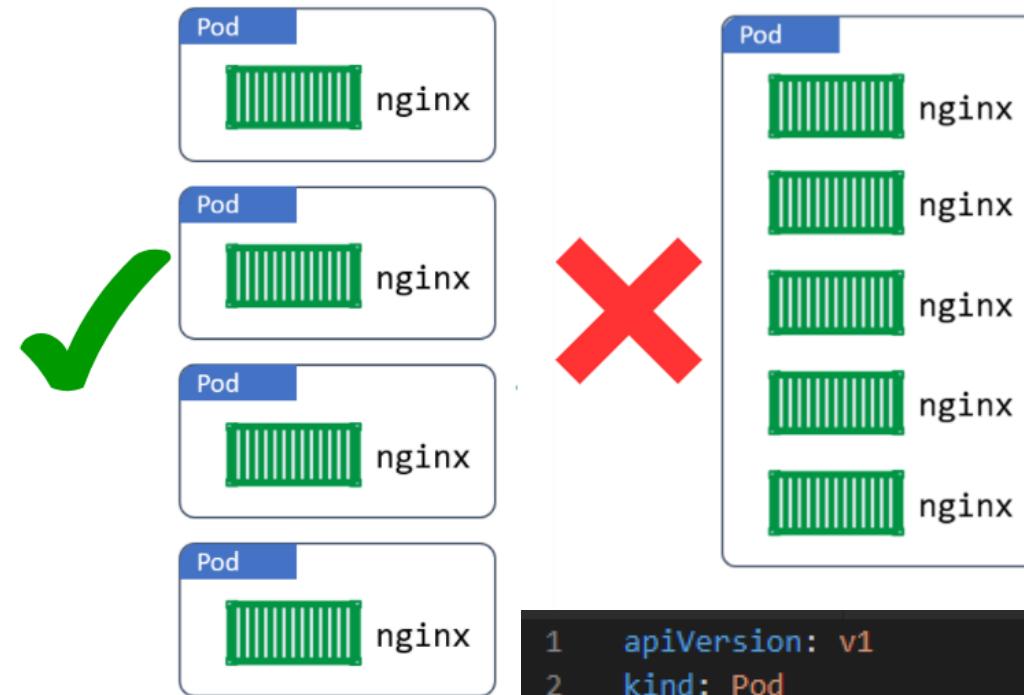
- Each pod gets its own IP address when starts
- Multiple containers in a Pod **share the same Pod environment**
 - Networking, as a unique cluster IP address
 - Shared storage, as Volumes
 - Information about how to run each container,
 - (e.g. the container image version or specific ports to use)



Pods are ephemeral (lasting for a very short time)!!

What is a Pod?

- **Scaling**: Pods are minimum unit of scaling
- **Deployment**: A single Pod can only be scheduled to a single node
- **Lifecycle**: New Pod is associated with a new ID and IP address



For an application to run on a Kubernetes cluster...

1. Being packaged as a container 
2. Being wrapped in a **Pod** 
3. Being deployed via a **declarative manifest file**

<https://kubernetes.io/docs/concepts/workloads/pods/>

CRICOS code 00025B

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: nginx
5 spec:
6   containers:
7     - name: nginx
8       image: nginx:1.14.2
9       ports:
10        - containerPort: 80
```

ReplicaSet

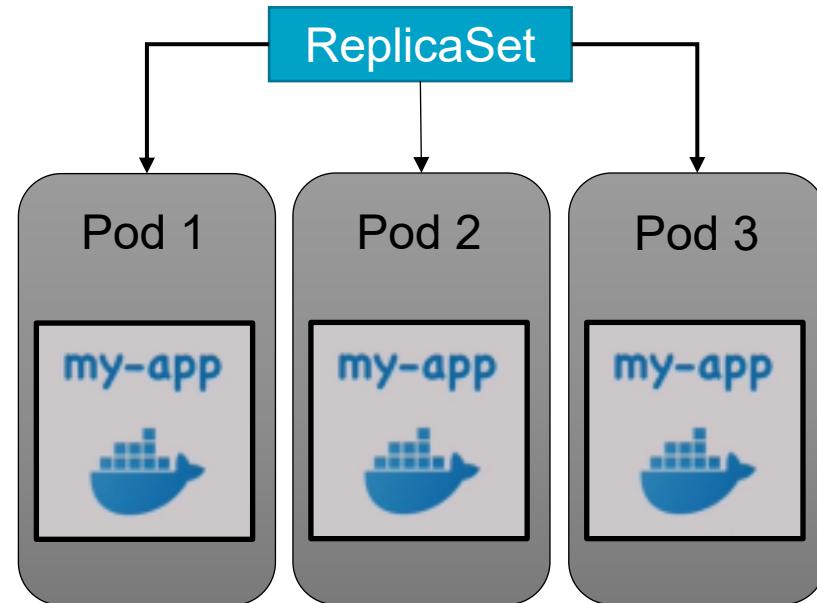
A ReplicaSet's purpose is to maintain a stable set of replica Pods running at any given time.

Field definitions:

- **#replicas** - indicating how many Pods it should be maintaining;
- a **selector** - specifying how to identify Pods it can acquire;
- a **pod template** - specifying the data of new Pods.

Dynamically creates and deletes Pods as needed to reach the desired number

- According to desired stats in the manifest file – yaml file.
- Run **kubectl create/apply - f**
 - E.g. `kubectl apply -f repset.yaml`



```

1  apiVersion: apps/v1
2  kind: ReplicaSet
3  metadata:
4    name: frontend
5    labels:
6      app: guestbook
7      tier: frontend
8  spec:
9    # modify replicas according to your case
10   replicas: 3
11   selector:
12     matchLabels:
13       tier: frontend
14   template:
15     metadata:
16       labels:
17         tier: frontend
18     spec:
19       containers:
20         - name: php-redis
21           image: gcr.io/google_samples/gb-frontend:v3
  
```

Demo: Deploy Replicsets

Replicaset YAML file

- `apiVersion`, `kind`, `metadata`
- `spec`:
 - `replicas`: desired number of replicas
 - `selector`: select which pods should be included
 - `template`: required field, same schema as a Pod specification

```
1  apiVersion: apps/v1
2  kind: ReplicaSet
3  metadata:
4    name: frontend
5    labels:
6      app: guestbook
7      tier: frontend
8  spec:
9    # modify replicas according to your case
10   replicas: 3
11   selector:
12     matchLabels:
13       tier: frontend
14   template:
15     metadata:
16       labels:
17         tier: frontend
18     spec:
19       containers:
20         - name: php-redis
21           image: gcr.io/google_samples/gb-frontend:v3
```

Demo: Deploy Replicasets

Run “kubectl describe rs cc-demo-replicaset”

```
Name:          cc-demo-replicaset
Namespace:     default
Selector:      tier=frontend
Labels:        app=guestbook
               tier=frontend
Annotations:   <none>
Replicas:      3 current / 3 desired
Pods Status:   3 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  tier=frontend
  Containers:
    php-redis:
      Image:      gcr.io/google_samples/gb-frontend:v3
      Port:       <none>
      Host Port: <none>
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
  Events:
    Type  Reason          Age   From           Message
    ----  -----          ----  ----           -----
    Normal SuccessfulCreate 4m28s replicaset-controller  Created pod: cc-demo-replicaset-hb5xz
    Normal SuccessfulCreate 4m28s replicaset-controller  Created pod: cc-demo-replicaset-kpv6r
    Normal SuccessfulCreate 4m28s replicaset-controller  Created pod: cc-demo-replicaset-q9mjl
```

Demo: Deploy Replicasets – Self-healing

- Destroy a Pod in the cluster as a simulation of failure

```
1 POD_NAME=$(kubectl get pods -o name \  
2 | tail -1)  
3 kubectl delete $POD_NAME
```

```
(base) ✘ sen@Sens-M1Pro ~ ➔ POD_NAME=$(kubectl get pods -o name | tail -1)  
(base) ✘ sen@Sens-M1Pro ~ ➔ kubectl delete $POD_NAME  
pod "cc-demo-replicaset-sqnf5" deleted  
(base) ✘ sen@Sens-M1Pro ~ ➔ kubectl get pods  
NAME READY STATUS RESTARTS AGE  
cc-demo-replicaset-2d42j 1/1 Running 0 2m1s  
cc-demo-replicaset-6pkrp 1/1 Running 0 35s  
cc-demo-replicaset-v48xb 1/1 Running 0 6s
```

Demo: Deploy Replicasets – Scaling

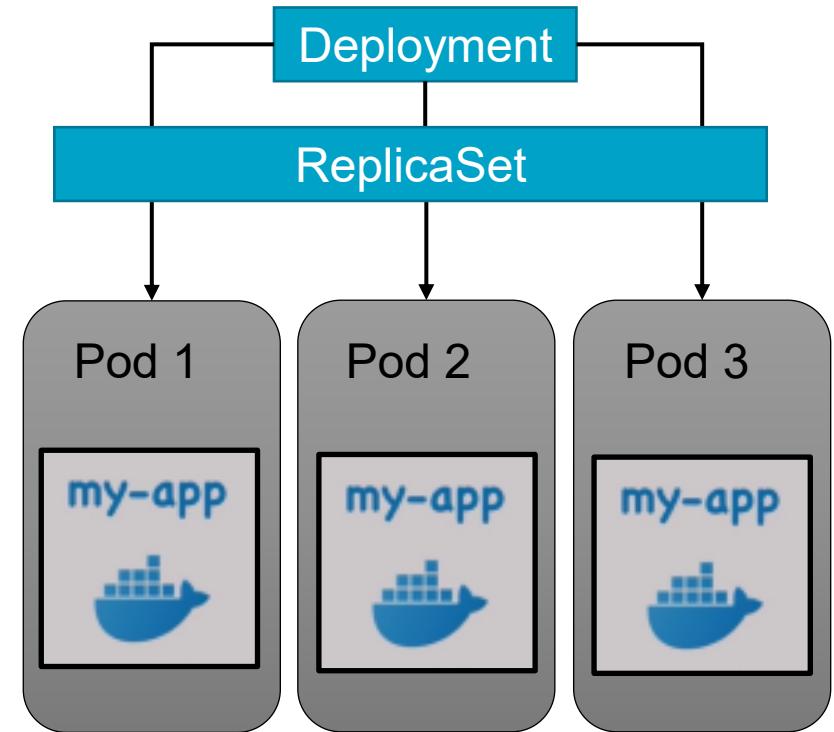
```
(base) sen@Sens-M1Pro ~ ➔ kubectl get rs
NAME          DESIRED   CURRENT   READY   AGE
cc-demo-replicaset   3        3        3      11m
(base) sen@Sens-M1Pro ~ ➔ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
cc-demo-replicaset-2d42j   1/1    Running   0          4m20s
cc-demo-replicaset-6pkrp   1/1    Running   0          2m54s
cc-demo-replicaset-v48xb   1/1    Running   0          2m25s
```

```
(base) sen@Sens-M1Pro ~ ➔ kubectl scale --replicas=5 rs cc-demo-replicaset
replicaset.apps/cc-demo-replicaset scaled
(base) sen@Sens-M1Pro ~ ➔ kubectl get rs
NAME          DESIRED   CURRENT   READY   AGE
cc-demo-replicaset   5        5        5      12m
(base) sen@Sens-M1Pro ~ ➔ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
cc-demo-replicaset-2d42j   1/1    Running   0          5m26s
cc-demo-replicaset-6pkrp   1/1    Running   0          4m
cc-demo-replicaset-82hr5   1/1    Running   0          14s
cc-demo-replicaset-jcxcj   1/1    Running   0          14s
cc-demo-replicaset-v48xb   1/1    Running   0          3m31s
```

Deployment

A Deployment provides declarative updates for Pods and ReplicaSets.

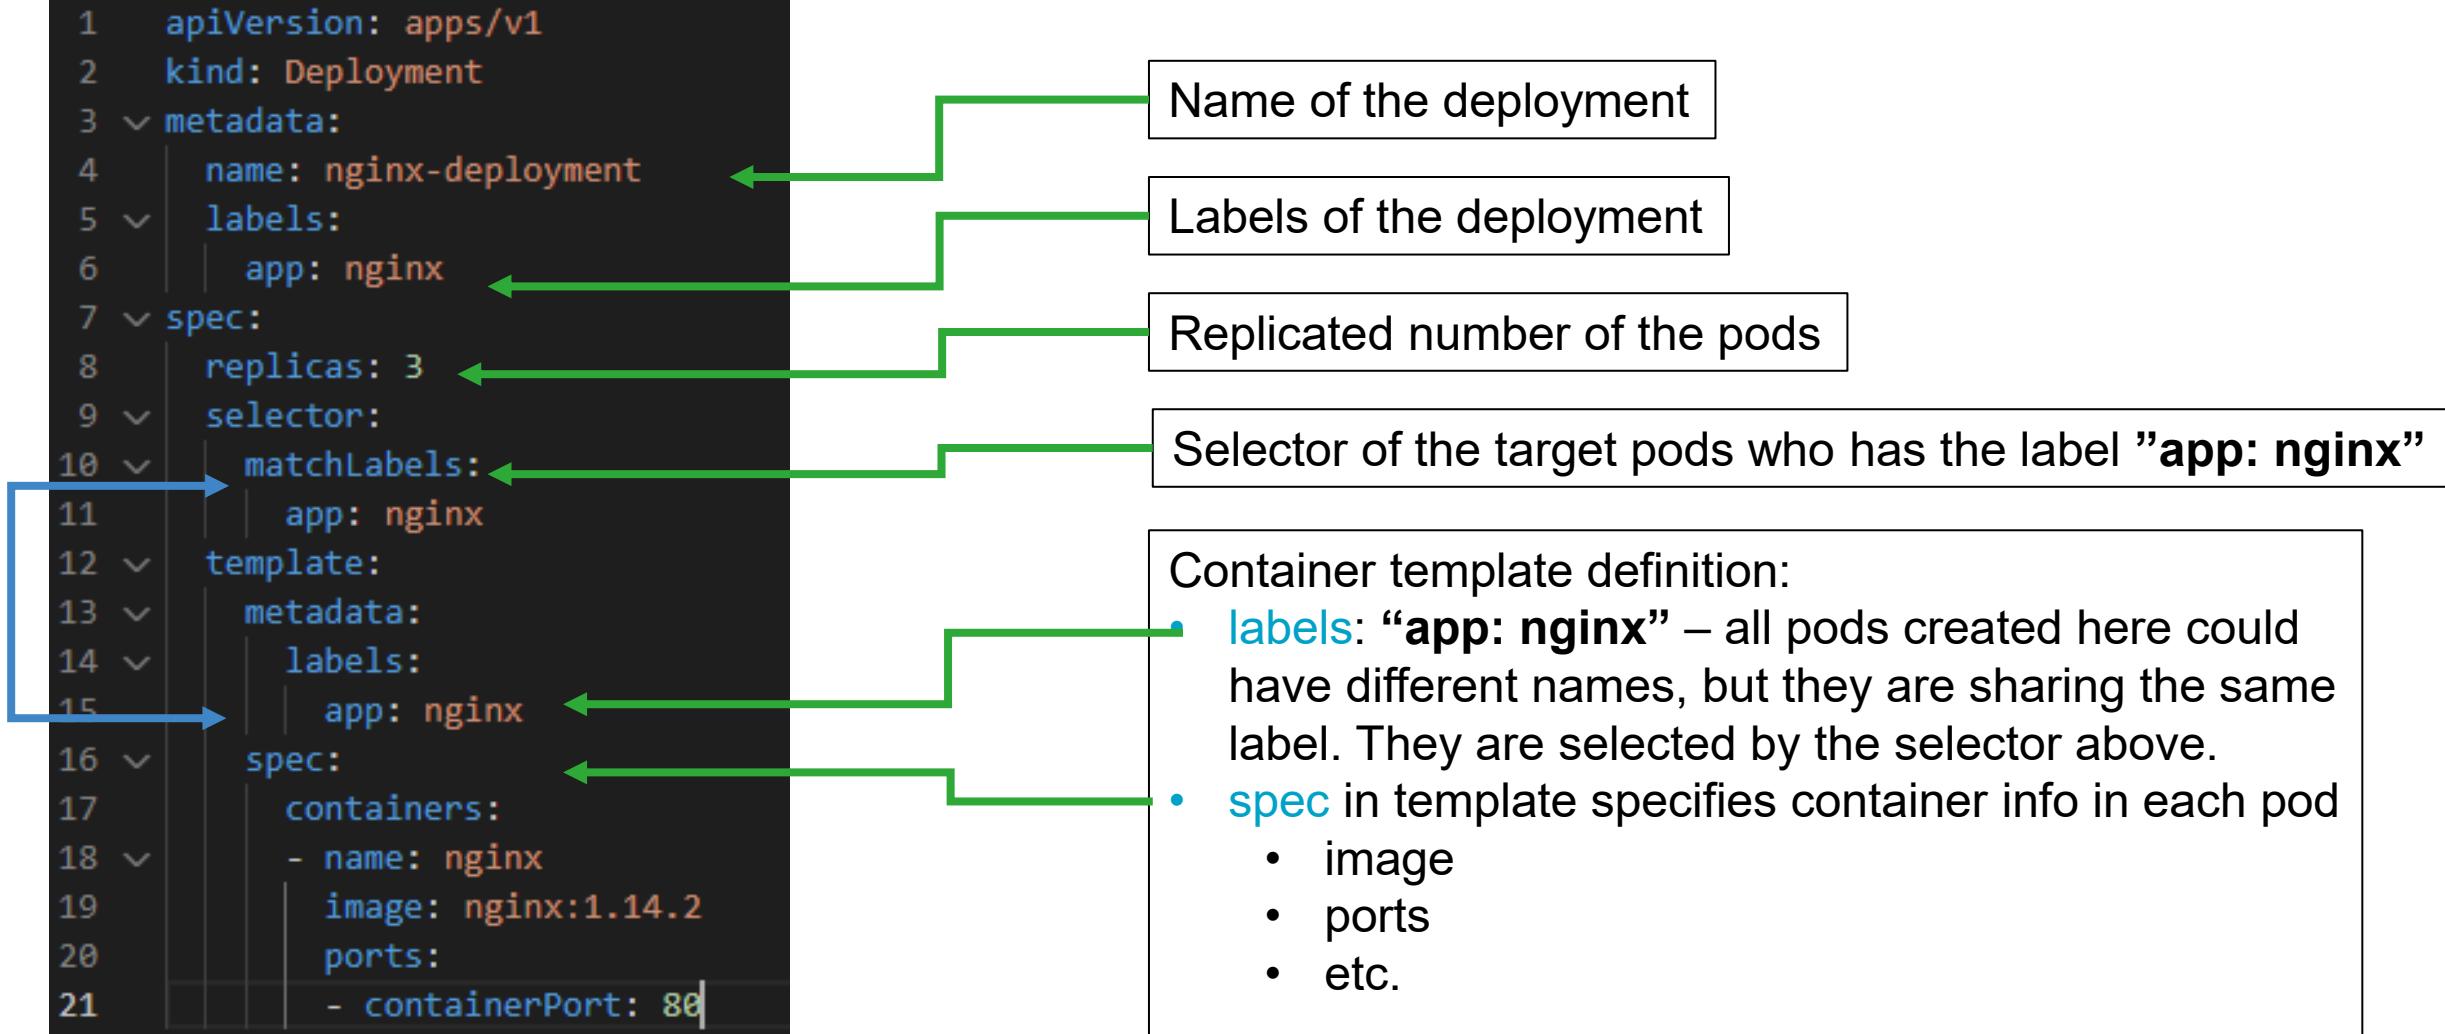
- Blueprint for pods (including replica number of the pod)
- You create deployments rather than pods (in practice no directly manipulating pods)
- ReplicaSet is to maintain a stable set of replica Pods (identical) running at any given time.
- Deployment can manage ReplicaSet (you may never need to manipulate ReplicaSet):
 - Use a Deployment directly and define your application in it.
 - Deployment can support for rolling updates allowing you to update a PodTemplateSpec and have the changes gradually rolled out in a controlled manner. It also supports rollbacks to previous versions.



Replica=3 for
the deployment

Deployment

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: nginx-deployment
5    labels:
6      app: nginx
7  spec:
8    replicas: 3
9    selector:
10   matchLabels:
11     app: nginx
12   template:
13     metadata:
14       labels:
15         app: nginx
16     spec:
17       containers:
18         - name: nginx
19           image: nginx:1.14.2
20           ports:
21             - containerPort: 80
```



Name of the deployment

Labels of the deployment

Replicated number of the pods

Selector of the target pods who has the label "app: nginx"

Container template definition:

- **labels**: "app: nginx" – all pods created here could have different names, but they are sharing the same label. They are selected by the selector above.
- **spec** in template specifies container info in each pod
 - image
 - ports
 - etc.

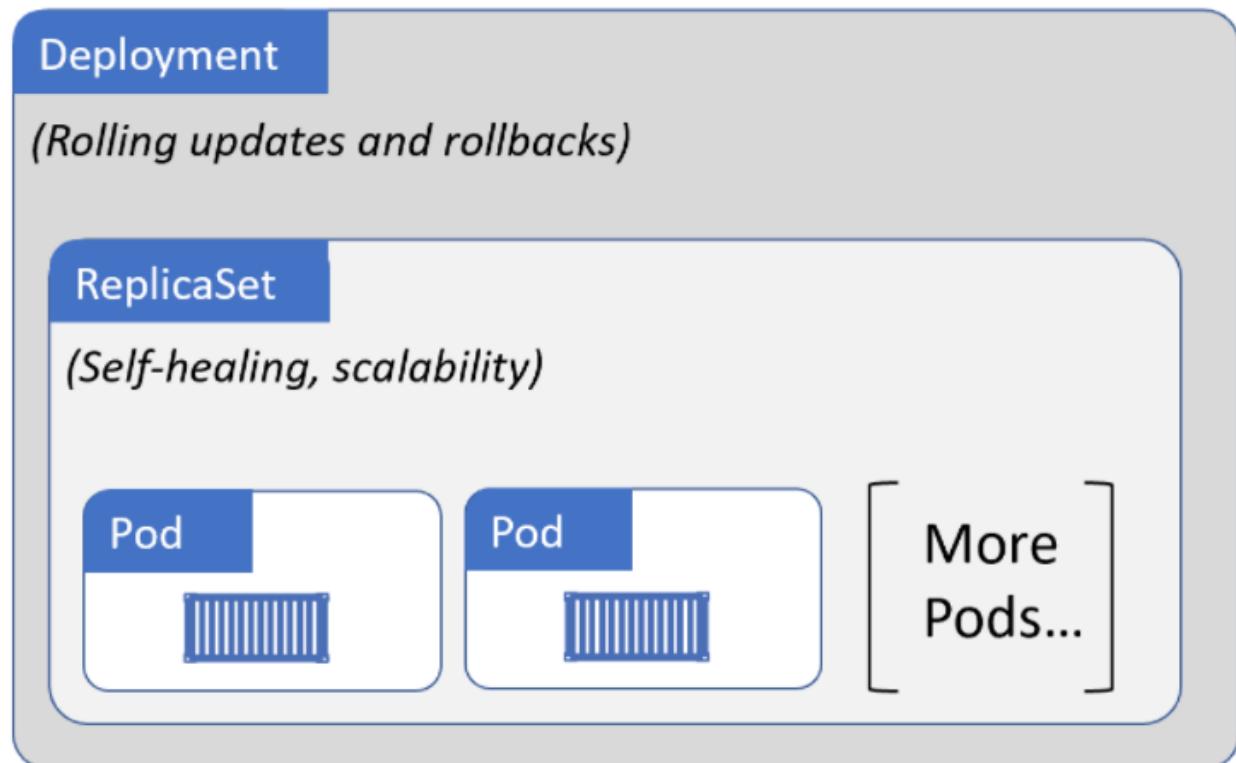
Deployments

Deployment controller

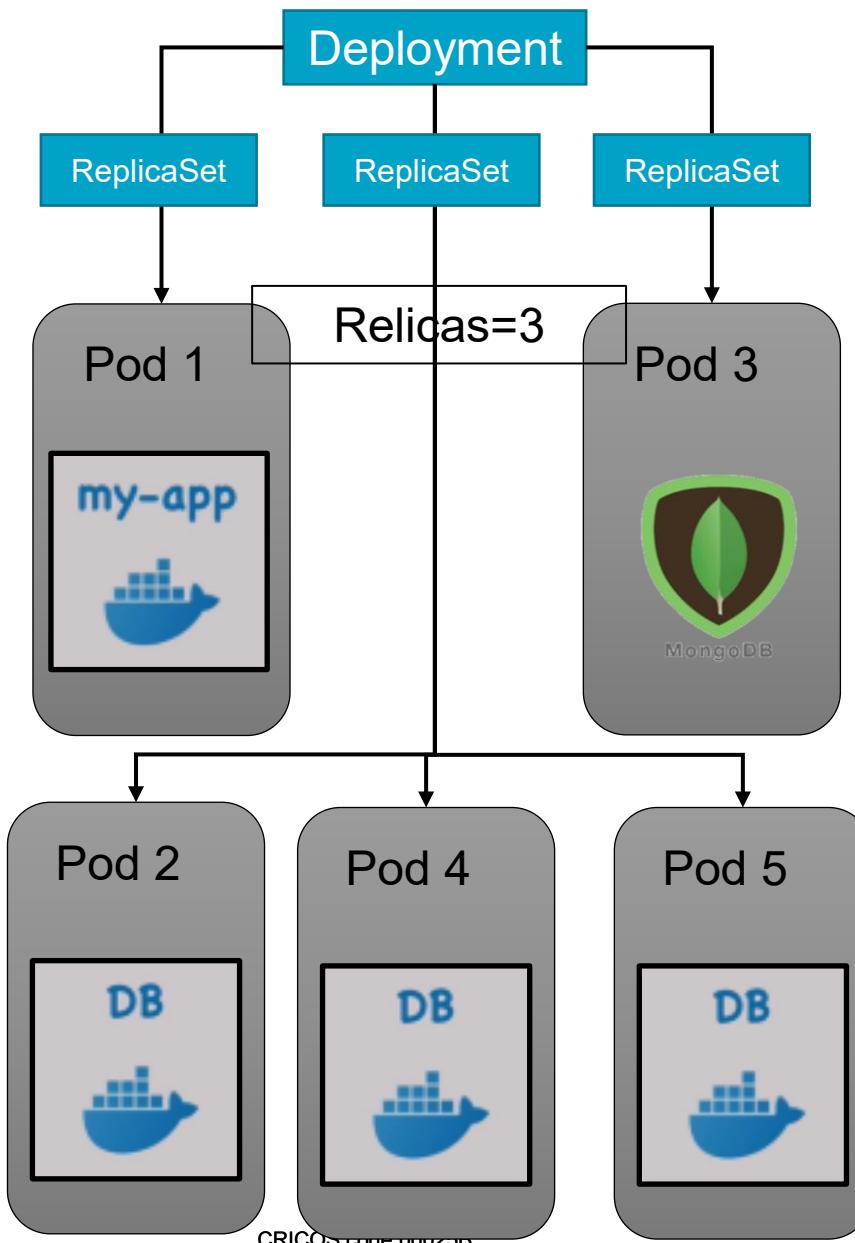
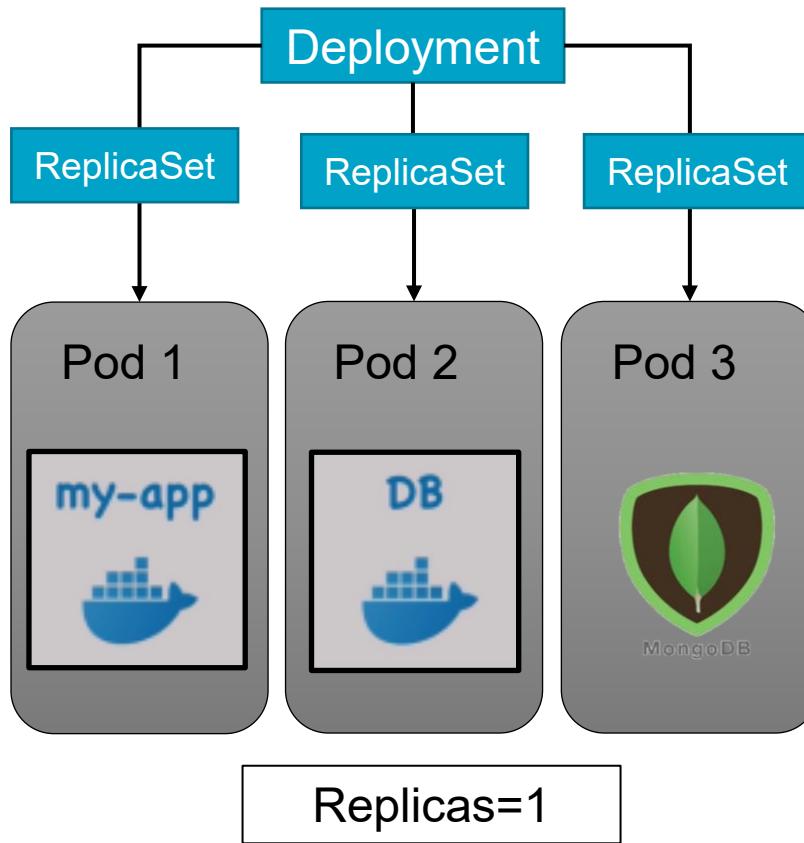
- Deployment: object type in K8s API
- Use **Replicaset**s provide self-healing and scaling

Replicaset controller

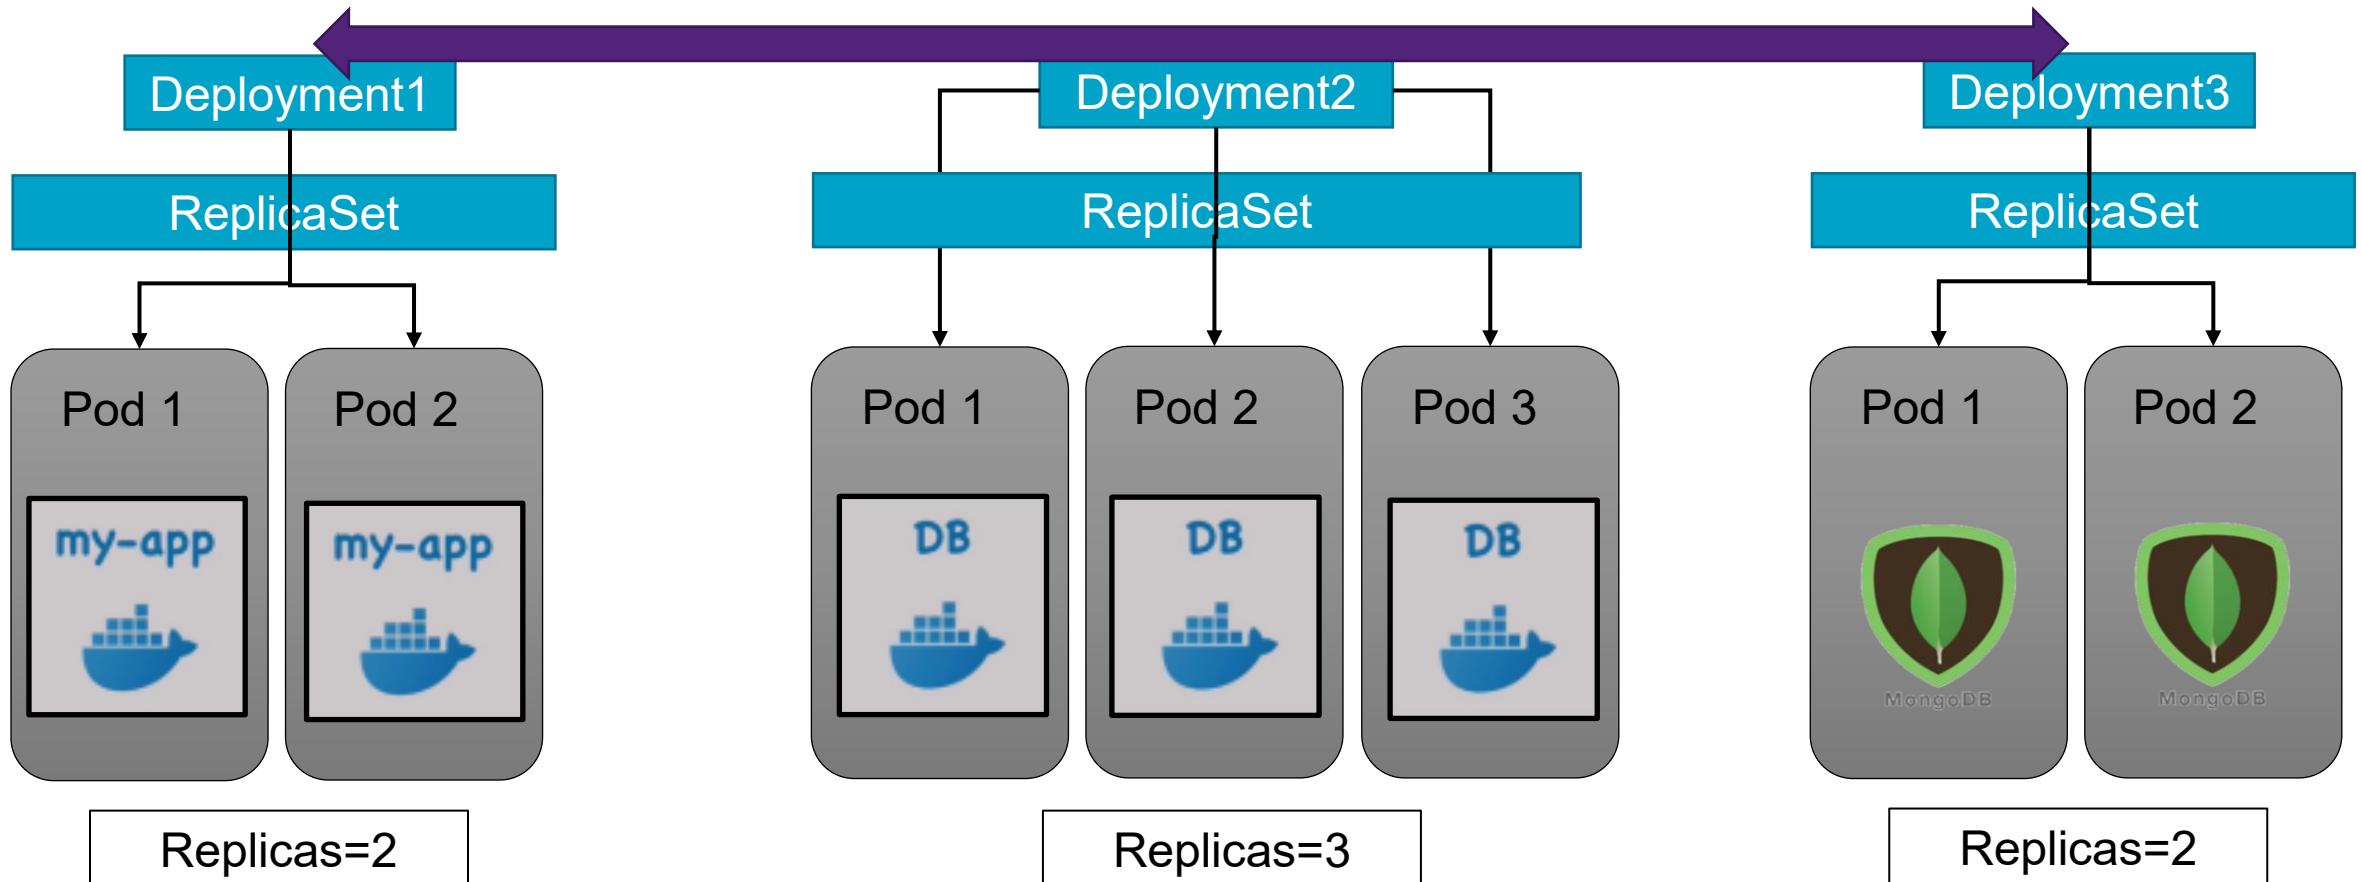
- Ensure the specified number of replicas of a service are always running



Deployment – Discussion 1



Deployment – Discussion 2

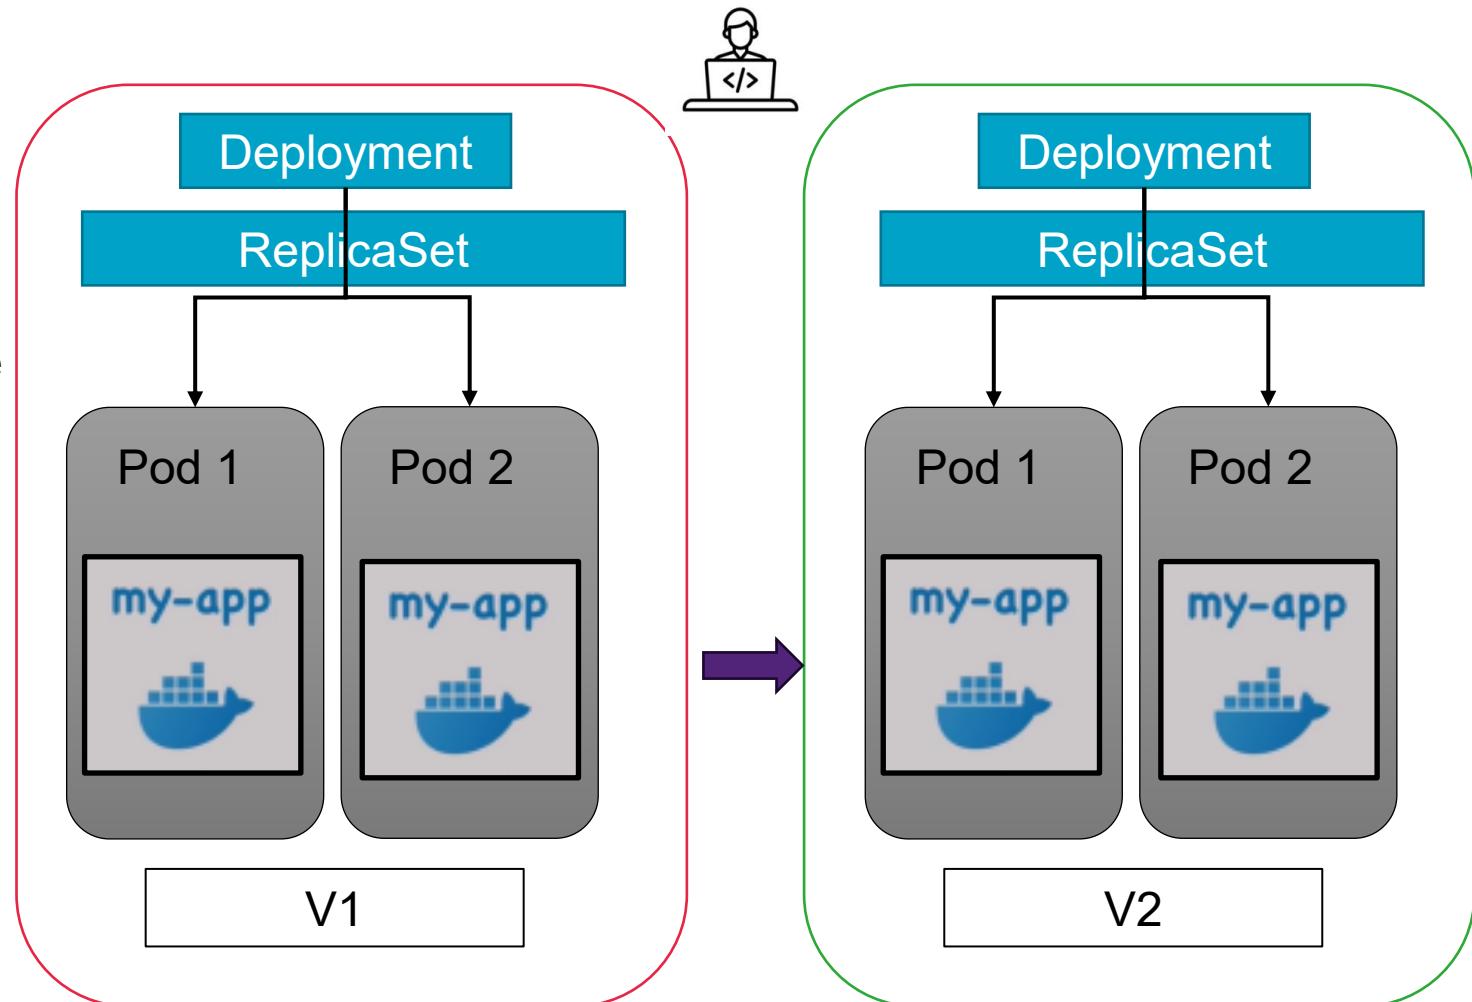


Note: The deployment is meant to be a microservice. For example: Deployment 1 could be a frontend microservice that has two replicas, while Deployment 2 could be a three-replicated DB microservice.

Deployment – Discussion 3

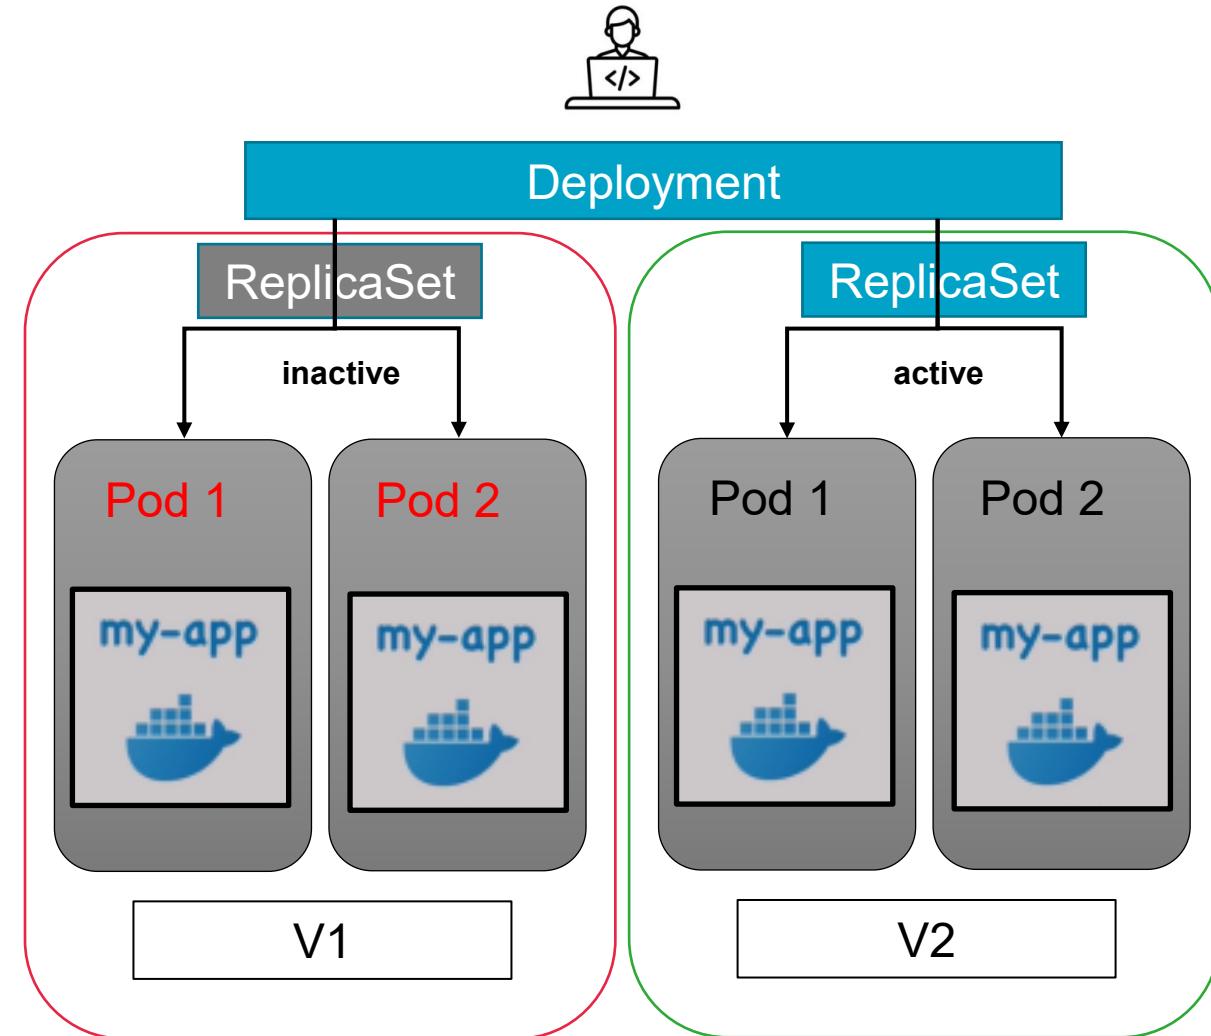
When deploying a newer version of a specific (micro)service: v2 \rightarrow v1 (in production), as a DevOps Engineer, you want to have the following features:

- 1. Rolling update**
the pods v1 to pods v2 in batches (the entire deployment won't go down);
- 2. Maintain the same number of replicas**
 $2 \times \text{pods_v1} \rightarrow 2 \times \text{pods_v2}$
- 3. Rollback can be conducted** when v2 is faulty (bugs, errors, etc.)
 $2 \times \text{pods_v2} \rightarrow 2 \times \text{pods_v1}$



Deployment – Discussion 4

- A Deployment can manage **multiple ReplicaSets**, but there is typically **only one active ReplicaSet**.
- The other **inactive ReplicaSets** are kept around with a reduced replica count (usually zero) **for the sake of rollback and revision history**.
- During a rolling update:
 - The new **ReplicaSet** (which matches the updated pod template in the Deployment) is scaled up.
 - The old **ReplicaSet** is scaled down (zero).



Demo: Deployments (1/4)

Deployment YAML file

- “**cc-demo-deployment**” is created, indicated by the **.metadata.name** field, which will be the basis for the **ReplicaSets** and **Pods**.
- creates a **ReplicaSet** that creates **THREE** replicated **Pods** (**.spec.replicas**).
- The **.spec.selector** defines how the created **ReplicaSet** finds which Pods to manage: selecting a label that is defined in the Pod template (**app: nginx**).
- The **template** field contains the following sub-fields:
 - **.metadata.labels** field: **app:nginx** using the **.**
 - **.template.spec** field: the Pods run one container (**nginx:1.14.2**) and the port of the container is **80**.

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: cc-demo-deployment
5    labels:
6      app: nginx
7  spec:
8    replicas: 3
9    selector:
10   matchLabels:
11     app: nginx
12   template:
13     metadata:
14       labels:
15         app: nginx
16     spec:
17       containers:
18         - name: nginx
19           image: nginx:1.14.2
20         ports:
21           - containerPort: 80
```

Demo: Deployments (2/4)

```
(base) sen@Sens-M1Pro ~ ➤ kubectl create -f deploy.yml
deployment.apps/nginx-deployment created
(base) sen@Sens-M1Pro ~ ➤ kubectl get deployment
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment  0/3        3            0          12s
(base) sen@Sens-M1Pro ~ ➤ kubectl rollout status deployment/nginx-deployment
deployment "nginx-deployment" successfully rolled out
(base) sen@Sens-M1Pro ~ ➤ kubectl get deployment
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment  3/3        3            3          52s
(base) sen@Sens-M1Pro ~ ➤ kubectl get rs
NAME           DESIRED   CURRENT   READY   AGE
cc-demo-replicaset      5         5         5       163m
nginx-deployment-6595874d85  3         3         3       61s
```

Demo: Deployments (3/4)

```
(base) sen@Sens-M1Pro ~ ➔ k get deployments
NAME          READY  UP-TO-DATE  AVAILABLE   AGE
nginx-deployment  5/5      5          5           115s
(base) sen@Sens-M1Pro ~ ➔ k set image deployment/nginx-deployment nginx=nginx:1.16.1
deployment.apps/nginx-deployment image updated
(base) sen@Sens-M1Pro ~ ➔ k rollout status deployment/nginx-deployment
Waiting for deployment "nginx-deployment" rollout to finish: 3 out of 5 new replicas have been updated...
Waiting for deployment "nginx-deployment" rollout to finish: 3 out of 5 new replicas have been updated...
Waiting for deployment "nginx-deployment" rollout to finish: 3 out of 5 new replicas have been updated...
Waiting for deployment "nginx-deployment" rollout to finish: 3 out of 5 new replicas have been updated...
Waiting for deployment "nginx-deployment" rollout to finish: 4 out of 5 new replicas have been updated...
Waiting for deployment "nginx-deployment" rollout to finish: 4 out of 5 new replicas have been updated...
Waiting for deployment "nginx-deployment" rollout to finish: 4 out of 5 new replicas have been updated...
Waiting for deployment "nginx-deployment" rollout to finish: 2 old replicas are pending termination...
Waiting for deployment "nginx-deployment" rollout to finish: 2 old replicas are pending termination...
Waiting for deployment "nginx-deployment" rollout to finish: 2 old replicas are pending termination...
Waiting for deployment "nginx-deployment" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "nginx-deployment" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "nginx-deployment" rollout to finish: 4 of 5 updated replicas are available...
deployment "nginx-deployment" successfully rolled out
(base) sen@Sens-M1Pro ~ ➔ k get deployments
NAME          READY  UP-TO-DATE  AVAILABLE   AGE
nginx-deployment  5/5      5          5           3m49s
(base) sen@Sens-M1Pro ~ ➔ k get rs
NAME          DESIRED  CURRENT  READY   AGE
cc-demo-replicaset     5         5        5  3h14m
nginx-deployment-6595874d85  0         0        0  3m53s
nginx-deployment-66b957f9d    5         5        5   53s
```

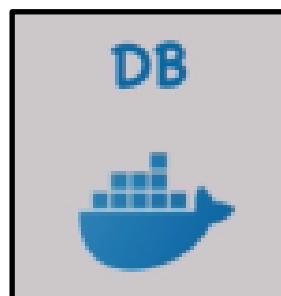
Demo: Deployments (4/4)

```
(base) sen@Sens-M1Pro ~ ➜ k describe deployment/nginx-deployment
Name: nginx-deployment
Namespace: default
CreationTimestamp: Mon, 28 Aug 2023 23:22:50 +1000
Labels: app=nginx
Annotations: deployment.kubernetes.io/revision: 2
Selector: app=nginx
Replicas: 5 desired | 5 updated | 5 total | 5 available | 0 unavailable
StrategyType: RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels: app=nginx
  Containers:
    nginx:
      Image: nginx:1.16.1
      Port: 80/TCP
      Host Port: 0/TCP
```

Run “`kubectl rollout history deployment/nginx-deployment`” to check out the revisions of the Deployment
Run “`kubectl rollout undo deployment/nginx-deployment --to-revision=2`” to rollback to a specific version

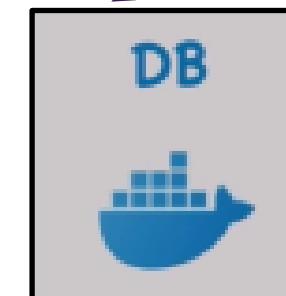
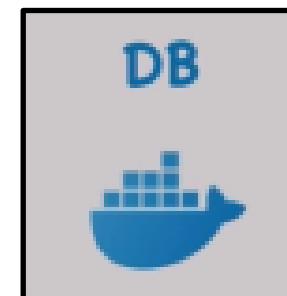
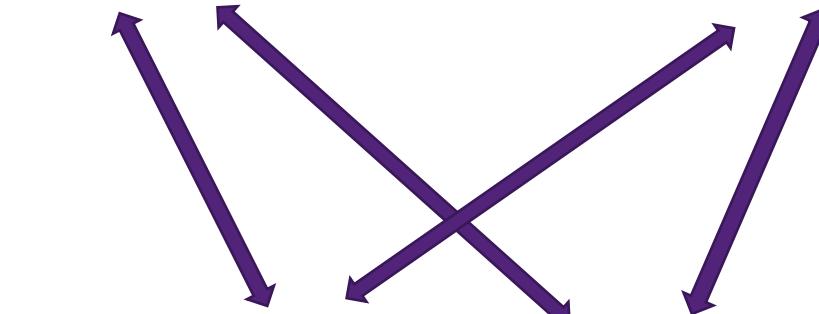
Running Applications on Pods

Pod: 10.0.0.1



Pod: 10.0.10.1

Pod 1: 10.0.0.1 Pod 2: 10.0.0.4 Pod 3: 10.0.0.2



Pod 4: 10.0.10.1 Pod 5: 10.0.10.3

Service and Ingress

Motivation

- There are no communication path
- Pods are unreliable
- New Pod gets a new address

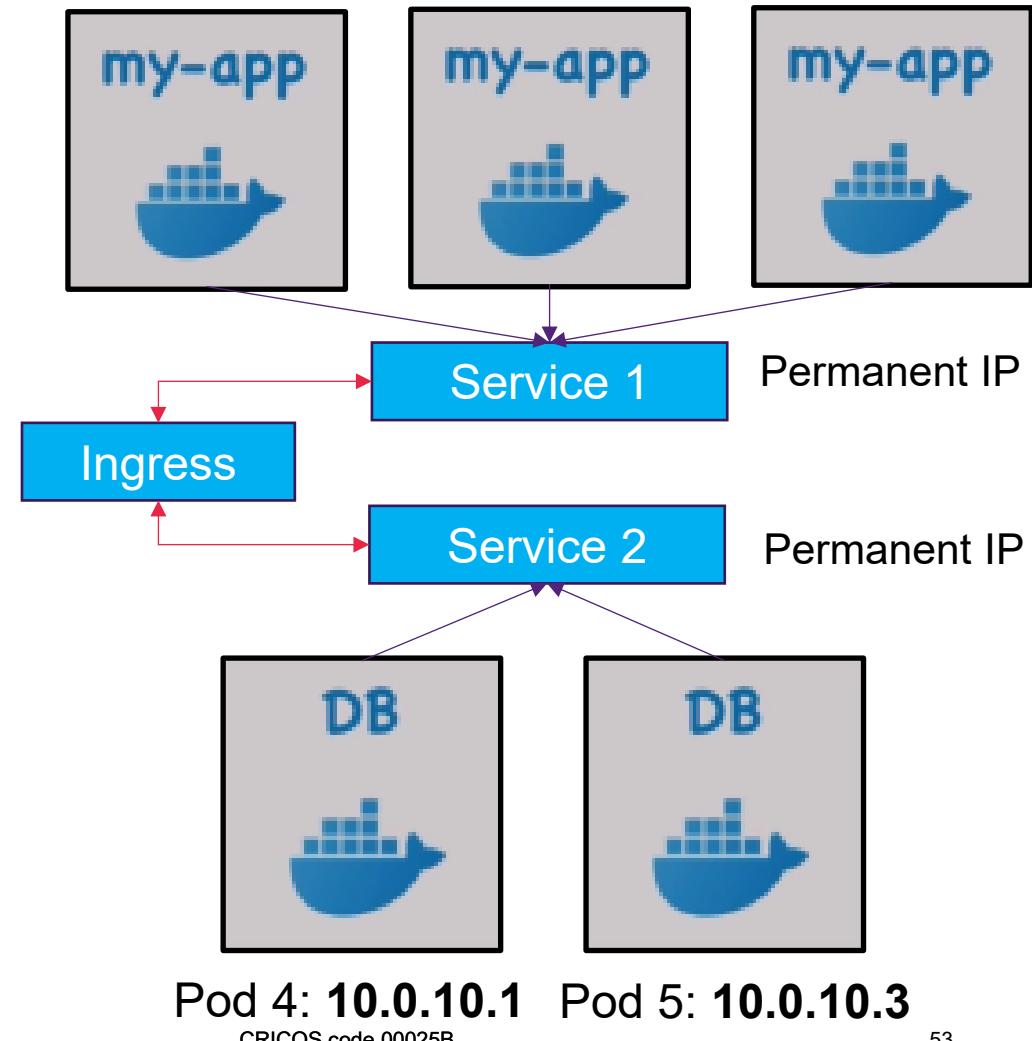
Services

- Abstraction way to expose an app running on a set of pods as a networking services
- Provide reliable networking for a set of Pods
- Load-balancing requests across abstracted pods

Ingress (DNS)

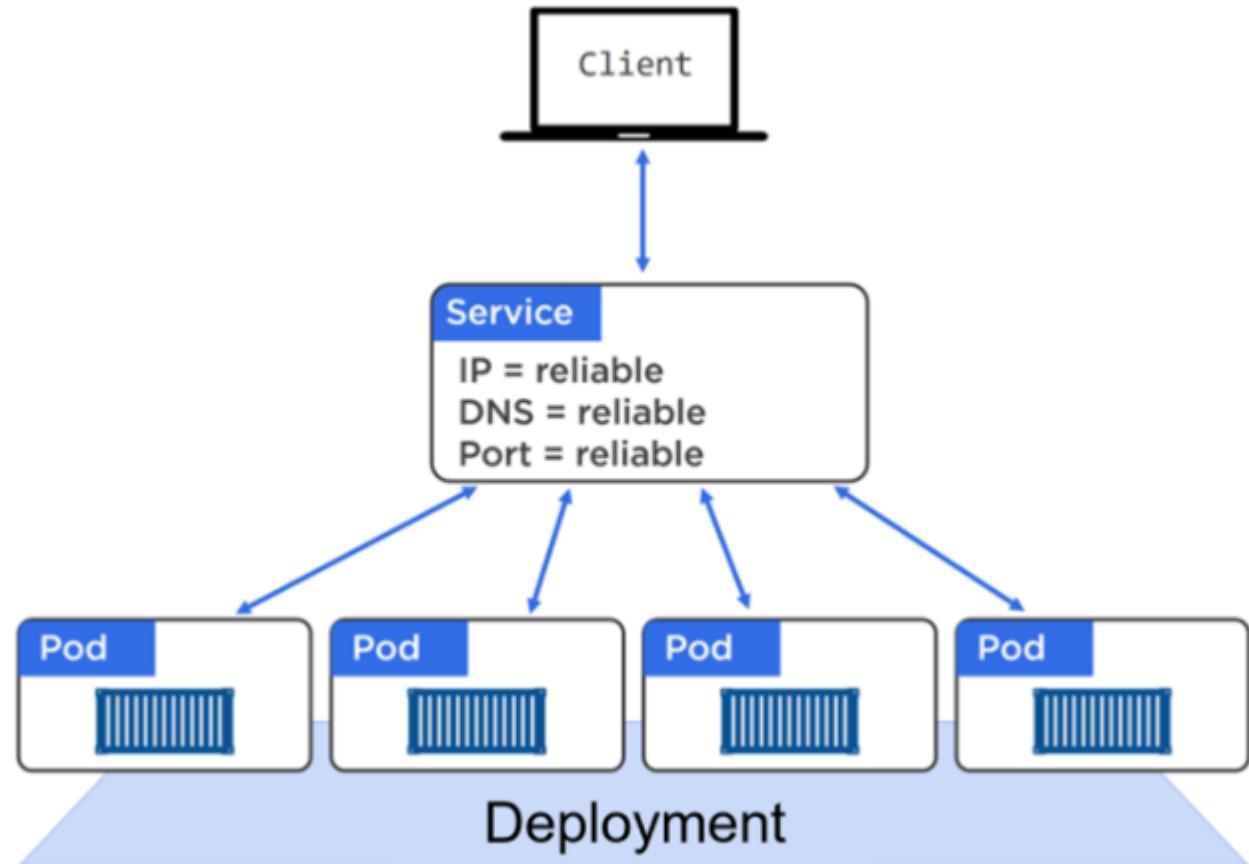
- manages external access to the services in a cluster, typically HTTP.

Pod 1: **10.0.0.1** Pod 2: **10.0.0.4** Pod 3: **10.0.0.2**



Service

- REST object in the API that you define in a manifest and POST to the API server
- Every Service gets its own stable IP address, its own stable DNS name and its own stable port.
- Leverage labels to dynamically select the Pods they will send traffic to.



Service

Create services by exposing ports

```
1  kubectl expose deployment cc-demo-deployment \
2      --name=cc-demo-service \
3      --target-port=28017 \
4      --type=NodePort
```

- Specified -> expose a deployment object
- Name it as ‘cc-demo-service’
- The port to expose (28017 for Mongodb)

- **ClusterIP** (default): expose the port only inside the cluster
- **NodePort**: expose the target port on every node to the outside
- **LoadBalancer**: only useful when combined with cloud provider’s load balancer

Service

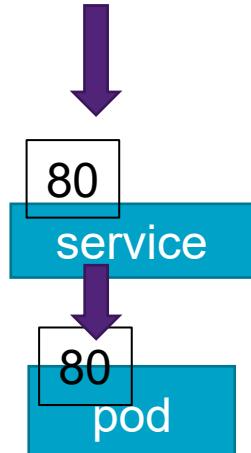
Create services with YAML file (kubectl apply -f service.yaml)

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: nginx
5    labels:
6      app.kubernetes.io/name: proxy
7  spec:
8    containers:
9      - name: nginx
10     image: nginx:stable
11     ports:
12       - containerPort: 80
13         name: http-web-svc
```

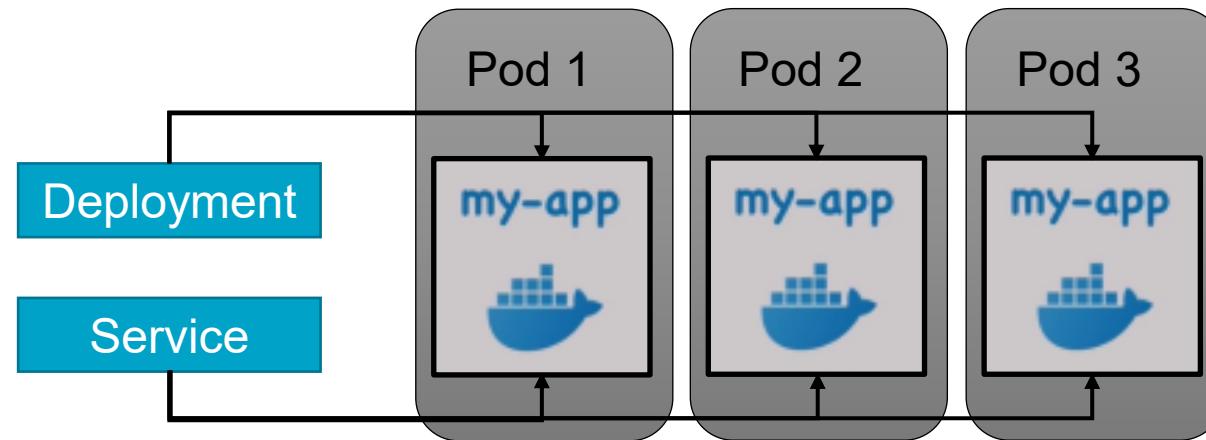
Pod.yaml

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: nginx-service
5  spec:
6    selector:
7      app.kubernetes.io/name: proxy
8    ports:
9      - name: name-of-service-port
10     protocol: TCP
11     port: 80
12     targetPort: http-web-svc
```

Service.yaml



Deployment vs Service



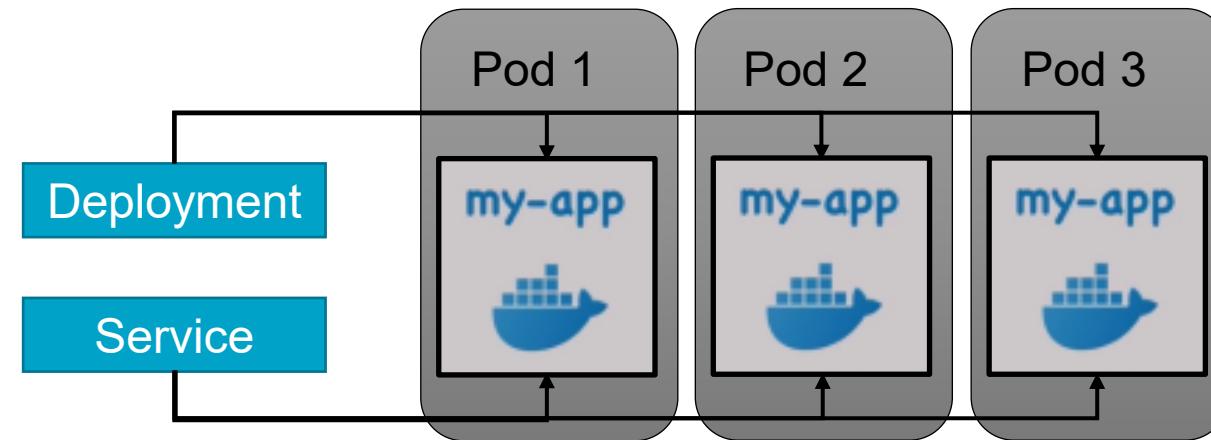
Both **Deployment** and **Service** are abstraction layers of Pods.

What are the differences between them?

Deployment vs Service

	Design Role	Functionality	Usage
Deployment	Pod abstraction Scalable (up or down) Blueprint (template)	<ul style="list-style-type: none"> Manages the desired state for your Pods and ReplicaSets. Allows you to perform rolling updates and rollbacks to your application. 	<ul style="list-style-type: none"> Ensures a certain number of instances of your application are running; to update your application, or to rollback an update. Deployment does NOT deal directly with networking or exposing your application.
Service	Pod abstraction Load-balancer (internal & external)	<ul style="list-style-type: none"> Provides a stable endpoint (IP address and port) to communicate with a set of Pods, typically exposed by a Deployment. Abstracts away the ephemeral nature of Pods, which might be created and destroyed frequently. 	<ul style="list-style-type: none"> Needs a persistent way to access your application, irrespective of the changes in the Pods (like creation, deletion, scaling). Service deals with networking and can expose your application internally (within the cluster) or externally (to the internet or outside world).

Deployment vs Service



Both **Deployment** and **Service** are abstraction layers of Pods.

What are the differences between them?

Simply speaking:

- a **Deployment** ensures your **desired application state** and **its updates**,
- a **Service** ensures your application remains **accessible** regardless of the dynamic changes in the underlying pods.
- They often work hand-in-hand: a **Deployment** maintains the application, and a **Service** provides stable access to it.

ConfigMap & Secret

There is information required when Pods are communicating:

- **Non-confidential**, such as the configuration of the Database;
- **Confidential data**, such as the password of database users

ConfigMap

- Used to store non-confidential data in **key-value** pairs.
- Consumed (by Pods) as **environment variables**, **command-line arguments**, or as **configuration files** in a volume.

Secret

- Similar to ConfigMap, but is used to store secret data (user name, password, tokens, keys, etc)
- Independently created (less risk of being exposed)
- base64 encoded (NOT encryption: binary -> ASCII).

```

1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: game-demo
5  data:
6    # property-like keys; each key maps to a simple value
7    player_initial_lives: "3"
8    ui_properties_file_name: "user-interface.properties"
9
10   # file-like keys
11   game.properties: |
12     enemy.types=aliens,monsters
13     player.maximum-lives=5
14   user-interface.properties: |
15     color.good=purple
16     color.bad=yellow
17     allow.textmode=true

```

Configmap.yaml

```

1  apiVersion: v1
2  kind: Secret
3  metadata:
4    name: mysql-secret
5  type: Opaque
6  data:
7    mysql-root-pwd: cm9vdHB3ZA==
8

```

Secret.yaml

Data Storage - Volumes

Data Storage for the app in the container (Pods):

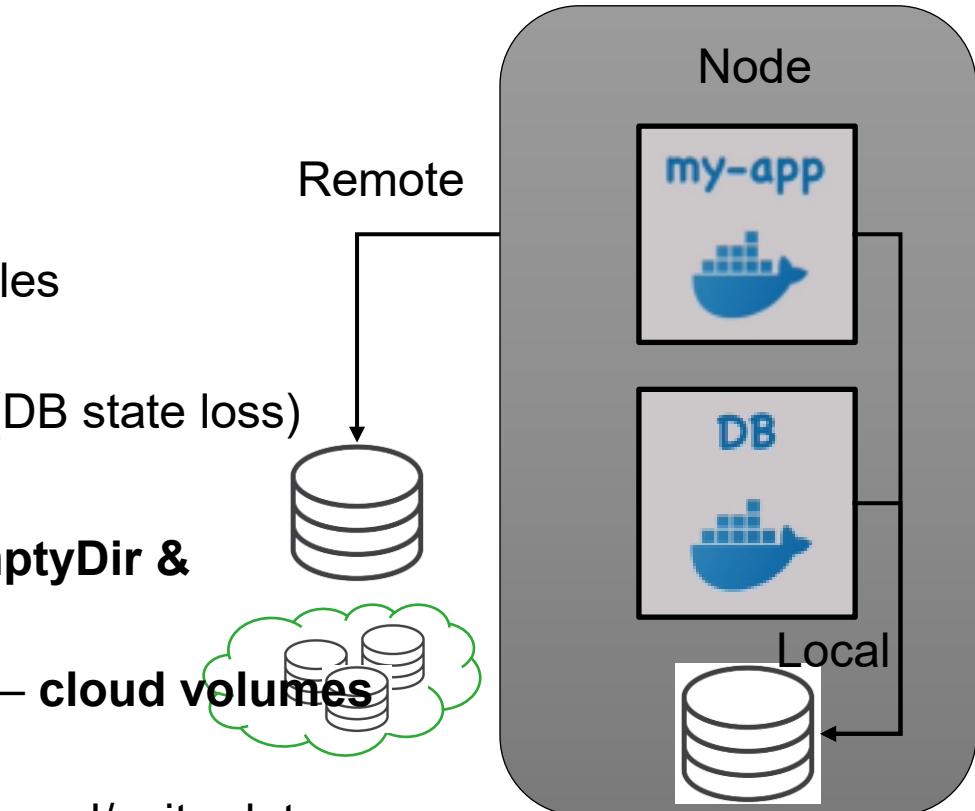
- In files (plain text, serialisation, encryption)
- In database files

Storing data directly in a container or a pod could have a loss of files (container/pod crashes) - **ephemeral**

Container/Pod is restarted with a clean state - no previous state (DB state loss)

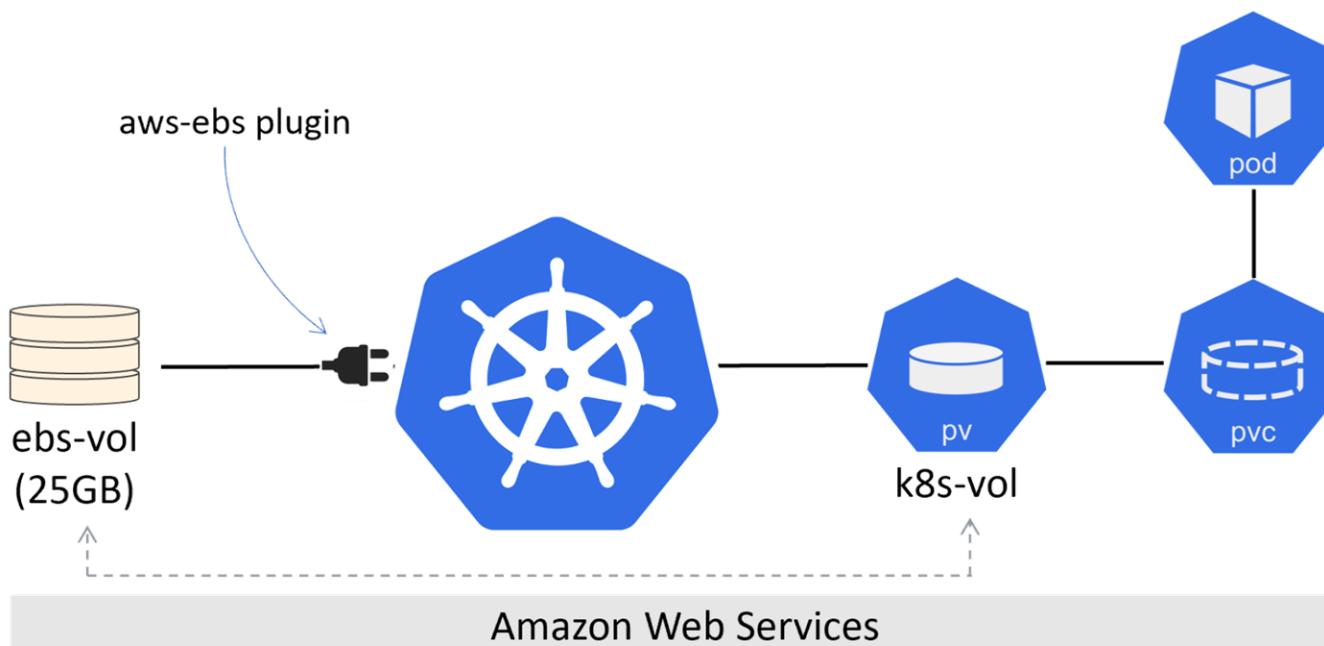
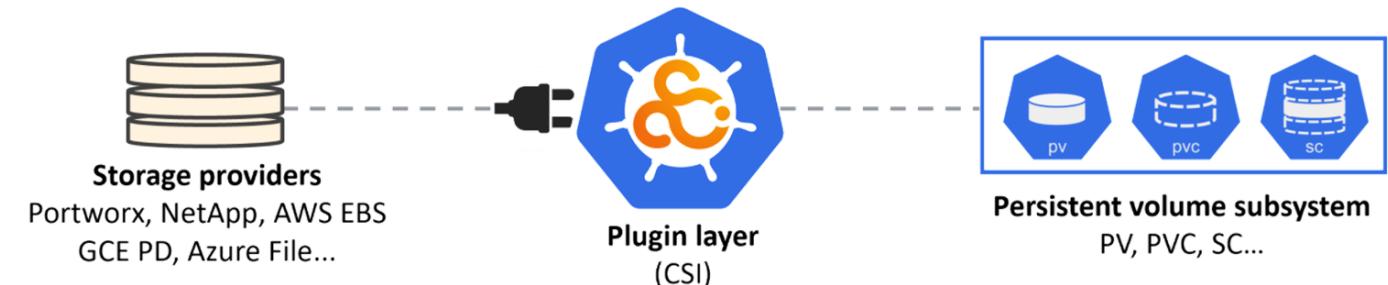
“Volumes” is to attach physical storage facility to persist the data

- Local storage (SSD or HDD) on the local machine (node) – **emptyDir & hostPath**
- Remote storage (on other Nodes in K8s or outside of the K8s) – **cloud volumes & NFS** (e.g. **awsElasticBlockStore** or **gcePersistentDisk**)
- K8s **DOES NOT** manage data persistence!! – mount a drive to read/write data
- K8s administrator should be responsible for data backup and security.



Kubernetes Storage

- Persistent Volume Subsystem:

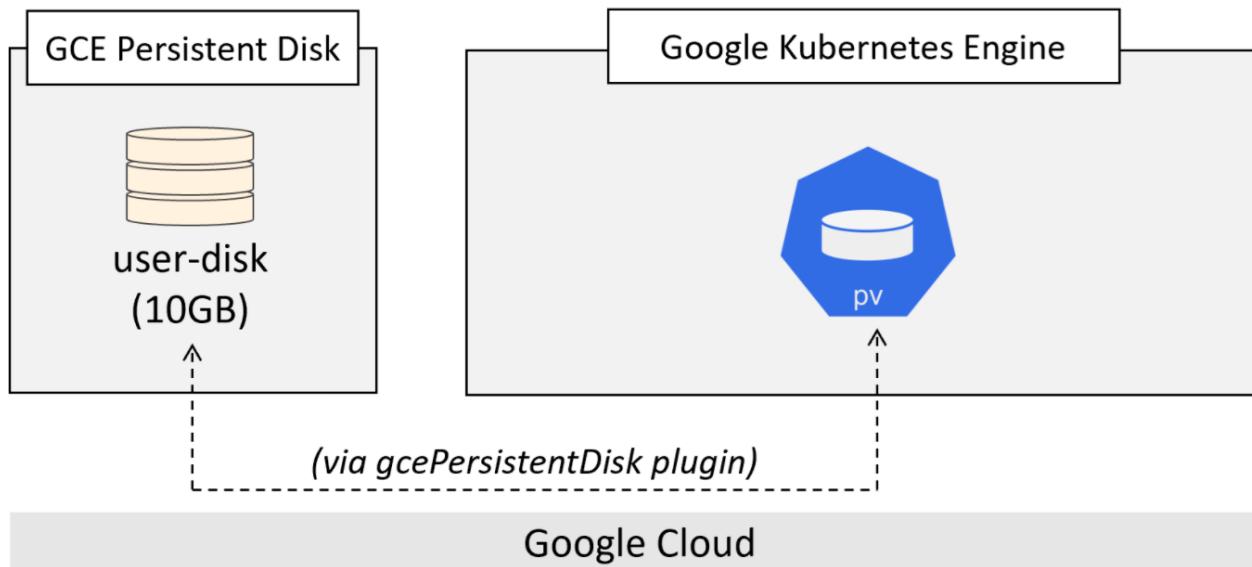
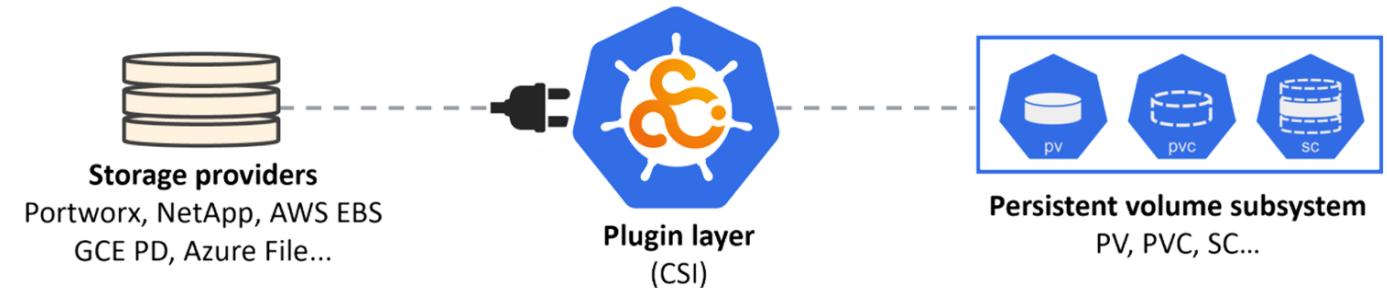


- Persistent Volume (PV):** how you map external storage onto the cluster; resource (**admin/ops**)
- Persistent Volume Claim (PVC):** a request for storage by a user; claim checks to the resource (**user/dev**)

<https://kubernetes.io/docs/concepts/storage/persistent-volumes/>

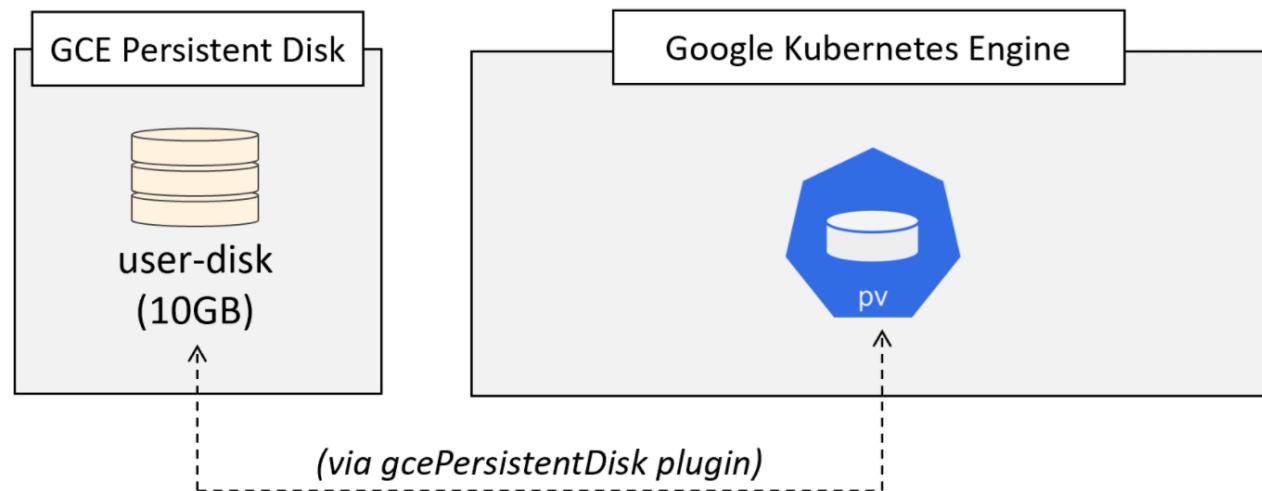
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes.html>

Kubernetes Storage



1. Create the PV.
2. Create the PVC.
3. Define the volume.
4. Mount it into a container.

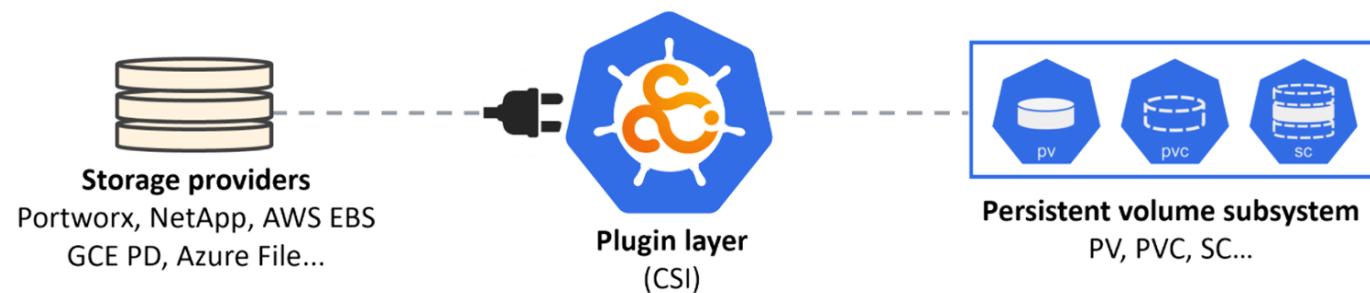
Kubernetes Storage



Google Cloud

- **ReadWriteOnce (RWO)**
- **ReadWriteMany (RWM)**
- **ReadOnlyMany (ROM)**

```
$ kubectl apply -f gke-pv.yml
```



```

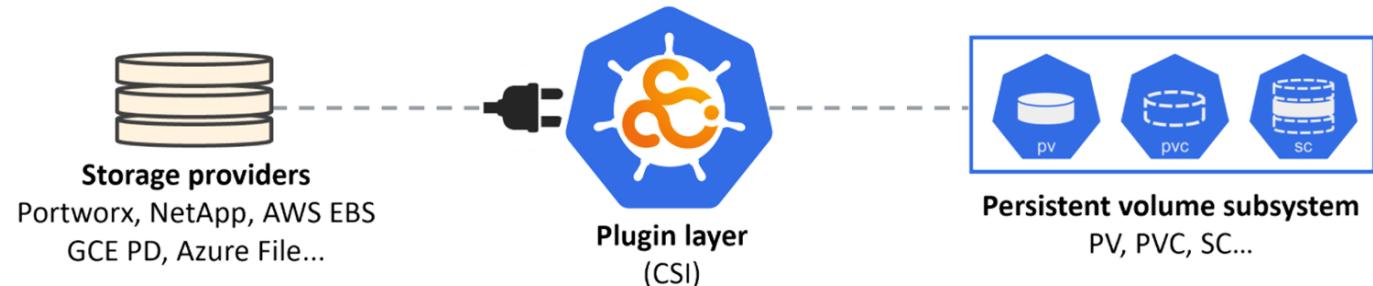
1  apiVersion: v1
2  kind: PersistentVolume
3  metadata:
4    name: pv1
5  spec:
6    accessModes:
7      - ReadWriteOnce
8    storageClassName: test
9    capacity:
10      storage: 10Gi
11    persistentVolumeReclaimPolicy: Retain
12    gcePersistentDisk:
13      pdName: uber-disk

```

gke-pv.yml

<https://kubernetes.io/docs/concepts/storage/persistent-volumes/>

Kubernetes Storage



Persistent Volume Claim: like a ticket that grants the Pod to the PV

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv1
spec:
  accessModes:
    - ReadWriteOnce
  storageClassName: test
  capacity:
    storage: 10Gi
  ...
  
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pvc1
spec:
  accessModes:
    - ReadWriteOnce
  storageClassName: test
  resources:
    requests:
      storage: 10Gi
  ...
  
```

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: volpod
5  spec:
6    volumes:
7      - name: data
8        persistentVolumeClaim:
9          claimName: pvc1
10   containers:
11     - name: ubuntu-ctr
12       image: ubuntu:latest
13       command:
14         - /bin/bash
15         - "-c"
16         - "sleep 60m"
17       volumeMounts:
18         - mountPath: /data
19           name: data
  
```

StatefulSet

Some applications have states, such as Database systems (MySQL, MongoDB, Redis, etc)

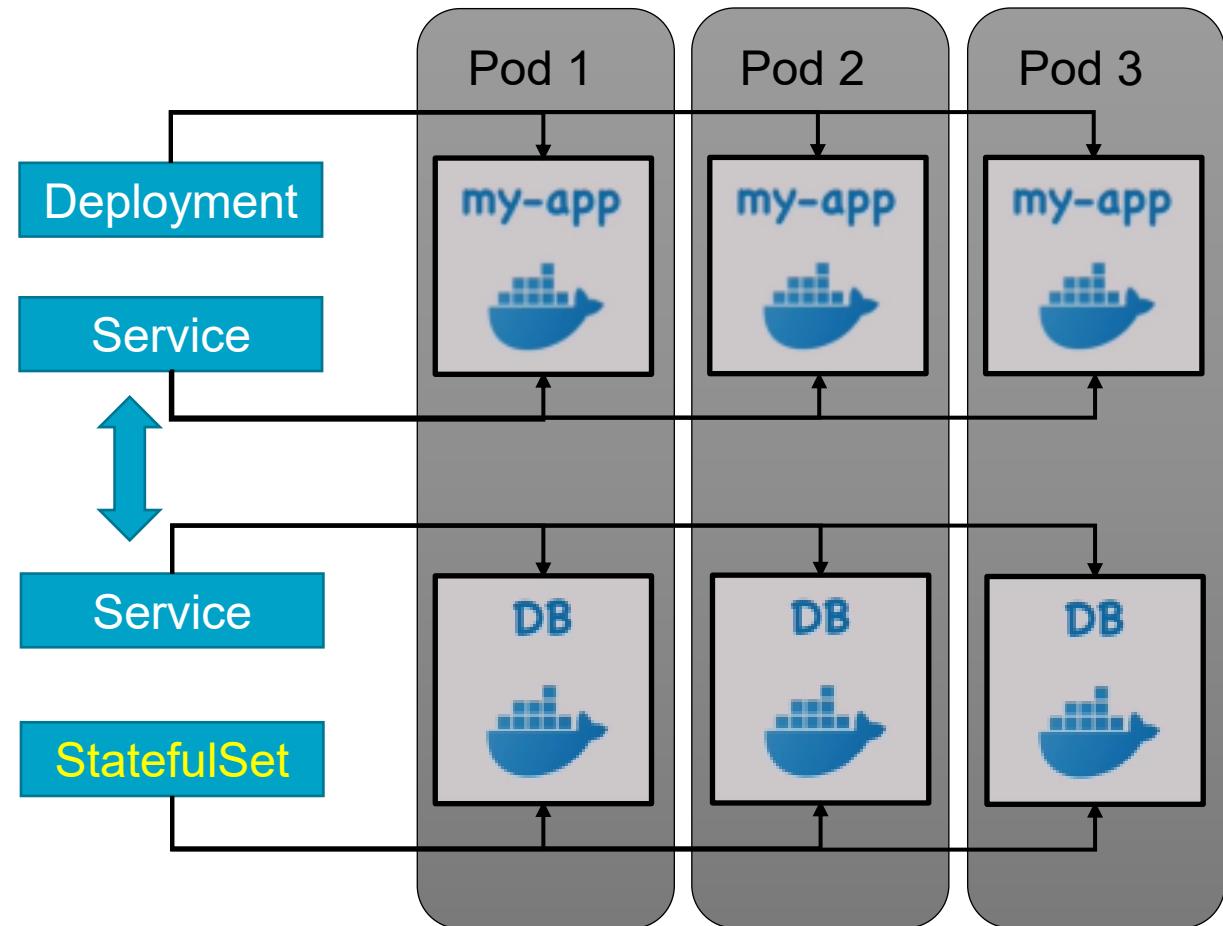
- States guarantee the data consistency.
- Stateful applications cannot be simply scaled up/down by Deployment

StatefulSet is to manage stateful applications.

- deploys and scales a set of Pods (DBs)
- provides guarantees about the ordering and uniqueness of these Pods.

Replicating DBs in K8s via StatefulSet is always difficult

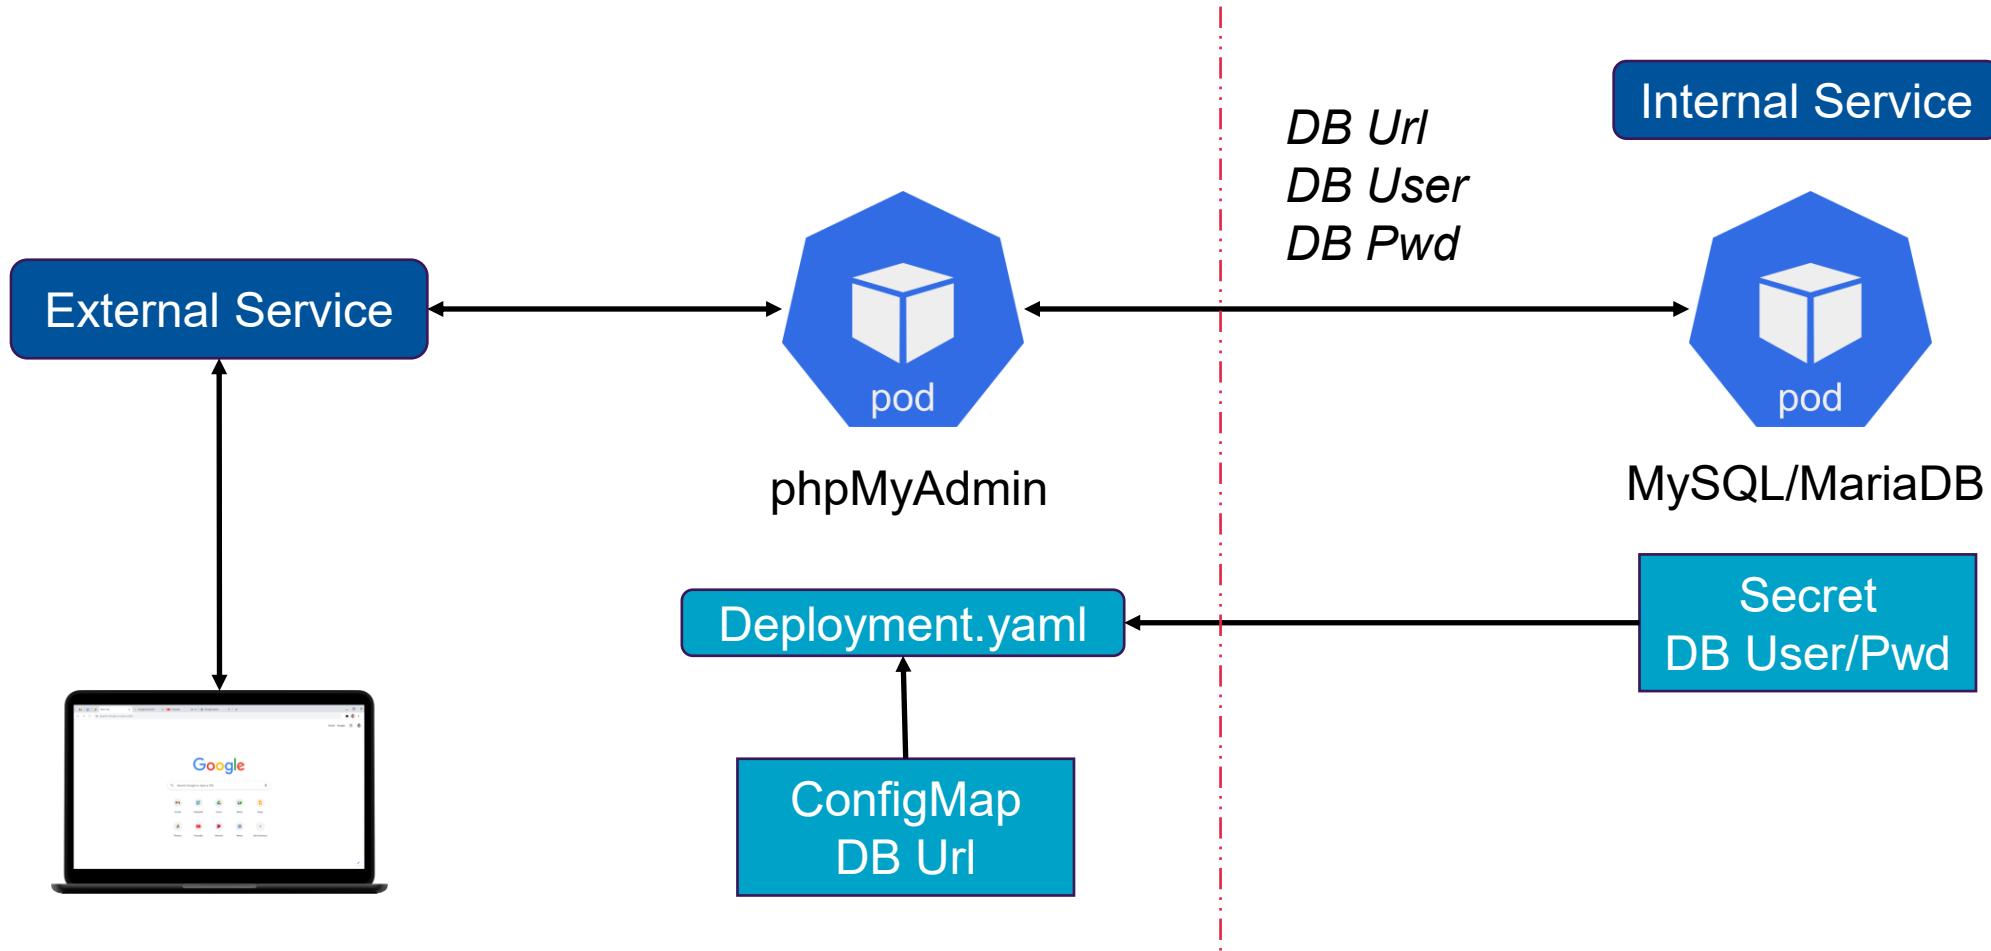
- DBs are often hosted outside of K8s cluster



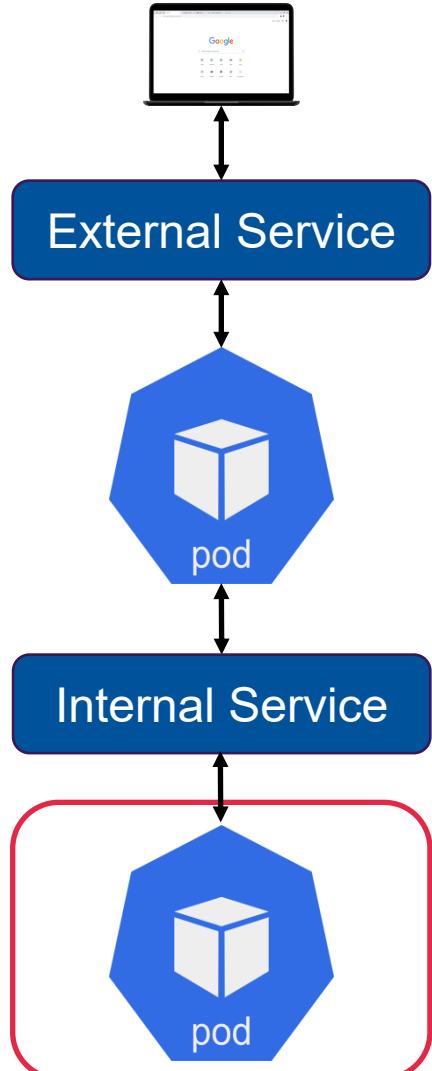
Outline

- Introduction to K8s
- K8s Architecture and its Components
 - Control Plane Node and Work Node
- Declarative Model in K8s
- Reconciliation Loop
- Minikube & Kubectl
- Demo of GKE - Cluster creation
- K8s Objects:
 - Pod & ReplicaSet & Deployment
 - & Service & Ingress & StatefulSet
 - & ConfigMap & Secrets & Volumes
- • A K8s deployment example - Mysql & phpMyAdmin

Deploying MySQL and phpMyAdmin on K8s - Framework



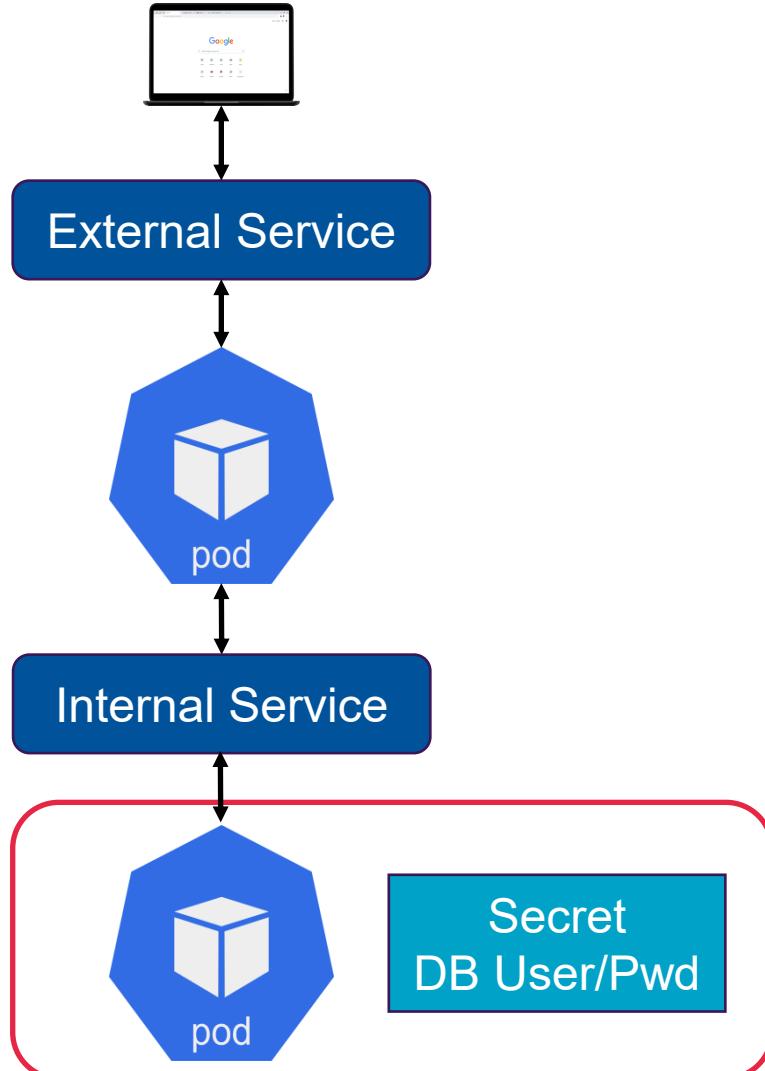
MySQL Pod Setup



mysql_deployment.yaml

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: mysql-deployment
5    labels:
6      app: mysql
7  spec:
8    replicas: 1
9    selector:
10   matchLabels:
11     app: mysql
12   template:
13     metadata:
14       labels:
15         app: mysql
16     spec:
17       containers:
18         - name: mysql-container
19           image: mysql
20           ports:
21             - containerPort: 3306
22           env:
23             - name: MYSQL_ROOT_PASSWORD
24               valueFrom:
25                 secretKeyRef:
26                   name: mysql-secret
27                     key: mysql-root-pwd
```

MySQL Secret



```
uqteaching@vm-global:~/lect6_demos$ echo -n 'mysql_root_pwd' | base64  
bXlzcWxfcm9vdF9wd2Q=
```

```
1  apiVersion: v1  
2  kind: Secret  
3  metadata:  
4    name: mysql-secret  
5  type: Opaque  
6  data:  
7    mysql-root-pwd: bXlzcWxfcm9vdF9wYXNzd29yZA==
```

mysql_secret.yaml

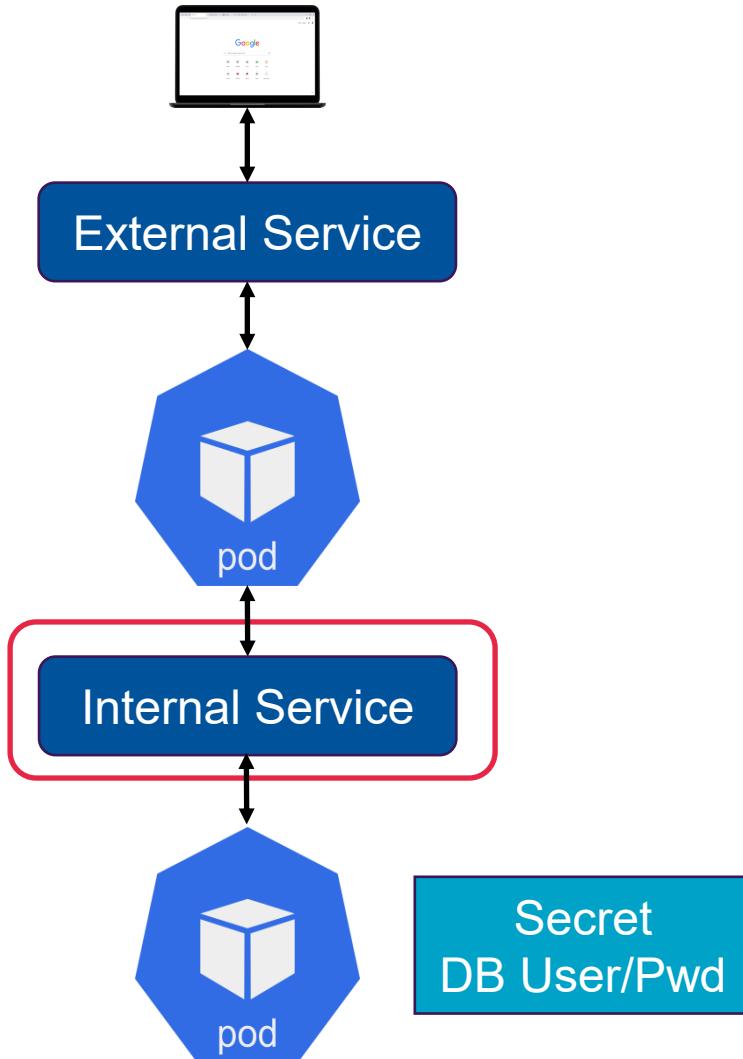
Note: Secret must be created before the Deployment

Create mysql secret with [kubectl apply](#)



```
uqteaching@vm-global:~/lect6_demos/mysql$ kubectl apply -f mysql_secret.yaml  
secret/mysql-secret created  
uqteaching@vm-global:~/lect6_demos/mysql$ kubectl get secret  
NAME        TYPE      DATA  AGE  
mysql-secret  Opaque    1     7s  
redis-secret  Opaque    2     5h7m
```

Internal Service



mysql_service.yaml

```

1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: mysql-service
5  spec:
6    selector:
7      app: mysql
8    ports:
9      - protocol: TCP
10     port: 3306
11     targetPort: 3306

```

mysql.yaml

```

1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: mysql-deployment
5    labels:
6      app: mysql
7  spec:
8    replicas: 1
9    selector:
10   matchLabels:
11     app: mysql
12   template:
13     metadata:
14       labels:
15         app: mysql
16     spec:
17       containers:
18         - name: mysql-container
19           image: mysql
20           ports:
21             - containerPort: 3306
22           env:
23             - name: MYSQL_ROOT_PASSWORD
24               valueFrom:
25                 secretKeyRef:
26                   name: mysql-secret
27                   key: mysql-root-pwd
28
29   ---|
30   apiVersion: v1
31   kind: Service
32   metadata:
33     name: mysql-service
34   spec:
35     selector:
36       app: mysql
37     ports:
38       - protocol: TCP
39         port: 3306
39         targetPort: 3306

```

MySQL Pod/Service/Deployment

```
uqteaching@vm-global:~/lect6_demos/mysql$ kubectl get all
NAME                                         READY   STATUS    RESTARTS   AGE
pod/mysql-deployment-56c94b9c87-jv657      1/1     Running   0          2m56s

NAME           TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
service/kubernetes   ClusterIP   10.96.0.1      <none>        443/TCP      6h56m
service/mysql-service   ClusterIP   10.97.169.71  <none>        3306/TCP      10s

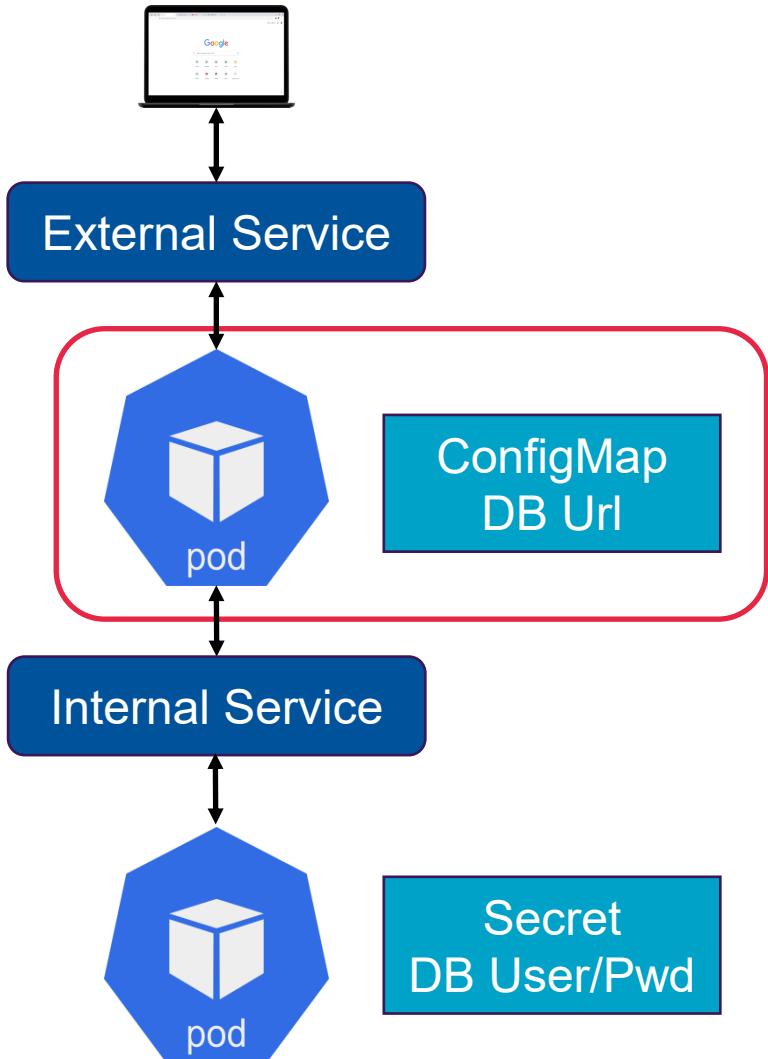
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/mysql-deployment      1/1       1           1          2m56s

NAME           DESIRED   CURRENT   READY   AGE
replicaset.apps/mysql-deployment-56c94b9c87  1         1         1        2m56s

uqteaching@vm-global:~/lect6_demos/mysql$ kubectl get pod
NAME                                         READY   STATUS    RESTARTS   AGE
mysql-deployment-56c94b9c87-jv657      1/1     Running   0          3m17s

uqteaching@vm-global:~/lect6_demos/mysql$
```

phpMyAdmin Pod



mysql_configmap.yaml

```

1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: mysql-configmap
5  data:
6    db_url: mysql-service
  
```

phpmyadmin_deployment.yaml

```

1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: phpmyadmin-deployment
5    labels:
6      app: phpmyadmin
7  spec:
8    replicas: 1
9    selector:
10   matchLabels:
11     app: phpmyadmin
12   template:
13     metadata:
14       labels:
15         app: phpmyadmin
16     spec:
17       containers:
18         - name: phpmyadmin-container
19           image: phpmyadmin
20           imagePullPolicy: IfNotPresent
21           ports:
22             - containerPort: 80
23               protocol: TCP
24             env:
25               - name: PMA_HOST
26                 valueFrom:
27                   configMapKeyRef:
28                     name: mysql-configmap
29                     key: db_url
30               - name: PMA_PORTS
31                 value: "3306"
32               - name: MYSQL_ROOT_PASSWORD
33                 valueFrom:
34                   secretKeyRef:
35                     name: mysql-secret
36                     key: mysql-root-pwd
  
```

phpMyAdmin Pod/Service/Deployment

```
uqteaching@vm-global:~/lect6_demos/mysql$ kubectl get all
NAME                                         READY   STATUS    RESTARTS   AGE
pod/mysql-deployment-56c94b9c87-jv657        1/1     Running   0          35m
pod/phpmyadmin-deployment-6f8d4f6854-cgfvz   1/1     Running   0          103s

NAME           TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
service/kubernetes   ClusterIP   10.96.0.1       <none>        443/TCP     7h28m
service/mysql-service   ClusterIP   10.97.169.71    <none>        3306/TCP     32m
service/phpmyadmin-service   LoadBalancer  10.109.234.164 <pending>     8081:30001/TCP 4s

NAME           READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/mysql-deployment   1/1      1           1           35m
deployment.apps/phpmyadmin-deployment   1/1      1           1           103s

NAME           DESIRED   CURRENT   READY   AGE
replicaset.apps/mysql-deployment-56c94b9c87  1         1         1      35m
replicaset.apps/phpmyadmin-deployment-6f8d4f6854  1         1         1      103s
```

Proxy for Minikube Running on VM

```
uqteaching@vm-global:~/lect6_demos/mysql$ minikube service phpmyadmin-service
|-----|-----|-----|-----|
| NAMESPACE |     NAME      | TARGET PORT |          URL       |
|-----|-----|-----|-----|
| default   | phpmyadmin-service |        8081 | http://192.168.49.2:30001 |
|-----|-----|-----|-----|
💡 Opening service default/phpmyadmin-service in default browser...
👉 http://192.168.49.2:30001
```

```
upstream app_server_30001 {
    server 192.168.49.2:30001;
}
server {
    listen 30001;
    location /proxy {
        proxy_pass http://app_server_30001/;
    }
}
```

```
uqteaching@vm-global:~/lect6_demos/mysql$ sudo vim /etc/nginx/conf.d/upstream.conf
uqteaching@vm-global:~/lect6_demos/mysql$ sudo nginx -t
nginx: the configuration file /etc/nginx/nginx.conf syntax is ok
nginx: configuration file /etc/nginx/nginx.conf test is successful
uqteaching@vm-global:~/lect6_demos/mysql$ sudo systemctl reload nginx
uqteaching@vm-global:~/lect6_demos/mysql$
```

Kubernetes on the Platforms

Major cloud providers offer managed Kubernetes services to simplify cluster management tasks for users. Here's a comparison of Kubernetes services across three major cloud providers: AWS, GCP, and Azure:

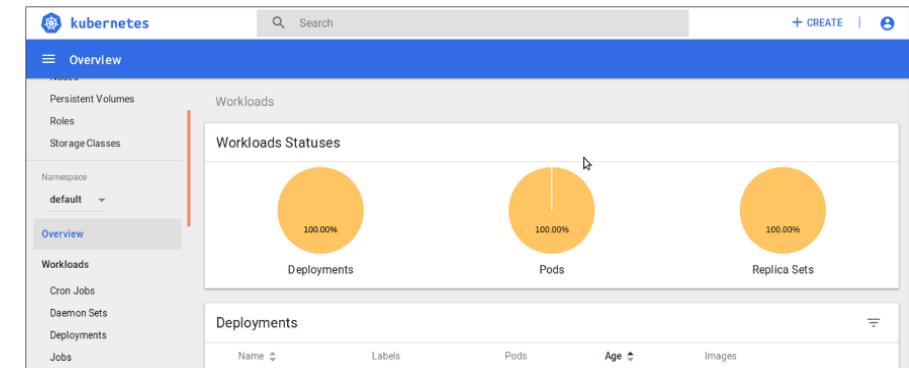
	Service Name	Integration with Native Services	Networking	Auto Scaling	Pricing
GCP	Google K8s Engine (GKE)	GKE has deep integrations with Google Cloud services like BigQuery , Cloud Spanner , etc.	Uses Google's global VPC for connectivity, providing strong network capabilities out of the box.	Yes	Charges for worker node usage ; control plane (master node) usage is FREE .
AWS	Amazon Elastic K8s Services (EKS)	EKS integrates with many AWS services like Elastic Load Balancer (ELB), Identity and Access Management (IAM), and CloudWatch.	Supports VPC Peering, allowing private communication between different VPCs.	Yes	Charges for control plane and worker node usage*. Charges for resource usage*.
Azure	Azure K8s Service (AKS)	AKS integrates with Azure services such as Azure Active Directory , Azure Monitor , and Azure Logic Apps .	Offers Azure Virtual Network, ensuring private communication between different services.	Yes	Does not charge for the Kubernetes management infrastructure; users pay only for the virtual machines and associated resources**.

All three cloud providers offer extensions for developers, robust support options, and Service Level Agreements (SLAs) to ensure high availability and security.

* Amazon EKS pricing model contains two major components: customers pay \$0.10 per hour for each configured cluster, and pay for the AWS resources (compute and storage) that are created within each cluster to run Kubernetes worker nodes.
**<https://azure.microsoft.com/en-au/pricing/details/kubernetes-service/#pricing>

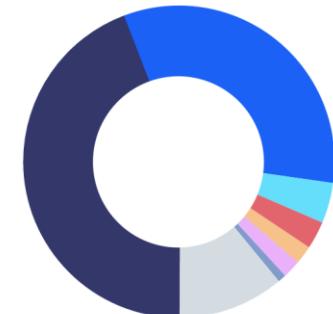
Kubernetes vs Docker Swarm

Features	Kubernetes	Docker Swarm
Installation & Cluster Configuration	Installation is complicated; but once setup, the cluster is very strong	Installation is very simple; but cluster is not very strong
GUI	GUI is the Kubernetes Dashboard	There is no GUI
Scalability	Highly scalable & scales fast	Highly scalable & scales 5x faster than Kubernetes
Auto-Scaling	Kubernetes can do auto-scaling	Docker Swarm cannot do auto-scaling
Load Balancing	Manual intervention needed for load balancing traffic between different containers in different Pods	Docker Swarm does auto load balancing of traffic between containers in the cluster
Rolling Updates & Rollbacks	Can deploy Rolling updates & does automatic Rollbacks	Can deploy Rolling updates, but not automatic Rollbacks
Data Volumes	Can share storage volumes only with other containers in same Pod	Can share storage volumes with any other container
Logging & Monitoring	In-built tools for logging & monitoring	3rd party tools like ELK should be used for logging & monitoring



`kubectl autoscale deployment php-apache --cpu-percent=50 --min=1 --max=10`

Which container orchestration platform do you primarily use?



Source: quarterly report on developer trends in the cloud by Digital Ocean

Review

- Introduction to K8s
- K8s Architecture and its Components
 - Control Plane Node and Work Node
- Declarative Model in K8s
- Reconciliation Loop
- Minikube & Kubectl
- K8s Objects:
 - Pod & ReplicaSet & Deployment
 - & Service & Ingress & StatefulSet
 - & ConfigMap & Secrets & Volumes

What's Next?

Getting more hands-on practice

- Play with Kubernetes classroom: <https://training.play-with-kubernetes.com/kubernetes-workshop/>
- Practice creating a K8s cluster on GCP and deploying a stateless application (practical session)

Getting prepared for certification

- <https://cloud.google.com/certification/cloud-architect>

Remember to stop GKE clusters when not in use...

```
uqyluo@cloudshell:~ (mythic-dynamo-300704)$ gcloud container clusters resize cluster-demo --zone=australia-southeast1-a --num-nodes=0
Pool [default-pool] for [cluster-demo] will be resized to 0.

Do you want to continue (Y/n)? y
```

What's Next?

A few more things to explore...

- Kubernetes Dashboard (<https://kubernetes.io/docs/tasks/access-application-cluster/web-ui-dashboard/>)
- ConfigMaps (<https://kubernetes.io/docs/concepts/configuration/configmap/>)
- Threat Modeling (<https://kubernetes.io/docs/tasks/administer-cluster/securing-a-cluster/>)

Tutorial & Practical for Week 6

Tutorial 5 (Week 6)

1. Please describe what Kubernetes (K8s) is.
2. What are the motivations of using Kubernetes?
3. Please describe the components and their functions in a Kubernetes cluster?
4. **Please generally summarise the workflow of Kubernetes**

Practical 5 (Week 6)

1. Use Docker commands to deploy a PHP-based web development environment with four containers.
2. Creating K8s cluster
3. Deploy a stateful application on K8s
4. Consultation for Programming Task 1

Next (Week 6) Topic:

Databases in Cloud Computing I (Dr Zhi Chen)