

Cloud Computing (INFS3208)

Lecture 8: Introduction to Spark Framework
-- An in-Memory Analytics Tool for Data Science

Lecturer: AsPr Sen Wang

School of Electrical Engineering and Computer Science

Faculty of Engineering, Architecture and Information Technology

The University of Queensland

PhD Recruitment (1/4)

Research focuses on state-of-the-art topics in AI and Machine Learning. Successful candidates will have the opportunity to contribute to one or more of the following areas:

- **Multimodal Learning:** Developing models that can understand and reason about the world from multiple data sources, such as text, images, and audio.
- **AI for Edge Devices:** Designing efficient and powerful AI algorithms (e.g., model compression, quantization, and neural architecture search) that can run on resource-constrained devices like smartphones and IoT sensors.
- **Reinforcement Learning (RL):** Creating intelligent agents that learn to make optimal decisions through interaction and feedback, with applications in robotics, game theory, and autonomous systems.
- **Large Language Models (LLMs) and their Applications:** Exploring the frontiers of LLMs, including topics like instruction-tuning, retrieval-augmented generation (RAG), efficient fine-tuning, and developing novel agent-based systems.

PhD Recruitment (2/4)

Essential:

- Domestic applicants and those onshore international students who have completed a program at UQ in Semester 2 2025.
- Applicants must be onshore at the time that offers are issued.
- An outstanding academic record, typically a First Class Honours degree or a Master's degree with a significant research component in Computer/Data Science, Engineering, Mathematics, or a related discipline.
- A strong mathematical foundation and excellent programming skills (Python is essential; experience with frameworks like PyTorch or TensorFlow is highly valued).
- Excellent written and verbal communication skills in English.
- Strong motivation to conduct high-impact research and publish at top-tier conferences.

Desirable:

- Previous research experience (e.g., an honours thesis, research-based project, or publications).
- Demonstrated experience in one of the research areas listed above.

PhD Recruitment (3/4)

This position is contingent upon the applicant successfully securing a research scholarship through the UQ Graduate School. The UQ Graduate School Scholarship (UQGSS) provides a competitive living stipend and a full tuition fee waiver.

Please take note of the key dates for the upcoming UQ scholarship round:

- **Applications Open:** Monday, 25 August 2025
- **Applications Close:** Monday, 22 September 2025 (All documents must be submitted)
- **Scholarship Outcomes:** From Monday, 15 December 2025
- **Acceptance Deadline:** Wednesday, 7 January 2026

It is crucial to prepare your application well in advance of the deadline.

PhD Recruitment (4/4)

Interested candidates should first submit an **Expression of Interest (EOI)** via email to me at **sen.wang@uq.edu.au** with the subject line "PhD Application EOI".

Please include the following in your EOI:

- A **Cover Letter** introducing yourself, outlining your research interests, and explaining why you are a suitable candidate.
- A detailed **Curriculum Vitae (CV)**.
- Your **Academic Transcripts**.

Shortlisted candidates will be contacted for an interview and will be supported in preparing their formal application to the UQ Graduate School.

[Associate Professor Sen Wang | UQ Experts](#)

Recap


- Introduction to Vector Databases
- Fundamentals of Vector Databases
 - What are vectors?
 - Vector Similarity Measures
- Indexing Techniques for Vector Databases
 - Introduction to Indexing
 - Types of Indexing Techniques
- Vector Databases in the Cloud
- Applications:
 - Retrieval Augmented Generation
 - Anomaly Detection

Outline

- • Background
 - Apache Spark and its Characteristics
 - Resilient Distributed Dataset (RDD) and its operations: Transform & Action
 - Lazy evaluation and RDD Lineage graph
 - RDD Persistence and Caching
 - Terms in Spark
 - Directed Acyclic Graph (DAG)
 - Narrow and Wide Dependencies
 - Shuffle
 - How Spark works


Background – Digital Revolution

- Digital Revolution (aka the Third Industrial Revolution), is the shift from **mechanical** and **analogue** electronic technology to **digital** electronics.
- Milestones:
 - Digital computers (the late 1950s to the late 1970s)
 - Personal computers (1970s to 1980s)
 - WWW (1990s)
 - Cellular phones (2000s)
 - IoT and Cloud Computing (2010s)
 - Generative AI (2020s) – chatGPT, falcon, Llama 2, etc.



Background – Big Data Era

- “Big Data” has been in use since 1990s.
- Data sets with sizes beyond the ability of commonly used software tools to **capture, curate, manage, and process** data within a tolerable elapsed time.
- Reasons of Big Data:
 - Hardware development: Storage (more cheaper), CPUs (more cores)
 - Internet bandwidth: 56kbps vs 1000Mbps
 - Data generation:
 - Transactional data (stock, shopping records in Woolworths/Coles)
 - User-centric data (videos/images)
 - Sensor-based data (cameras)



1 Bit	=	0 or 1
8 Bits	=	1 Byte
1024 Bytes	=	1 KB
1024 KBs	=	1 MB
1024 MBs	=	1 GB
1024 GBs	=	1 TB
1024 TBs	=	1 PB
1024 PBs	=	1 EB
1024 EBs	=	1 ZB
1024 ZBs	=	1 YB

Background – Big Data Era

- Characteristics of Big Data: 5Vs:
 - **Volume** (quantity) – from 4.4 trillion gigabytes to 44 trillion (more than doubles every two years).
 - **Variety** (type) – structured vs non-structured
 - **Velocity** (speed)
 - **Veracity** (quality)
 - **Value**
- Data must be **processed** with advanced tools to reveal **meaningful information**.
 - Data mining and machine learning
 - Cloud computing



1 Bit	=	0 or 1
8 Bits	=	1 Byte
1024 Bytes	=	1 KB
1024 KBs	=	1 MB
1024 MBs	=	1 GB
1024 GBs	=	1 TB
1024 TBs	=	1 PB
1024 PBs	=	1 EB
1024 EBs	=	1 ZB
1024 ZBs	=	1 YB

Background – Big Data Technologies

To solve big data problems, two types of technologies are required:

- Distributed storage
 - Huge volumes of data are stored on clusters of storage nodes
 - Distributed file systems & Databases (Cluster Relational DBs & NoSQL DBs)
 - GFS, BigTable (Google) and HDFS/HBase (open-source in Apache Hadoop)
- Distributed computing
 - Clusters of computing nodes process data in a parallel manner
 - Distributed computing models/frameworks
 - MapReduce in Hadoop and Apache Spark
- MapReduce is a framework for processing parallelizable problems across large datasets using a large number of computers (nodes)
 - Map/Reduce
 - More I/O expensive than Spark (Hard disks vs Memories)



Outline

- Background
- • Apache Spark and its Characteristics
- Resilient Distributed Dataset (RDD) and its operations: Transform & Action
- Lazy evaluation and RDD Lineage graph
- RDD Persistence and Caching
- Terms in Spark
- Directed Acyclic Graph (DAG)
- Narrow and Wide Dependencies
- Shuffle
- How Spark works

What Is Apache Spark?

- Apache Spark is an open-source distributed general-purpose cluster-computing framework.
- Spark provides an interface for programming entire clusters with implicit data parallelism and fault tolerance.
- Spark was initially started by Matei Zaharia at UC Berkeley's AMPLab in 2009, and open sourced in 2010 under a BSD license.
- Spark codebase was later donated to the Apache Software Foundation, which is now labelled as one of most active Apache projects.
- Spark aims to be fast: **in-memory computing**
- Spark vs MapReduce
 - Much faster (in-memory, low latency)
 - Wider range of workloads (e.g. iterative algorithms)



Dr Matei Zaharia




Characteristics of Spark – Speed

Apache Spark achieves high performance in terms of processing speed

- Process data mainly in memory of working nodes
- Prevent unnecessary I/O operations on disks
- Use a state-of-the-art Directed Acyclic Graph (DAG) scheduler

Performance:

- Sorting Benchmark in 2014
 - 100TB data
 - 206 nodes in 23 mins (Spark)
 - 2,000 nodes in 72 mins (MapReduce)
 - 3x speed performance and 1/10 resources
- Machine Learning algorithm
 - 100x faster than Hadoop for logistic regression



Characteristics of Spark – Ease of Use

Spark has provided many operators making it easy to build parallel apps.

- *map / reduce*
- *filter / groupByKey / join*
- and more.


Highly accessible with supported languages:

- Scala (interactive, fast)
- Python (interactive, slow)
- Java (non-interactive, fast)
- R and SQL shells (interactive, slow)

Characteristics of Spark – A Unified Stack

Spark contains multiple closely integrated components:

- Spark Core:
 - contains the basic functionality of Spark,
 - **task scheduling**,
 - **memory management**,
 - **fault recovery**,
 - **interacting** with storage systems,
 - and more.
 - Responsible for ***resilient distributed data (RDDs)*** definitions
 - RDDs are Spark's main programming abstraction.
 - RDDs represent a collection of items distributed across many compute nodes (in memory)
 - RDDs are central in Spark programming




Characteristics of Spark – A Unified Stack

Spark contains multiple closely integrated components:

- Spark SQL is for working with structured data.
 - It allows querying data via SQL
 - it supports many sources of data, including Hive tables, JSON, etc.
 - Except for SQL interface to Spark, Spark SQL allows to intermix SQL queries in Python, Java, and Scala, all within a single application

```
spark.sql("SELECT * FROM people").show()
```



Spark Stack

- When running SQL from within another programming language the results will be returned as a **Dataset/DataFrame**.

Characteristics of Spark – A Unified Stack

Features of Spark SQL

- Integration With Spark:
 - Spark SQL queries are integrated with Spark programs, allowing to query structured data inside Spark programs, using SQL or a DataFrame API. Usable in Java, Scala, Python and R.
- Uniform Data Access:
 - DataFrames and SQL support a common way to access a variety of data sources, like Hive, JSON, and JDBC.
 - Joins the data across these different sources. This is very helpful to accommodate all the existing users into Spark SQL.
- Hive Compatibility:
 - Spark SQL runs unmodified Hive queries on current data.
 - It rewrites the Hive front-end and meta store, allowing full compatibility with current Hive data, queries, and UDFs.

Characteristics of Spark – A Unified Stack

Features of Spark SQL

- Standard Connectivity:
 - JDBC or ODBC: JDBC and ODBC are the industry norms for connectivity for business intelligence tools.
- Performance And Scalability:
 - Due to the internal optimizers, Spark SQL can effectively and efficiently read from multiple sources (files, HDFS, JSON/Parquet files, existing RDDs, Hive, etc.) and conduct fast execution of existing Hive queries.
 - Much higher performance compared to Hadoop.






Figure: Runtime of Spark SQL vs Hadoop. Spark SQL is faster

Source: Cloudera Apache Spark Blog

Characteristics of Spark – A Unified Stack

Spark contains multiple closely integrated components:

- Spark Streaming:
 - a component that enables **scalable, high-throughput, fault-tolerant** stream processing of live streams of data.
 - **Logfiles** generated by web servers
 - **Web clicks & Advertising**
 - Internet of Things: **sensors**
- Data can be obtained from **many sources** like Kafka, Flume, or TCP sockets
- Spark Streaming **receives** live input data streams and divides the data into **smaller** batches, which are then **processed** by the Spark engine to generate the final stream of results in batches.
- Final results can be stored in **file systems, databases**, or even moved to live **dashboards** to be displayed



Characteristics of Spark – A Unified Stack

Features of Spark Streaming

- **Dynamic load balancing**
 - Dividing the data into **small micro-batches** allows for **fine-grained allocation** of computations to resources.
- **Fast failure recovery**
 - In Spark, the computation is divided into small tasks that can run **anywhere** without affecting correctness.
 - Failed tasks are **evenly distributed** on all the other nodes in the cluster to perform the recomputations and recover from the failure **faster** than the traditional approach.


Characteristics of Spark – A Unified Stack

- **Interactive analytics on streaming data**
 - Arbitrary Apache Spark **functions** can be applied to **each batch** of streaming data. Since the batches of streaming data are stored in the Spark's **worker memory**, it can be **interactively queried**.
- **Advanced analytics like machine learning and interactive SQL**
 - **Rich libraries** like MLlib (machine learning) and SQL are available to micro-batches data.
 - Machine learning models generated **offline** with MLlib can apply to streaming data.
- **Performance**
 - Spark Streaming's ability to batch data can handle second-level tasks in many scenarios.
 - Spark Streaming can achieve latencies in subsecond (less than 1 sec).

Characteristics of Spark – A Unified Stack

Spark contains multiple closely integrated components:



- Spark MLlib:
 - MLlib is Spark's machine learning (ML) library. Its goal is to make practical machine learning **scalable** and **easy**.
 - ML Algorithms: common learning algorithms such as **classification**, **regression**, **clustering**, **association rule mining**, and **collaborative filtering**
 - Featurization: **feature extraction**, transformation, **dimensionality reduction**, and **selection**
 - some lower-level ML primitives, including a generic **gradient descent optimization** algorithm.
 - Utilities: e.g. model **evaluation** and **data import** tools.
- All of these methods are designed to scale out across a cluster.



Characteristics of Spark – A Unified Stack

Spark contains multiple closely integrated components:

- Spark GraphX:
 - is a new component in Spark for graphs and graph-parallel computation.
 - At a high level, GraphX extends the Spark RDD by introducing a new Graph abstraction:
 - a directed multigraph with properties attached to each vertex and edge.
 - To support graph computation, GraphX exposes a set of fundamental operators (e.g., subgraph, joinVertices, and aggregateMessages).
 - In addition, GraphX includes a growing collection of graph algorithms and builders to simplify graph analytics tasks.



Characteristics of Spark – Runs Everywhere


- Under the hood, Spark is designed to efficiently **scale up** from one to many thousands of compute nodes.
- To achieve this while maximizing flexibility, Spark can run over a variety of *cluster managers*:
 - Hadoop YARN,
 - Apache Mesos,
 - Kubernetes,
 - and a simple cluster manager included in Spark itself (Standalone Scheduler).
- Can access diverse data sources:
 - HDFS, Cassandra, Hbase, Apache Hive, and hundreds of other data sources



Running Spark

Option 1: Create a Cluster using **Docker Swarm / K8s / GKE** and Run Spark on it.

You start with a general-purpose container orchestration platform (your GKE cluster). Then, you deploy and manage your Spark application on it, just like any other containerized microservice. You are responsible for configuring how Spark runs, managing its resources, and setting up its integrations.




You have **strong DevOps and K8s expertise**. Your team is comfortable with kubectl, Docker, and YAML. You require **fine-grained control and customization** over your Spark environment and dependencies.

Running Spark

Option 2: Create a dedicated service (DataProc/EMR/HDInsight) for administering Spark/Hadoop.

You press a button, and Google gives you a perfectly configured, optimized, and ready-to-go cluster for your Spark jobs. It handles the installation, integration with other GCP services (like BigQuery, Cloud Storage), monitoring, and scaling.



Your team (e.g., data scientists, analysts) **wants to focus on** writing Spark code, not managing infrastructure. You need **deep, out-of-the-box integration** with the GCP data ecosystem (BigQuery, Cloud Storage, etc.).

Demo - Spark on DataProc

- Create a DataProc Cluster enabling a specific version of Apache Spark
- Use [NYC Taxi and Limousine Commission \(TLC\) Trips Dataset](#) (big dataset) available in BigQuery
 - tlc_yellow_trips_2022
 - **Table Size:** Approximately **8.85 GB**
 - **Number of Rows:** Approximately **39.7 million rows**
- **Task 1:** Big Data Query and Aggregation – **Spark's SQL**
 - Find the **average trip distance** and **total fare amount** by passenger **count #** for trips with a positive fare amount.
 - Aggregate big data using GROUP BY
- **Task 2:** Machine Learning for Fare Prediction – **Spark's Mllib**
 - Extract feature, train a regression model, and test the model's performance
 - Machine Learning model is applied
- **Both tasks CANNOT be processed on a single laptop due to Big Data!!**



Demo – Create a Dataproc Cluster (1/3)

- Open console and check project ID:

```
uqteaching@cloudshell:~$ gcloud projects list
PROJECT_ID: inf3208-437208
NAME: INFS3208
PROJECT NUMBER: 644867479445
```

- Set a valid project:

```
uqteaching@cloudshell:~$ gcloud config set project inf3208-437208
Updated property [core/project].
```

- Enable the Dataproc, BigQuery, and Cloud Storage APIs:

```
uqteaching@cloudshell:~ (inf3208-437208)$ gcloud services enable dataproc.googleapis.com \
    compute.googleapis.com \
    storage.googleapis.com \
    bigquery.googleapis.com
Operation "operations/acat.p2-644867479445-f197f253-b8f5-4d4f-90bd-410d1612b971" finished successfully.
uqteaching@cloudshell:~ (inf3208-437208)$ █
```

Demo – Create a Dataproc Cluster (2/3)

- Create a GCS Bucket (Cloud Storage): a Cloud Storage bucket is used for staging files and storing job output. Note that the name of the bucket **must be globally unique**.


`gsutil mb gs://your-unique-bucket-name/`

```
uqteaching@cloudshell:~ (inf3208-437208)$ gsutil mb gs://uq-inf3208-dataproc-demo-bucket/  
Creating gs://uq-inf3208-dataproc-demo-bucket/...  
uqteaching@cloudshell:~ (inf3208-437208)$ █
```

- Create a Cluster: `gcloud dataproc clusters create taxi-demo-cluster \`
`--region=australia-southeast1 --num-workers=3 \`
`--worker-machine-type=n2-standard-4 \`
`--master-machine-type=n2-standard-4 \`
`--image-version=2.1-debian11 \`
`--master-boot-disk-size=50GB --worker-boot-disk-size=50GB`

Demo – Create a Dataproc Cluster (3/3)

- After the creation of the cluster, you will see the corresponding VMs in GCE, one master node and three worker nodes.



The screenshot shows the Google Cloud VM instances page. At the top, there are navigation links for 'Google Cloud' and 'INFS3208'. A search bar on the right contains the placeholder 'Search (/) for resources, docs, products and more'. Below the header, there are buttons for 'Create instance', 'Import VM', and 'Refresh'. The main area is divided into tabs: 'Instances' (which is selected), 'Observability', and 'Instance schedules'. Under the 'Instances' tab, the section title 'VM instances' is followed by a 'Filter' input field with the placeholder 'Enter property name or value'. A table lists four VM instances:

Status	Name	Zone	Recommendations	In use by	Internal IP	External IP	Connect
<input checked="" type="checkbox"/>	taxi-demo-cluster-m	australia-southeast1-b			10.152.0.19 (nic0)	34.40.140.194 (nic0)	SSH <input type="button" value="▼"/>
<input checked="" type="checkbox"/>	taxi-demo-cluster-w-0	australia-southeast1-b			10.152.0.21 (nic0)	34.87.221.241 (nic0)	SSH <input type="button" value="▼"/>
<input checked="" type="checkbox"/>	taxi-demo-cluster-w-1	australia-southeast1-b			10.152.0.20 (nic0)	34.87.197.176 (nic0)	SSH <input type="button" value="▼"/>
<input checked="" type="checkbox"/>	taxi-demo-cluster-w-2	australia-southeast1-b			10.152.0.22 (nic0)	35.201.5.188 (nic0)	SSH <input type="button" value="▼"/>

DELETE the cluster to avoid incurring further charges.

```
gcloud dataproc clusters delete taxi-demo-cluster --region=australia-southeast1
```

Demo – Big Data Query and Aggregation (T1)

Create a local Python file named `taxi_query_job.py`:

- Read the NYC Yellow Taxi trip data for the year 2022 from a public BigQuery table.
- Perform a `GROUP BY` operation to aggregate data.
- Write the aggregated result to your GCS bucket in Parquet format.
- Output:

passenger_count	trip_count	avg_distance	total_fare
1	25702970	3.3944167300000	364911278.980000000
2	5298235	3.9966738580000	85679423.800000000
3	1389687	3.8205254920000	22654264.660000000
4	620920	4.0977473590000	10860404.080000000
5	633917	3.4099003020000	8873865.220000000
6	425746	3.4838078340000	6051653.970000000
7	205	1.6563414630000	14402.060000000
8	130	3.8932307690000	9166.670000000
9	39	3.7094871790000	3171.090000000

```
query = """
    SELECT
        passenger_count,
        COUNT(*) as trip_count,
        AVG(trip_distance) as avg_distance,
        SUM(fare_amount) as total_fare
    FROM
        trips
    WHERE
        passenger_count > 0 AND fare_amount > 0
    GROUP BY
        passenger_count
    ORDER BY
        passenger_count
"""

aggregated_data = spark.sql(query)
```

Demo – Machine Learning for Fare Prediction (T2)

Create a local Python file named [taxi_ml_job.py](#):

- Read taxi data from [BigQuery](#).
- Perform data cleaning and feature engineering.
- Train a Gradient-Boosted Tree (GBT) regression model.
- Evaluate the model's performance.

Output:

```
-----  
Model training complete.  
Root Mean Squared Error (RMSE) on test data = 1.9940  
-----
```

```
# --- 2. FEATURE ENGINEERING ---  
# Combine feature columns into a single vector column  
feature_cols = ["passenger_count", "trip_distance", "fare_amount", "PULocationID", "DOLocationID"]  
assembler = VectorAssembler(inputCols=feature_cols, outputCol="features")  
model_df = assembler.transform(model_df)  
  
# --- 3. TRAIN the Model ---  
# Split data into training and testing sets  
(training_data, test_data) = model_df.randomSplit([0.8, 0.2], seed=42)  
  
print(f"Training on {training_data.count()} records, testing on {test_data.count()} records.")  
  
# Create and train the GBT Regressor model  
gbt = GBTRRegressor(featuresCol="features", labelCol="label", maxIter=10)  
model = gbt.fit(training_data)  
  
# --- 4. EVALUATE the Model ---  
# Make predictions on the test data  
predictions = model.transform(test_data)  
  
# Evaluate the model using Root Mean Squared Error (RMSE)  
evaluator = RegressionEvaluator(  
    labelCol="label", predictionCol="prediction", metricName="rmse")  
rmse = evaluator.evaluate(predictions)
```


Outline

- Background
- Apache Spark and its Characteristics
- Resilient Distributed Dataset (RDD) and its operations: Transform & Action
- Lazy evaluation and RDD Lineage graph
- RDD Persistence and Caching
- Terms in Spark
- Directed Acyclic Graph (DAG)
- Narrow and Wide Dependencies
- Shuffle
- How Spark works

Resilient Distributed Dataset (RDD)

Resilient Distributed Dataset (RDD)


- RDD is a fundamental **data structure** of Spark.
- RDD is a **read-only** (i.e. immutable) **distributed** collection of objects/elements.
- **Distributed**: Each dataset in RDD is divided into **logical partitions**, which are computed by many worker nodes (computers) in the cluster.
- **Resilient**: RDD can be **self-recovered** in case of failure (support rebuild if a partition is destroyed).
- **Datasets**: JSON file, CSV file, text file etc.



Resilient Distributed Dataset (RDD)

Resilient Distributed Dataset (RDD)

- In Spark, all work is expressed as
 - **creating** new RDDs or,
 - **transforming** existing RDDs or,
 - **action** on RDDs to compute a result.
- Data manipulation in Spark is heavily based on RDDs.
- RDDs can contain any type of **Python**, **Java**, or **Scala** objects, including **user-defined** classes.
- Spark **automatically distributes** the data contained in RDDs across your cluster and **parallelizes** the operations you perform on them.




RDD Operations

There are two types of RDD Operations: **Transformation** and **Action**


Transformations

- Transformations are operations on RDDs that return a **new RDD**.
- Transformed RDDs are computed lazily (only when you use them in an action)
- Many transformations are **element-wise** (working on one element at a time)
- introduce **dependencies** between RDDs to generate **lineage graph**



Example: a huge logfile (100TB), *log.txt*, with lines of messages, including normal, warning, and error messages. We want to display the warning and error messages.

- use a **filter()** transformation from **inputRDD** to **errorsRDD** (only transform error lines) **Direct Acyclic Graph**
- use a **filter()** transformation from **inputRDD** to **warningsRDD** (only transform warning lines)
- use a **union()** transformation to display the final results.
- **filter()** has **one** ascendant, while **union()** has **two** ascendants.
 - Transformations can actually operate on **any number of RDDs**.



RDD Operations



There are two types of RDD Operations:

Actions

- trigger job **execution** that forces the **evaluation** of all the transformations
- must return a **final value**
- the values of action are stored to **drivers** or to the **external storage system**.
- It brings **laziness** of RDD into motion.

To print out some information about the **badLinesRDD**:

- **count()** action – **returns** the count as a number
- **take()** action – **collects** a number of elements from the RDD



Spark – Transformation

Common Transformations (15+) supported by Spark

Transformation	Meaning
map(func)	Return a new distributed dataset formed by passing each element of the source through a function func.
filter(func)	Return a new dataset formed by selecting those elements of the source on which func returns true.
flatMap(func)	Similar to map, but each input item can be mapped to 0 or more output items (so func should return a Seq rather than a single item).
union(otherDataset)	Return a new dataset that contains the union of the elements in the source dataset and the argument.
sortByKey([ascending], [numPartitions])	When called on a dataset of (K, V) pairs where K implements Ordered, returns a dataset of (K, V) pairs sorted by keys in ascending or descending order, as specified in the boolean ascending argument.
join(otherDataset, [numPartitions])	When called on datasets of type (K, V) and (K, W), returns a dataset of (K, (V, W)) pairs with all pairs of elements for each key. Outer joins are supported through leftOuterJoin, rightOuterJoin, and fullOuterJoin.

Spark – Action

Common Actions (10+) supported by Spark

Action	Meaning
collect()	Return all the elements of the dataset as an array at the driver program. This is usually useful after a filter or other operation that returns a sufficiently small subset of the data.
count()	Return the number of elements in the dataset.
reduce(<i>func</i>)	Aggregate the elements of the dataset using a function <i>func</i> (which takes two arguments and returns one). The function should be commutative and associative so that it can be computed correctly in parallel.
first()	Return the first element of the dataset (similar to take(1)).
take(<i>n</i>)	Return an array with the first <i>n</i> elements of the dataset.
foreach(<i>func</i>)	Run a function <i>func</i> on each element of the dataset. This is usually done for side effects such as updating an Accumulator or interacting with external storage systems.

Outline

- Background
- Apache Spark and its Characteristics
- Resilient Distributed Dataset (RDD) and its operations: Transform & Action
- • Lazy evaluation and RDD Lineage graph
- RDD Persistence and Caching
- Terms in Spark
- Directed Acyclic Graph (DAG)
- Narrow and Wide Dependencies
- Shuffle
- How Spark works

Lazy Evaluation

Transformations on RDDs are evaluated or computed in a **lazy manner**:

- Spark will not begin to execute until it sees an **action**.

Lazy evaluation means that when a transformation on an RDD is called, the operation is **not immediately performed**.

Spark internally **records** metadata to indicate that this operation has been requested.

Spark can decide what **the best way** is to perform a series of transformations that are recorded.


Spark uses lazy evaluation to **reduce** the number of **passes** (storage on disk)

- *Hadoop MapReduce, developers often have to spend a lot of time considering how to group together operations to minimize the number of MapReduce passes.*



RDD Lineage Graph

- Because of the lazy nature of RDD, dependencies between RDDs are logged in a **lineage graph** (or RDD operator graph or RDD dependency graph).
- Lineage graph can be regarded as a **family tree** of a given RDD.
- To get lineage graph:
 - METHOD: `toDebugString:String`
- When you run into an **action**, this local plan is submitted to an optimizer, which is going to do optimization and implement it into a **physical plan** (DAG) containing **stages**.
- Spark logs all transformations with a **graph structure** which can be optimized by graph optimization technology.
- Lineage graph can be used to **re-build** the RDDs



Outline

- Background
- Apache Spark and its Characteristics
- Resilient Distributed Dataset (RDD) and its operations: Transform & Action
- Lazy evaluation and RDD Lineage graph
- • RDD Persistence and Caching
- Terms in Spark
- Directed Acyclic Graph (DAG)
- Narrow and Wide Dependencies
- Shuffle
- How Spark works

RDD Persistence and Caching

- RDD persistence is an **optimization** technique in which **saves** the result of RDD evaluation.
- To significantly reduce computation overhead, we use **RDD persistence** to save the intermediate result so that we can use it further if required.
- Methods:
 - **cache()**: store all the RDD in-memory (default storage level: MEMORY_ONLY - executor's JVM heap as serialized Java/Scala objects).
 - **persist(level)**: can cache in memory, on disk, or off-heap memory.
- It is a key tool for **iterative** and **interactive** algorithms.
 - persistence process **speeds up** the further computation many times.
- The cache memory of the Spark is **fault tolerant**

RDD Persistence and Caching

When to use persistence and caching?

Caching is recommended in the following situations:


- For RDD re-use in **iterative** machine learning applications
 - E.g. in machine learning, some computed dataset is reused in each iteration (gradient-based)
- For RDD re-use in **standalone** Spark applications
 - In standalone Spark applications, multiple actions will perform on the same RDD
- When **too many transformations** on RDD or some computations are **very expensive**
 - Too many transformations will have a very long chain (long list of records)
 - Some expensive computations (transformations) take time to finish
 - caching can help in reducing the cost of **recovery** in the case one executor fails

Outline

- Background
- Apache Spark and its Characteristics
- Resilient Distributed Dataset (RDD) and its operations: Transform & Action
- Lazy evaluation and RDD Lineage graph
- RDD Persistence and Caching
- • Terms in Spark
 - Directed Acyclic Graph (DAG)
 - Narrow and Wide Dependencies
 - Shuffle
 - How Spark works


Terms in Spark

- Job:
 - A piece of code which **reads** some input from HDFS or local, **performs** some computation on the data and writes some output data.
- Stages:
 - Jobs are divided into **stages**.
 - E.g. Map or Reduce stages (similar with Hadoop).
 - Stages are divided based on **computational boundaries**.
- Tasks:
 - Each stage has some **tasks**, one task per partition. One task is executed on one **partition** of data on one **executor**.



Terms in Spark

- Spark Driver (program driver)
 - A separate **process** to execute user applications
 - creates **SparkContext** to schedule jobs execution and negotiate with cluster manager
- Executors
 - run **tasks** scheduled by driver
 - store **computation results** in memory, on disk or **off-heap memory**
 - interact with storage systems
- Master: The machine on which the **Driver program** runs
- Slave: The machine on which the **Executor program** runs



Terms in Spark

- **SparkContext**
 - The first step to create Apache Spark SparkContext, which is the [main entry point](#) to spark functionality
 - Configurable [parameters](#) of SparkContext for applications.
 - Spark use some of them to allocate resources on a cluster by executors (e.g. memory size, and cores).
 - Once SparkContext is created, it can be used to [create](#) RDDs (e.g. [textfile](#) method), broadcast variable, accumulator, and run jobs until SparkContext is stopped.
 - Functionalities:
 - Get the [current status](#) of application and set the [configuration](#)
 - Cancel a job/stage and [closure cleaning](#) in Spark
 - [Access](#) various services and programmable [dynamic](#) allocation (request/kill executors)
 - [Access](#) persistent RDDs and [unpersist](#) RDDs
 - etc.


Cluster Manager

- Apache Spark is an in-memory computing engine to process Big Data and it can run a **cluster**, which has **master** and **slave** nodes (worker nodes).
- Cluster manager is to divide resources across applications and works as an **external service** for acquiring resources on the cluster.
- Spark supports pluggable cluster management, which handles starting executor processes, including:
 - Standalone Cluster Manager
 - Hadoop YARN
 - Apache Mesos



Cluster Manager – Spark Standalone

- Standalone mode is a simple cluster manager **incorporated** with Spark.
 - It makes it **easy** to setup a cluster that Spark itself manages and can run on Linux, Windows, or Mac OSX.
 - Often it is the simplest way to run Spark application in a clustered environment.
- It has **masters** and number of **workers** with configured amount of memory and CPU cores.
- Spark allocates resources based on the **core** and an application will grab all the cores in the cluster by default.
- To check the application, each Apache Spark application has a **Web User Interface**
 - provides **information** of executors, storage usage, running task in the application.
 - **monitors** cluster and job statistics
 - checks detailed **log** output for each job.
 - can **reconstruct** the application's UI after the application exits, if an application has logged event for its lifetime.



The screenshot shows the Apache Spark Web User Interface for a running application. The main page displays the application details:

ID:	app-20160218212811-0000
Name:	Spark shell
User:	jacek
Cores:	Unlimited (2 granted)
Executor Memory:	1024.0 MB
Submit Date:	Thu Feb 18 21:28:11 EST 2016
State:	RUNNING

[Application Detail UI](#)

Executor Summary

ExecutorID	Worker	Cores	Memory	State	Logs
0	worker-20160218212802-10.20.3.164-61455	2	1024	RUNNING	stdout stderr

Cluster Manager – Apache Mesos



- Apache Mesos is an **open-source** cluster manager developed in AMPLab at UC Berkeley.
- Apache Mesos **abstracts** CPU, memory, storage, and other compute resources away from machines (physical or virtual), enabling **fault-tolerant** and **elastic distributed systems** to easily be built and run effectively.
- Mesos runs on **nodes** of a cluster and provides an APIs to applications for managing and scheduling resources.
- Mesos is one of the cluster configurations in which Spark can operate.
- Mesos vs Spark Standalone vs YARN
 - There are differences in terms of **high availability, security, monitoring** (depends on your goals)
 - They have provided **almost all** the features as a cluster manager
 - Standalone is the **easiest** for beginners
 - Mesos has **richer resource scheduling** capabilities (allows fine-grained sharing options)
 - YARN can **share** and **configure** the same pool of cluster resources between all frameworks that are running on YARN

Outline

- Background
 - Apache Spark and its Characteristics
 - Resilient Distributed Dataset (RDD) and its operations: Transform & Action
 - Lazy evaluation and RDD Lineage graph
 - RDD Persistence and Caching
 - Terms in Spark
- • Directed Acyclic Graph (DAG)
- Narrow and Wide Dependencies
 - Shuffle
 - How Spark works

Directed Acyclic Graph (DAG)

- Directed Acyclic Graph (DAG) is a set of vertices and edges
 - **vertices** represent the RDDs;
 - **edges** represent the operation to be applied on RDD.
- DAG is a **finite** directed graph (finite vertices and edges) with **no directed cycles**.
- It contains a sequence of vertices such that every edge is directed from **earlier** to **later** in the sequence.
- DAG operations can do **better global optimization** than other systems like MapReduce.
- On the calling of Action, the created DAG submits to DAG Scheduler which further splits the graph into the **stages** of the job.
- The DAGScheduler splits the Spark RDD into **stages** based on applied transformation.
- DAG vs RDD Lineage Graph: Lineage is a portion of a DAG that leads to the creation of that particular RDD.

Lineage Graph v.s. Directed Acyclic Graph (DAG)

Lineage Graph:

- Captures the sequence of transformations applied to the data.
- A lineage graph is a record of how a particular piece of data was derived. For each RDD (Resilient Distributed Dataset), there's a lineage graph that traces back the sequence of transformations (like `map`, `filter`, etc.) that produced it from the source data.
- The lineage graph aids in recovery during data loss. If a partition of an RDD is lost, Spark can recompute it from the source data using the lineage graph, thus ensuring fault tolerance without data replication.
- **A lineage graph represents individual transformations and their dependencies.**

Directed Acyclic Graph (DAG):

- Models the entire execution plan for a Spark job.
- When an action (like `count`, `collect`, etc.) is called in Spark, it constructs a DAG to model the stages and tasks required for that action. This DAG represents the optimized execution plan for the computation.
- The DAG scheduler divides the DAG into stages. Each stage contains as many transformations as possible that have narrow dependencies. Wide dependencies introduce stage boundaries. This means that transformations which can be done in parallel without shuffling data are grouped into a single stage.
- **The DAG represents the entire job, and Spark executes the job by executing these stages in topological order. If one stage fails, only that stage is recomputed, leveraging the DAG's ability to identify the minimum set of tasks to recompute.**

Outline



- Background
- Apache Spark and its Characteristics
- Resilient Distributed Dataset (RDD) and its operations: Transform & Action
- Lazy evaluation and RDD Lineage graph
- RDD Persistence and Caching
- Terms in Spark
- Directed Acyclic Graph (DAG)
- • Narrow and Wide Dependencies
- Shuffle
- How Spark works

Narrow and Wide Dependencies

There are two types of transformations:

- **Narrow transformation (dependencies)**




- each partition of the parent RDD is used by **at most one** partition of the child RDD
- allow for **pipelined execution** on one cluster node
- failure recovery is **more efficient** as only lost parent partitions need to be recomputed
- *Example: map, flatmap, filter, sample, union, etc.*



Narrow and Wide Dependencies

There are two types of transformations:

- **Wide transformation (dependencies)**
 - multiple child partitions may depend on one parent partition
 - require data from all parent partitions to be available and to be shuffled across the nodes
 - a complete re-computation is needed, if some partition is lost from all the ancestors
 - Example: `groupByKey()` and `reduceByKey()`.



Outline

- Background
- Apache Spark and its Characteristics
- Resilient Distributed Dataset (RDD) and its operations: Transform & Action
- Lazy evaluation and RDD Lineage graph
- RDD Persistence and Caching
- Terms in Spark
- Directed Acyclic Graph (DAG)
- Narrow and Wide Dependencies
- • Shuffle
- How Spark works

Shuffle

Example: Given 100k sales records of 100 stores, how to calculate how many items have been sold for each store?

Solution – Database manner:

- use [groupby](#) with a key dividing the entire data sheet into different groups: e.g. all records of sold items for storeID=1 will be in one group.
- Use [aggregate](#) function (count()) to calculate how many rows within each group: e.g. 10k rows for storeID=1 and 12k rows for storeID=99.
- Return the values as the final results

It's feasible for modern databases, such as MS SQL Server, MySQL, PostgreSQL, etc.

What if we have much more data? E.g. 1,000 times more?

Calculate how many items are sold for each store?


trasactionID	itemID	itemPrice	storeID
00000001	131	13.49	1
00000002	211	55.99	74
00000003	432	109.99	54
...
99999999	123	10.0	41
100000000	411	3.0	99

Shuffle

- In the traditional database, we **divide** the entire data into different groups and **calculate** result of each group.
- In RDD, data are **distributed** across multiple partitions on nodes, traditional groupby operation can cause expensive communication cost.
- In Spark, we can **group** the records with the same key to one partition for further operation (e.g. reduceByKey, groupByKey).
- **Example:** there are 10 worker nodes and each node calculates records of 10 stores: Node1 – Store 1 to Store 10, Node 2 – Store 11 to Store 20, ...
- To do this computation, we need to **re-distribute (shuffle)** data so that it's grouped differently across partitions.
- In Spark, the huge database is distributedly stored in RDDs across partitions in different nodes, some **transformation** operations will trigger shuffle (re-distributing) data.
- Shuffle typically involves copying data across executors and machines, making the shuffle a **complex** and **costly** operation.

Calculate how many items are sold for each store?

trasactionID	itemID	itemPrice	storeID
00000001	131	13.49	1
00000002	211	55.99	74
00000003	432	109.99	54
...
99999999	123	10.0	41
100000000	411	3.0	99





Shuffle

- Shuffle is an **expensive** operation since it involves disk I/O, data serialization, and network I/O.
- Shuffle also generates a large number of **intermediate** files on disk.
- Transformations that ***cause communications across nodes when repartitioning*** will cause a shuffle in Spark:
 - join
 - groupByKey
 - reduceByKey
 - combineByKey
 - etc.
- One direct method to check if the shuffle will occur: **.toDebugString**

you should always assume that all wide transformations will trigger a shuffle.


How Stage Generated

- When the program runs into an **action** (like collect), the graph of transformations is submitted to a DAG Scheduler.
- DAG scheduler aims to **optimize** the performance by dividing transformation graph into different stages.
- DAG scheduler **pipelines** transformation operations together to optimize the graph.
- E.g. Many map operators can be scheduled in a single stage.



How Stage Generated

- When the transformations can trigger **shuffle**, DAG Scheduler will divide the transformations into different stages.
- In the following example, **reduceByKey** triggers shuffle and DAG Scheduler divide the pipeline into **two** stages.



How Stage Generated

The DAG scheduler divides operator graph into stages according to the [types of dependencies](#):

- [Narrow](#) dependency – all the transformations can be [pipelined](#) into one stage;
- [Wide](#) dependency – [divide](#) the transformations into different stages.

The final result of a DAG scheduler is [a set of stages](#).

The stages are passed on to the [Task Scheduler](#).


The [task scheduler](#) launches tasks via cluster manager. (Spark Standalone or Hadoop YARN or Apache Mesos).

The [task scheduler](#) doesn't know about dependencies among stages.

How Stage Generated

Let's look at another complicated example:

- RDDs: A – G
- Partitions: 1 – 19
- Narrow transformations:
 - map, union
- Wide transformation:
 - groupByKey, join




Outline

- Background
- Apache Spark and its Characteristics
- Resilient Distributed Dataset (RDD) and its operations: Transform & Action
- Lazy evaluation and RDD Lineage graph
- RDD Persistence and Caching
- Terms in Spark
- Directed Acyclic Graph (DAG)
- Narrow and Wide Dependencies
- Shuffle
- • How Spark works

How Spark Works

Based on aforementioned RDD contents, let us have a look at the summary of the RDD running process in the Spark architecture

1. **Create** an RDD object;
2. `SparkContext` is responsible for calculating the dependencies between RDDs and **building** DAGs;
3. DAG Scheduler is responsible for decomposing the DAG graph into **multiple stages**, each stage containing multiple tasks,
4. Task scheduler launches tasks to distribute across the worker nodes via **cluster manager** (Standalone or Mesos or YARN). The task scheduler **does not know** about dependencies among stages.




Fault Tolerance

Apache Spark fault tolerance property :

- RDD has a capability of handling if any loss/fault/failure occurs.
- Transformations applied to RDDs:
 - Narrow dependency: if one partition fails, simply re-calculate the corresponding partition in the parent RDD.
 - Wide dependency: if one partition fails, need to re-calculate all computations in the parent RDD.
- lineage graph (logical execution plan)
 - Too long or some transformations are too expensive (checkpoint: write RDDs on Disks)
 - Checkpoint for wide dependent transformation.

Spark Deployment

Deployment architecture comparison



Spark vs. Hadoop

Parameters	Spark	Hadoop
Data Storage	Spark stores data in-memory.	Hadoop stores data on disk.
Fault tolerance	Spark's data storage model, resilient distributed datasets (RDD) guarantees fault tolerance.	It uses replication to achieve fault tolerance.
Line of code	Apache Spark is project of 20,000 Line of code.	Hadoop 2.0 has 1,20,000 Line of code
Speed	It is Faster due to In-memory computation.	It is relatively slower than Spark.
OS Support	<ul style="list-style-type: none">• Linux• Windows• Mac OS	<ul style="list-style-type: none">• Linux
High level language	<ul style="list-style-type: none">• Scala• Python• Java• R	<ul style="list-style-type: none">• Java
Streaming data	Spark can be used to process as well as modify real-time data with Spark streaming.	With Hadoop Map-Reduce one can process batch of stored data.
Machine Learning	Spark has its own set of Machine learning libraries (MLib).	Hadoop requires interface with other Machine learning library. Eg: Apache Mahout.

Pros and Cons of Apache Spark

Advantages:

- Spark is **fast** in data processing (In-memory computation technology).
- Spark has significantly **less** computation resources than MapReduce.
- Spark has **more complicated** computing operations than MapReduce.
- Spark support many languages: **Java**, **Python**, **R** and **Scala**.
- Spark can **easily integrate** with almost all Big data technologies (incorporated with Hadoop ecosystem).
- It is **fault tolerant** (RDD) and easily scalable.
- Spark provides 128-bit **encryption** and SSL support for its network.

<https://youthgiri.com/it-world/advantages-and-disadvantages-of-spark/12936/.html/>

Pros and Cons of Apache Spark

Disadvantages:

- In-memory computing needs **large memory** to store all the data which makes hardware very pricy.
- Spark does **not support** genuine real-time processing (micro-batch). It processes data in the micro-batch which we can make it as small as 1 second (second-level). Apache Storm can perform at millisecond level.
- Spark **doesn't have its own file system**. It uses the file system of other technology like HDFS, Hive, etc. (Small file problem)

Reading Materials

1. <https://data-flair.training/blogs/rdd-lineage/>
2. <http://datastrophic.io/core-concepts-architecture-and-internals-of-apache-spark/>
3. <http://web.utk.edu/~wfeng1/spark/introduction.html>
4. https://www.tutorialspoint.com/apache_spark/index.htm
5. <https://techvidvan.com/tutorials/spark-tutorial/>

Next (Week 9) Topic:

Spark Applications