

# Cloud Computing (INFS3208)

## Lecture 11: Hadoop Ecosystem

Lecturer: Dr Sen Wang

School of Electrical Engineering and Computer Science

Faculty of Engineering, Architecture and Information Technology

The University of Queensland

# Caution

- If your GCP credit runs out, please contact Mr Jason Liang;
- All assignments/practical materials can be implemented/deployed locally;
- No extension request will be considered by the school due to the insufficient GCP credit.
- Read more details about the extension policy and late submission policy

## Submission guidelines

Online Submission.

## Deferral or extension

You may be able to [apply for an extension](#).

The maximum extension allowed is 7 days. Extensions are given in multiples of 24 hours.

Marked assignments with feedback and/or detailed solutions with feedback will be released to students within 14-21 days, where the earlier time frame applies if there are no extensions.

## Late submission

A [penalty](#) of 10% of the maximum possible mark will be deducted per 24 hours from time submission is due for up to 7 days. After 7 days, you will receive a mark of 0.



# Recap


- Background of Distributed File Systems
  - Big Data, Data Centre Technology, and Storage Hardware
- Distributed File System
  - File System
  - Server/Client System
  - Sun's Network File System (NFS)
- Clustered File System (CFS)
  - Google File System (GFS)
  - Hadoop Distributed File System (HDFS)
  - HDFS Shell Commands

# Outline

- • Introduction to Hadoop
  - What is Hadoop & History
  - Hadoop Ecosystem
  - Hadoop Computation Model: MapReduce
    - MapReduce Components & Paradigm
    - MapReduce Workflow
    - MapReduce Examples:
      - Word Count
      - Find the Highest/Averaged Temperature
      - Find Word Length Distribution
      - Find Mutual Friends
      - Inverted Indexing
  - Besides the Core: Hive and Pig

# Search Engines in 1990s

1995



The screenshot shows the Yahoo! homepage from 1995. The main feature is a large purple "YAHOO!" logo. Below it is a search bar with the word "YAHOO!" in it. A red banner on the left side reads "Excite". The page includes links for Mail, News, Sports, Finance, Celeb, Style, Movies, Weather, Answers, and Mobile. At the bottom, there's a sidebar with links for stocks, cars, weddings, and "Excite Seeing Tours".

*MetaCrawler Parallel Web Search Service*

by [Erik Selberg](#) and [Oren Etzioni](#)

Try the new [MetaCrawler Beta!](#)  
 If you're searching for a person's home page, try [Abov!](#)

• Examples • Beta Site • Add Site • About •

Search for:   
 as a Phrase    All of these words    Any of these words

For better results, please specify:  
 Search Region:  Search Sites:

Performance parameters:  
 Max wait:  1 minutes   Match type:

[ [About](#) | [Help](#) | [Problems](#) | [Add Site](#) | [Search](#) ]  
[swebmaster@metacrawler.com](mailto:swebmaster@metacrawler.com)




The screenshot shows the MetaCrawler Parallel Web Search Service interface from 1996. It features a search bar at the top with the placeholder "look for all the words". Below the search bar are dropdown menus for "Date" (set to "in the last week") and "Content" (set to "North America (.com)"). There are also checkboxes for "Include media type" (Image, Audio, Video, Photoshop) and "Return Results" (10, full descriptions). The interface has a green and blue color scheme with various links and advertisements for services like WIRED News, WIRED Magazine, and WIRED Entertainment.

1996

# Google search engines



1998



2025

**Google Search works in three stages:**

**1. Crawling:** Google downloads text, images, and videos from pages it found on the internet with automated programs called crawlers.

**2. Indexing:** Google analyzes the text, images, and video files on the page, and stores the information in the Google index, which is a large database.

**3. Serving search results:** When a user searches on Google, Google returns information that's relevant to the user's query.



Distributed File Systems (DFSs)  
Big Data Computation Models  
(MapReduce)

# What is Hadoop?



- Apache Hadoop's MapReduce and HDFS components were **inspired** by Google papers on **MapReduce** and **Google File System** (GFS).
- is a **collection** of open-source software utilities that facilitate using a **network of** many computers to solve problems involving **massive amounts** of data and computation.
- provides a software framework for distributed storage and processing of big data using the MapReduce programming model.
- originally designed for clusters of **cheap commodity** hardware with fundamental assumptions:
  - Abstract and facilitate the storage and processing of large and/or rapidly growing data sets
    - **Structured** and **non-structured** data
    - **Simple** programming models
  - High **scalability** and **availability**
  - **Fault-tolerance** (failures are common)
  - **Move computation** rather than data


# Hadoop Core and Sub-modules



- The **core** of Apache Hadoop consists of:
  - **Storage**: Hadoop Distributed File System (HDFS),
  - **Computation model**: MapReduce programming model.
- The **base** Apache Hadoop framework is composed of the following modules:
  - **Hadoop Common** – contains libraries and utilities needed by other Hadoop modules;
  - Hadoop Distributed File System (**HDFS**) – a **distributed file-system**;
  - Hadoop **MapReduce** – an implementation of the MapReduce programming model for big data processing.
  - **Hadoop YARN** – a cluster manager;
  - Hadoop **HBase** - a NoSQL
  - Hadoop **Hive** - a data warehouse that supports high-level SQL-like query language

# Additional Packages in Hadoop

- The term **Hadoop** is often used for both **base modules** and **sub-modules** and also the **ecosystem**,
- Some **additional** software packages that can be installed on top of or alongside Hadoop are also included:
  - Apache **Pig**: is a high-level platform for creating programs that run on Apache Hadoop
  - Apache **Hive**: is a data warehouse software project built on top of Apache Hadoop for providing data query and analysis
  - Apache **HBase**: is an open-source, distributed, versioned, non-relational database inspired by Google BigTable
  - Apache **Phoenix**: is an open source, massively parallel, relational database engine supporting OLTP for Hadoop using Apache HBase as its backing store.
  - Apache **Spark**: is an in-memory computing platform
  - etc.



# Why use Hadoop?

- Need to process Multi Petabyte Datasets
- Data may not have strict schema
- Expensive to build reliability in each application
- Nodes fails everyday
- Need common infrastructure
- Very Large Distributed File System
- Assumes Commodity Hardware on heterogeneous OS
- Optimized for Batch Processing

# Brief History of Hadoop

Designed to answer the question: “**How to process big data with reasonable cost and time?**”

Hadoop was created by [Doug Cutting](#) and has its origins in [Apache Nutch](#), an open-source web search engine.

Inspired by Google’s GFS and MapReduce papers, development started and then moved to the new Hadoop subproject in [Jan 2006](#).

The name of Hadoop was named after [a toy elephant](#) of Doug Cutting’s son.




The initial code just includes 5,000 lines of code for HDFS and about 6,000 lines for MapReduce.

## Milestones:

- 2006/April – Hadoop won the sorting competition (1.8T on 188 nodes in 47.9 hours)
- 2013 – Hadoop 1.1.2 and Hadoop 2.0.3 alpha.
- 2014 – Hadoop 2.3 and became top-level Apache Project
- 2019 – Apache Hadoop 3.2
- 2022 – Apache Hadoop 3.3.4
- 2023 – Apache Hadoop 3.3.6
- 2024 – Apache Hadoop 3.4.0
- 2025 – Apache Hadoop 3.4.2




Doug Cutting



# Hadoop Ecosystem

- 2003 – Google File System (GFS)
- 2004 – MapReduce computation model (Google)
- 2006 – Hadoop was born
- 2006/April – HDFS + MapReduce in Hadoop 0.1.0
- 2006/April – Hadoop won sorting competition (1.8T on 188 nodes in 47.9 hours)
- 2006 – 2007 – Yahoo contributed to Hadoop
- 2008/March – HBase released
- 2008/Sept – Pig released
- 2009/June – Sqoop released
- 2010/Oct – Hive released
- 2011/Feb – Zookeeper released
- 2012 – Hadoop YARN released
- 2014/May – Spark released



# Hadoop in the Wild


- Hadoop is in use at most organizations that handle big data:
  - Yahoo!
  - Facebook
  - Amazon
  - Netflix
  - and more...
- Main applications using Hadoop:
  - **Advertisement** (Mining user behavior to generate recommendations)
  - **Searches** (group related documents)
  - **Security** (search for uncommon patterns)

- **Adobe**

- We use Apache Hadoop and Apache HBase in several areas from social services to structured data storage and processing for internal use.
- We currently have about 30 nodes running HDFS, Hadoop and HBase in clusters ranging from 5 to 14 nodes on both production and development. We plan a deployment on an 80 nodes cluster.
- We constantly write data to Apache HBase and run MapReduce jobs to process then store it back to Apache HBase or external systems.
- Our production cluster has been running since Oct 2008.

Big Data and Hadoop Market Size Forecast Worldwide 2017-2022

**Size of Hadoop and Big Data Market Worldwide From 2017 To 2022  
(in billion U.S. dollars)**




# Outline

- Introduction to Hadoop
  - What is Hadoop & History
  - Hadoop Ecosystem
- • Hadoop Computation Model: MapReduce
  - MapReduce Components & Paradigm
  - MapReduce Workflow
  - MapReduce Examples:
    - Word Count
    - Find the Highest/Averaged Temperature
    - Find Word Length Distribution
    - Find Mutual Friends
    - Inverted Indexing
- Besides the Core: Hive and Pig

# MapReduce: A Real-World Analogy

## Coins Deposit




**Mapper:** Categorize coins by their face values  
**Reducer:** Count the coins in each face value *in parallel*

# MapReduce Architecture: Components

- **Job Client:**
  - MapReduce program by users will be submitted to JobTracker via Client
  - Users can display job running status through interfaces in Client
- **JobTracker:**
  - Monitor resources and coordinate jobs
  - Monitor health of all the TaskTrackers (transfer jobs to other nodes once failure found)
  - Monitor execution percentage of jobs and resources availability
- **TaskTracker:**
  - Periodically heartbeat with resource information job execution status to JobTracker
  - Receive and execute commands from JobTracker (start new tasks or kill existing tasks)
- **Task:**
  - Map Task and Reduce Task
  - Initiated by TaskTracker


# MapReduce Architecture: Components

- **Client:**
  - MapReduce program by users will be submitted to JobTracker via Client
  - Users can display job running status through interfaces in Client



# MapReduce Architecture: Components


- **JobTracker:**
  - Monitor resources and coordinate jobs
  - Monitor health of all the TaskTrackers (transfer jobs to other nodes once failure found)



# MapReduce Architecture: Components


- **TaskTracker:**

- Periodically heartbeat with resource information job execution status to JobTracker
- Receive and execute commands from JobTracker (start new tasks or kill existing tasks)




# MapReduce Architecture: Components

- **Task:**
  - Map Task and Reduce Task
  - Initiated by TaskTracker




# MapReduce Architecture Ver 1: Workflow



# MapReduce Paradigm

- Map-Reduce is the data processing component of Hadoop.
- Map-Reduce programs transform lists of input data elements into lists of output data elements.
- A Map-Reduce program will do map and reduce tasks asynchronously
- MapReduce Job (an execution of a Mapper and Reducer across a data set) consists of
  - the input data (submitted by users),
  - the MapReduce Program (program logic written by users),
  - and configuration info.



# Inputs and Outputs

- The MapReduce framework operates exclusively on **<key, value>** pairs;
- The framework views the **input** to the job as a set of **<key, value>** pairs and produces a set of **<key, value>** pairs as the **output** of the job.
- Input and Output types of a MapReduce job:

Function	Input	Output	Note
Map	$\langle k_1, v_1 \rangle$ e.g. $\langle \text{lineNum}, \text{"a b c b c a a c c"} \rangle$	$\text{List}(\langle k_2, v_2 \rangle)$ e.g. $(\langle \text{"a"}, 1 \rangle, \langle \text{"b"}, 1 \rangle, \langle \text{"c"}, 1 \rangle, \langle \text{"b"}, 1 \rangle, \langle \text{"c"}, 1 \rangle, \langle \text{"a"}, 1 \rangle, \langle \text{"a"}, 1 \rangle, \langle \text{"c"}, 1 \rangle, \langle \text{"c"}, 1 \rangle)$	1. Convert splits of data into a list of <b>&lt;key, value&gt;</b> pairs. 2. Each input $\langle k_1, v_1 \rangle$ will output a list of key/value pairs as intermediate results
Reduce	$\langle k_2, \text{List}(v_2) \rangle$ e.g. $\langle \text{"a"}, \langle 1, 1, 1 \rangle \rangle$	$\langle k_3, v_3 \rangle$ e.g. $\langle \text{"a"}, 3 \rangle, \langle \text{"b"}, 2 \rangle, \langle \text{"c"}, 4 \rangle$	The value of $\langle k_2, \text{List}(v_2) \rangle$ in the intermediate result, $\text{List}(v_2)$ , represents the values of the same key $k_2$ .

# Inputs and Outputs

Map takes one pair of data with a type in one data domain, and returns a list of pairs in a different domain:  $\text{Map}(k_1, v_1) \rightarrow \text{list}(k_2, v_2)$

- $k_1$  stands for line number while  $v_1$  stands for contents
- E.g.  $(1, \text{"Hello, world. Welcome to MapReduce world!"}) \rightarrow (\text{"Hello"}, 1), (\text{"world"}, 1), (\text{"Welcome"}, 1), (\text{"to"}, 1), (\text{"MapReduce"}, 1), (\text{"world"}, 1)$
- The Map function is applied in parallel to every pair (keyed by  $k_1$ ) in the input dataset
  - Each line will be applied with Map function
  - $(2, \text{"Hello, MapReduce. Welcome to world!"}) \rightarrow (\text{"Hello"}, 1), (\text{"world"}, 1) \dots$
  - $(3, \text{"Hello, Spark. Spark is better than MapReduce."}) \rightarrow (\text{"Hello"}, 1), (\text{"Spark"}, 1) \dots$
  - (...)
- a list of pairs (keyed by  $k_2$ ) will be generated.  $k_2$  here is a word not a line number.
- After that, the MapReduce framework collects all pairs with the same key ( $k_2$ ) from all lists and groups them together, creating one group for each key.
  - $(\text{"Hello"}, \langle 1, 1, 1 \rangle), (\text{"world"}, \langle 1, 1, 1, 1 \rangle)$ , etc.

# Inputs and Outputs

The Reduce function is then applied in parallel to each group, which in turn produces a collection of values in the same domain:  $\text{Reduce}(k_2, \text{list } (v_2)) \rightarrow (k_3, v_3)$


- E.g. ("Hello", <1, 1, 1>) → ("Hello", 3), ("world", <1, 1, 1, 1>) → ("world", 4)
- Each Reduce call typically produces either one value  $v_3$  or an empty return.

```
function map(String name, String document):
    // name: document name
    // document: document contents
    for each word w in document:
        emit (w, 1)

function reduce(String word, Iterator partialCounts):
    // word: a word
    // partialCounts: a list of aggregated partial counts
    sum = 0
    for each pc in partialCounts:
        sum += pc
    emit (word, sum)
```


# Divide and Conquer

- No communications between Map tasks
- No communications between Reduce tasks
- Need shuffle process to transfer data to reducer



# Example I – Word Count

- Given a file consists of the following words: “**Dear, Bear, River, Car, Car, River, Deer, Car and Bear**” and The overall MapReduce process will look like:




## Example I – word count

- The input is split into three different data-sets that are distributed among the map nodes.
- The words are tokenized and a value of 1 is assigned to each word which is hardcoded. This is done assuming that each word will appear only once.
- After this list of key-value pairs is created where the key is the word and value is the number one. For e.g. the first data-set has three key-value pairs: Bear, 1; River, 1; Dear, 1.
- After the mapping process is completed the shuffling and sorting tasks are executed so that all the tuples that have the same key are combined together and sent to the reducer.
- On completion of the sorting and shuffling tasks, each reducer will have a unique key and the corresponding list of values. For example, Car, [1,1,1]; Deer, [1,1]... etc.
- The reducer will now count the values present in the list of each unique key and generate the key, value pair for that key, where the value will now correspond to the count.
- The output is written in the key, value format to the output file.


# DataFlow

- Data are stored in HDFS across two nodes as input files.
- **InputFormat** defines how these input files are split and read. It selects the files or other objects that are used for input. **InputFormat** creates **InputSplit**.
- **InputSplit** logically represents the data which will be processed by an individual **Mapper**.
- One map task is created for each split; thus the number of map tasks will be **equal** to the number of InputSplits.
- The split is divided into records and each record will be processed by the **mapper**.





# DataFlow

- **RecordReader (RR)** **communicates** with the **InputSplit** and **converts** the data into key-value pairs suitable for reading by the mapper.
- RecordReader (RR) will **actually read** data in HDFS according to the split information and pass key-value pairs to mappers.
- **Mapper** processes each input record (from RecordReader) and generates **new** key-value pair, which is **completely different** from the input pair.
- The output of Mapper is also known as **intermediate output** which is written to the **local** disk.
- The output of the Mapper is **not stored** on HDFS as this is temporary data and writing on HDFS will create unnecessary copies (due to replication setting, e.g. replicas=3).
- Mappers output is passed to the **combiner** for further process



# DataFlow

- **Combiner** (aka mini-reducer) performs local aggregation on the mappers' output, which helps to minimize the data transfer between **mapper** and **reducer** (we will see reducer below).



With combiners

6 key/value  
intermediate data

Improved Overall  
Performance



# DataFlow

- **Partitioner** takes the output from **combiners** and performs partitioning.
- Keys of the intermediate data will be **partitioned** according to hash function.
- All keys within the same partition will go to the same reducer.
- The total number of Partitioner depends on the number of reducers.

```


<"apple",1>
<"able",1>
<"accept",1>
<"boy",2>
<"bear",1>
<"car",1>
<"country",2>
    
```

Hash function or  
Partition function



# DataFlow

- **Shuffling** is the process by which the intermediate output from **mappers** is transferred to the **reducer**.
- Reducer has three primary phases: **shuffle & sort** and **reduce**.
- **OutputFormat** determines the way these output key-value pairs are written in output files.
- The final results will be written in HDFS.
  - Note that the intermediate data from mappers are not immediately written into local storage, but into memory.
  - There is a mechanism, call **spill**, that periodically write intermediate data in memory into disk.



# Revisit Word Count Example:

- Assuming that there are four text files on different nodes, you want to count the word frequency using MapReduce model.
- Text 1 (node 1): the weather is good
- Text 2 (node 2): today is good
- Text 3 (node 3): good weather is good
- Text 4 (node 4): today has good weather


The weather is good

Today is good


Good weather is good

Today has good weather

# Revisit Word Count Example: Map phase



# Revisit Word Count Example: Reduce phase



Hash function: need to distribute records evenly over the partitions

## Example II: Find the Highest Temperature

**Problem:** Find the maximum monthly temperature for **each year** from weather reports

**Input:** A set of records with format as:

<Year/Month, Average Temperature of that month>

- (201707,37.8), (201706,32.2)
- (201508, 32.2), (201607,37.8)
- (201708, 26.7), (201606,26.7)

**Question:** write down the Map and Reduce function to solve this problem


- Assume we split the input by line:

Split 1: (201707,37.8), (201706,32.2)


Split 2: (201508, 32.2), (201607,37.8)

Split 3: (201708, 26.7), (201606,26.7)


# Example II: Find the Highest Temperature




# Example III: Find the Average Temperature (1)



# Example III: Find the Average Temperature (2)



# Example III: Find the Average Temperature (3)



## Example III: Find the Average Temperature (Pseudo code)

- **Map(key, value){**  
    // key: line number  
    // value: tuples in a line  
    for each tuple t in value:  
        Emit(t->year, t->temperature);}
- **Reduce (key, list of values){**  
    // key: year  
    // list of values: a list of <temperature sums, counts> tuples  
    int total\_temp = 0;  
    int total\_count=0;  
    for each v in values:  
        total\_temp= total\_temp+v->sum;  
        total\_count=total\_count+v->count;  
    Emit(key,total\_temp/total\_count);}  
  
    • **Combiner(key, list of values){**  
        // key: year  
        // list of values: a list of monthly temperature  
        int total\_temp = 0;  
        for each v in values:  
            total\_temp= total\_temp+v;  
        Emit(key,<total\_temp,size\_of(values)>);}

# Example IV: Find Word Length Distribution

- Given a set of documents, use Map-Reduce model to find the length distribution of all words contained in the documents.
- Example:
  - Input Doc: {This is a test data for the word length distribution problem}
  - Output: {(12,1), (7,1), (6,1), (4,4), (3,2), (2,1), (1,1)}

Word	Length	Word	Length
this	4	the	3
is	2	word	4
a	1	length	6
test	4	distribution	12
data	4	problem	7
for	3		

# Example IV: Find Word Length Distribution

Map(key, value){

// key: document name

// value: words in a document

for each word w in value:

Emit(length(w), w);}

Reduce(key, list of values){

// key: length of a word

// list of values: a list of words

with the same length

Emit(key, size\_of(values));}

{This is a test data for the word length distribution problem}

Map



(4, this), (2, is), (1, a), ..., (12, distribution), (7, problem)

Combiner



(4, <this, test, data, word>), (2, is), ..., (7, problem)

Reduce



{(12,1), (7,1), (6,1), (4,4), (3,2), (2,1), (1,1)}

## Example V: Find Mutual Friends

Given a group of people on online social media (e.g., Facebook), each has a list of friends, use Map-Reduce to find *mutual friends* between two persons

### People You May Know



Lives in Gold Coast, Queensland



2 mutual friends

 Add Friend

Remove



Research Assistant (RA) at **Nanyang Technological University, Singapore**




and 4 other mutual friends

 Add Friend

Remove

# Example V: Find Mutual Friends

Simple example:



Input:

A  $\rightarrow$  B,C,D

B  $\rightarrow$  A,C,D

C  $\rightarrow$  A,B

D  $\rightarrow$  A,B

MapReduce

Output:

(A ,B)  $\rightarrow$  C,D

(A,C)  $\rightarrow$  B

(A,D)  $\rightarrow$  ..

....


# Example V: Find Mutual Friends

```

Map(key, value){
    // key: person_id
    // value: the list of friends of the person
    for each friend f_id in value:
        Emit(<person_id, f_id>, value);}

Reduce(key, list of values){
    // key: <friend pair>
    // list of values: a set of friend lists related with the friend pair
    for v1, v2 in values:
        mutual_friends = v1 intersects v2;
        Emit(key, mutual_friends);}
  
```

$(A, \langle B,C,D \rangle)$ $\Rightarrow$ $\langle A,B \rangle, \langle B,C,D \rangle$ $\langle A,C \rangle, \langle B,C,D \rangle$ $\langle A,D \rangle, \langle B,C,D \rangle$	$(B, \langle A,C,D \rangle)$ $\Rightarrow$ $\langle B,A \rangle, \langle A,C,D \rangle$ $\langle B,C \rangle, \langle A,C,D \rangle$ $\langle B,D \rangle, \langle A,C,D \rangle$
---	---



# Example V: Find Mutual Friends

## Input:

A -> B,C,D  
B-> A,C,D  
C-> A,B  
D->A,B



## Map:

(A,B) -> B,C,D  
(A,C) -> B,C,D  
(A,D) -> B,C,D  
**(A,B) -> A,C,D**  
(B,C) -> A,C,D  
(B,D) -> A,C,D  
(C,A) -> A,B  
(C,B) -> A,B  
(D,A) -> A,B  
(D,B) -> A,B




## Reduce:

**(A,B)** -> C,D  
(A,C) -> B  
(A,D) -> B  
(B,C) -> A  
(B,D) -> A


*Suggest C&D  
to be Friends ☺*

## Example VI: Inverted Index for Large Collection of Documents

- **Inverted index** is a database index storing a mapping from content (e.g. words) to its locations (e.g. pages) in a document or a set of documents
- The purpose of an inverted index is to allow **fast full-text** searches.
- Inverted index algorithm becomes very challenging for documents on Internet:
  - Documents are in all **different formats**
  - **Millions** of Words and Short Phrases in Documents
  - **Billions** of short or long documents



## Example VI: Inverted Index for Large Collection of Documents



Algorithm:

**Step 1.** For each page, get a list of all words on that page.

**Step 2.** For each word from Step 1, get a list of pages on which it appears.

# Legacy vs. Modern: MapReduce and Its Successors

MapReduce is still in use in 2025, but

- MR is now **largely confined** to legacy systems and specific archival batch-processing contexts, rather than as the backbone for new, modern data analytics pipelines.
- MR strengths **remain relevant** for some organisations, especially those that have long-established MapReduce workflows or operate within sectors slow to migrate (e.g., traditional enterprises, government).
- MR has been increasingly replaced by faster, in-memory, stream-first processing frameworks like **Apache Spark**, which offer greater speed, flexibility, and cloud-readiness for real-time and AI workloads.

## Dominant Modern Alternatives

- **Apache Spark** is now the ***primary platform for distributed data processing***, favoured for its in-memory operations, speed, and comprehensive ecosystem supporting batch, streaming, ML, and interactive analytics.
- **Apache Flink** and **Apache Beam** are used for event-driven and streaming-first workloads.
- **Databricks Delta Engine** and **Snowflake** offer transactional consistency, concurrent access, and high-speed analytics beyond what MapReduce can provide.

# Outline

- Introduction to Hadoop
  - What is Hadoop & History
  - Hadoop Ecosystem
- Hadoop Computation Model: MapReduce
  - MapReduce Components & Paradigm
  - MapReduce Workflow
  - MapReduce Examples:
    - Word Count
    - Find the Highest/Averaged Temperature
    - Find Word Length Distribution
    - Find Mutual Friends
    - Inverted Indexing
- • Besides the Core: Hive and Pig

# Hive and Pig

Hadoop is great for large-data processing!

- But writing Java programs for everything is verbose and slow
- Not everyone wants to (or can) write Java code

Solution: develop higher-level data processing languages and convert to Hadoop jobs

- Hive: HQL is like SQL
  - Query language is HQL, a variant of SQL
  - Tables stored on HDFS as flat files
  - Developed by Facebook, now open-source
- Pig:
  - Pig Latin is a bit like Perl
  - Scripts are written in Pig Latin, a dataflow language
  - Developed by Yahoo!, now open-source
  - Roughly 1/3 of all Yahoo! internal jobs



Pig

Hive

MapReduce

# Hive

- Apache Hive supports analysis of **large datasets** stored in Hadoop's **HDFS** and compatible file systems such as **Amazon S3** filesystem.
- It provides a **SQL-like** query language called HiveQL with schema on read and transparently converts queries to MapReduce
- HiveQL does not strictly follow the full **SQL-92** standard but offers some extensions not found in SQL, such as multitable inserts.
- HiveQL has **full ACID properties**.
- Hive is a data warehouse that is good for **OLAP** rather than **OLTP**

```
1 DROP TABLE IF EXISTS docs;
2 CREATE TABLE docs (line STRING);
3 LOAD DATA INPATH 'input_file' OVERWRITE INTO TABLE docs;
4 CREATE TABLE word_counts AS
5 SELECT word, count(1) AS count FROM
6 (SELECT explode(split(line, '\s')) AS word FROM docs) temp
7 GROUP BY word
8 ORDER BY word;
```

## Online Transactional Processing (OLTP) vs Online Analytical Processing (OLAP)


- **Online transaction processing**, or **OLTP**, is a class of information systems that facilitate and manage transaction-oriented applications, typically for data entry and retrieval transaction processing.
  - Processes **operational** data
  - **Quick responses** to user requests
  - **Concurrently used** by many users
  - **Frequent updates** (UPDATE, INSERT, DELETE) e.g., booking airline tickets.
  - HBase is a NoSQL database for OLTP
- **Online Analytical Processing**, or **OLAP**, is a class of systems that is designed to response to **multi-dimensional analytical (MDA)** queries.
  - **No need** to promptly respond
  - Hive is a data warehouse suitable for OLAP

**Dimension:** Product, Region, Year  
**MDA example:**  
Show the total sales amount for each **product category** in **each region** for **2024**.

# Hive Data Model

Data in Hive organized into :

- **Tables**
  - Analogous to relational tables
  - Each table has a corresponding directory in HDFS
  - Data serialized and stored as files within that directory
- **Partitions**
  - Each table can be broken into partitions
  - Partitions determine distribution of data within subdirectories
- **Buckets**
  - Data in each partition divided into buckets
  - Based on a hash function of the column
  - Each bucket is stored as a file in partition directory




# Hive Data Model - Example

Data in Hive organized into :

/hive/warehouse/Student\_record/EECS

/hive/warehouse/Student\_record



# Apache Pig

- Apache Pig is a **high-level** platform for creating programs that run on Apache Hadoop.
- The language for this platform is called **Pig Latin**.
- Pig can execute its Hadoop jobs in **MapReduce** or **Apache Spark**.
- Pig Latin can be **regarded** as
  - High-level Map/Reduce commands pipeline
  - Very intuitive and friendly to users who dislike Java
- Pig Latin can be **extended** using user-defined functions (UDFs)
  - the user can write in Java, Python, JavaScript, etc
  - then call directly from the language.
- Limitations:
  - No loops and conditions (IF.. THEN)



# Pig Latin Script

```
-- inner join
A = LOAD 'student_data' AS (name: chararray, age: int, gpa: float);
B = LOAD 'voter_data' AS (name: chararray, age: int, registration: chararray, contributions: float);
C = COGROUP A BY name, B BY name;
D = FILTER C BY not IsEmpty(A);
E = FILTER D BY not IsEmpty(B);
F = FOREACH E GENERATE flatten(A), flatten(B);
DUMP F;
```

```
/*
public class IsEmpty extends FilterFunc {

    @Override
    public Boolean exec(Tuple input) throws IOException {
        try {
            Object values = input.get(0);
            if (values instanceof DataBag)
                return ((DataBag)values).size() == 0;
            else if (values instanceof Map)
                return ((Map)values).size() == 0;
            else {
                int errCode = 2102;
                String msg = "Cannot test a " +
                DataType.findTypeName(values) + " for emptiness.";
                throw new ExecException(msg, errCode, PigException.BUG);
            }
        } catch (ExecException ee) {
            throw ee;
        }
    }
}
```

# Java vs. Pig Latin




**Performance on par with raw Hadoop!**

# Hive v/s Pig

- **Similarities:**
  - Both High level Languages which work on top of map reduce framework
  - Can coexist since both use the under lying HDFS and map reduce
- **Differences:**
  - **Language**
    - Pig is a procedural ; (A = load 'mydata'; dump A)
    - Hive is Declarative (select \* from A)
  - **Work Type**
    - Pig more suited for adhoc analysis (on demand analysis of click stream search logs)
    - Hive a reporting tool (e.g. weekly Business Intelligence reporting)
  - **Users**
    - Pig – Researchers, Programmers (build complex data pipelines, machine learning)
    - Hive – Business Analysts

# Head-to-Head (the Bee, the Pig, the Elephant)

## Hive, PIG and Hadoop benchmark



Version: Hadoop – 0.18x, Pig:786346, Hive:786346

# Guest Lecture in Week 12

## Cloud Computing (INFS3208)

### Lecture 12: Security & Privacy

Guest Lecturer



**Cheng Jiang**  He/Him · 2nd  
Lead Platform Engineer @ Dabble  
Greater Brisbane Area · [Contact info](#)

#### Part I

- Guest Lectures' exclusive content

#### Part II

- Security in Cloud Computing
- Privacy in Cloud Computing

# References

1. <https://data-flair.training/blogs/hadoop-mapreduce-tutorial/>
2. CSE 40822/60822, Cloud Computing, <https://www3.nd.edu/~dthain/courses/cse40822/fall2018/>
3. MapReduce: Simplified Data Processing on Large Clusters
4. Hadoop Tutorial (<https://www.tutorialspoint.com/hadoop/index.htm>)